


# SIG-Scheduling Deep Dive

Da Ma (Huawei, @k82cn)  
Wei Huang (IBM, @Huang-Wei)

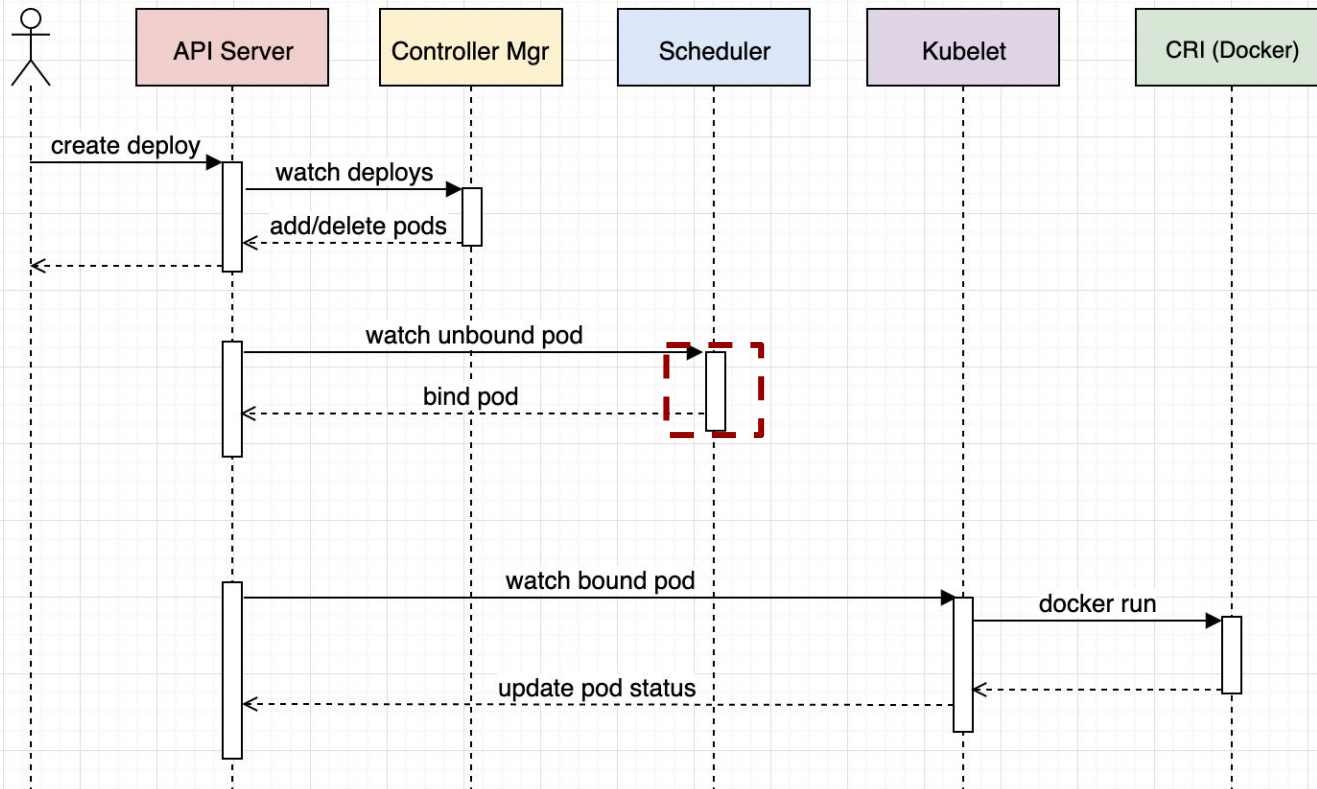


kubernetes

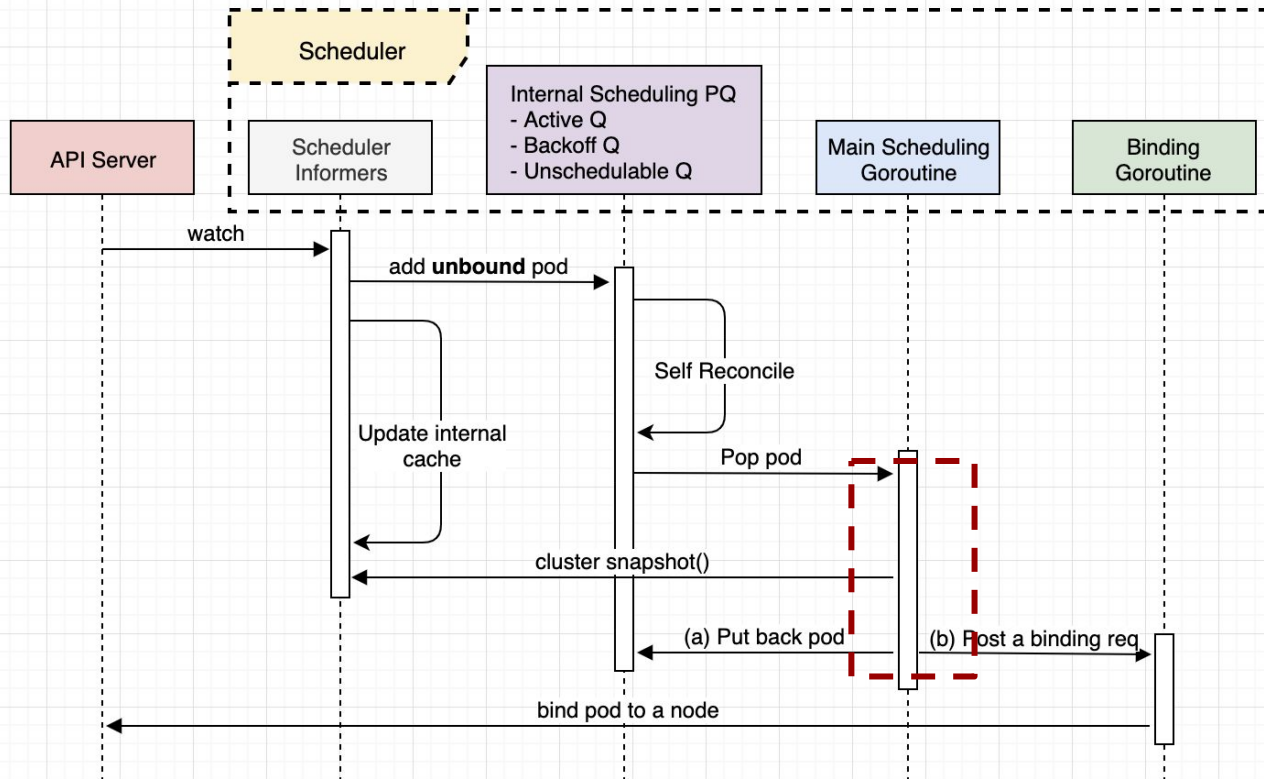
# Scheduler Overview




# Scope of Scheduler



# Detailed Scheduling Flow



# Predicates - filter out nodes



kubernetes

# Predicates



For incoming pod, gives a **YES** or **NO** answer whether it fits on a node or not

```
For each node, check if the node fits the incoming pod (running in
Parallel)
  For each sorted predicate: (running in batch)
    Run the predicate function:
      If succeeded, continue to next predicate
      If failed, record fail reason, break 1
    If failure can be resolved by Preemption, proceed with Preemption
  Aggregate fail reason(s), and return if it's a fit
Filter out failed nodes, and get a node list that fits for incoming node
```

<sup>1</sup> Prior to 1.10, the behavior is continue to next predicate. Starting 1.10, the default behavior was optimized to “break” the predicates loop (see PR [56926](#)). And it’s configurable by parameter “AlwaysCheckAllPredicates” - which is useful for debugging.




# Preemption (in Predicates phase)



```
When a pod can't be scheduled, starts to preempt: blue == dry-run operation)
For each node, delete pods which has lower priority. Based on this state, check:
  If incoming pod cannot be scheduled, return;
  If incoming pod can be scheduled:
    For each pod with lower priority than incoming pod, start with higher one:
      Add the pod. Check:
        If incoming pod can still be placed, keep current pod. Continue
        If not, remove current pod, and add current pod to a "candidates" list
    Return the "candidates" list
If len(candidates) !=0, pick a node which a series of candidate victims belong to,
and talk to APIServer to "reserve" that node (set NominatedNode) for incoming pod,
then delete the candidate victims.
```



# Priorities - rank the remaining nodes



kubernetes



# Priorities



For incoming pod, give a score (0~10) on each filtered node after Predicates phase. Each priority is defaulted weight 1.

Implemented using map/reduce pattern.

```
For each filtered node:
  For each priority:
    Calculate a score, then multiple its weight
  Final score for current node = sum(score * weight)
Get a final score list on all nodes
Pick up the node which has highest final score.
```


# Design Rationale of Scheduler



1. The scheduler is **NOT** responsible for managing life cycle of Pods.
2. The minimum scheduling unit is **POD** (tried EquivalenceCache, but not good as supposed to be)
3. Schedule **one pod** at a time (scheduler can have multiple replicas, but only one leader is running)
4. **Best Fit** vs. First Fit
5. **Predicates** and **Priorities**
6. **Configurable** (schedule config file)
7. **Pluggable** (new scheduler framework, scheduler extender, multiple schedulers)



# Recent Developments



# Default Scheduler



Performance improvements

GA: Priority and Preemption

Planned features

- Even Pod Spreading
- Scheduler Framework

2019 KubeCon EU Deep Dive by Bobby Salamat - <https://youtu.be/t189URLpG8E>

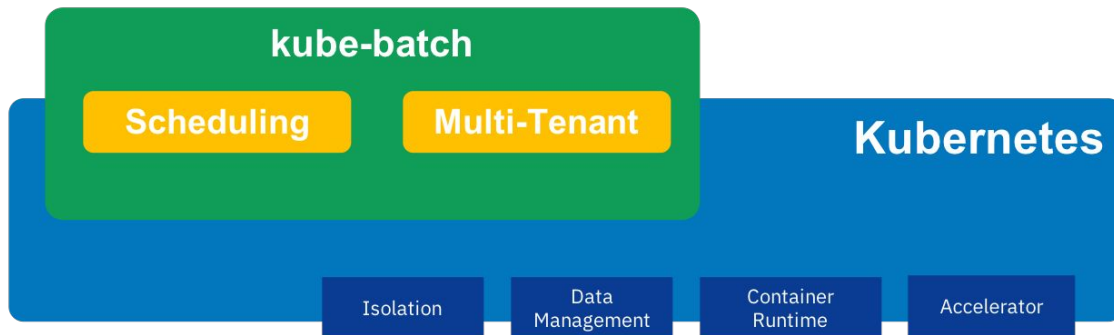
# Overview of kube-batch



Infra

**kube-batch** focus on:

- “Batch” scheduling
- Resource sharing between multi-tenant



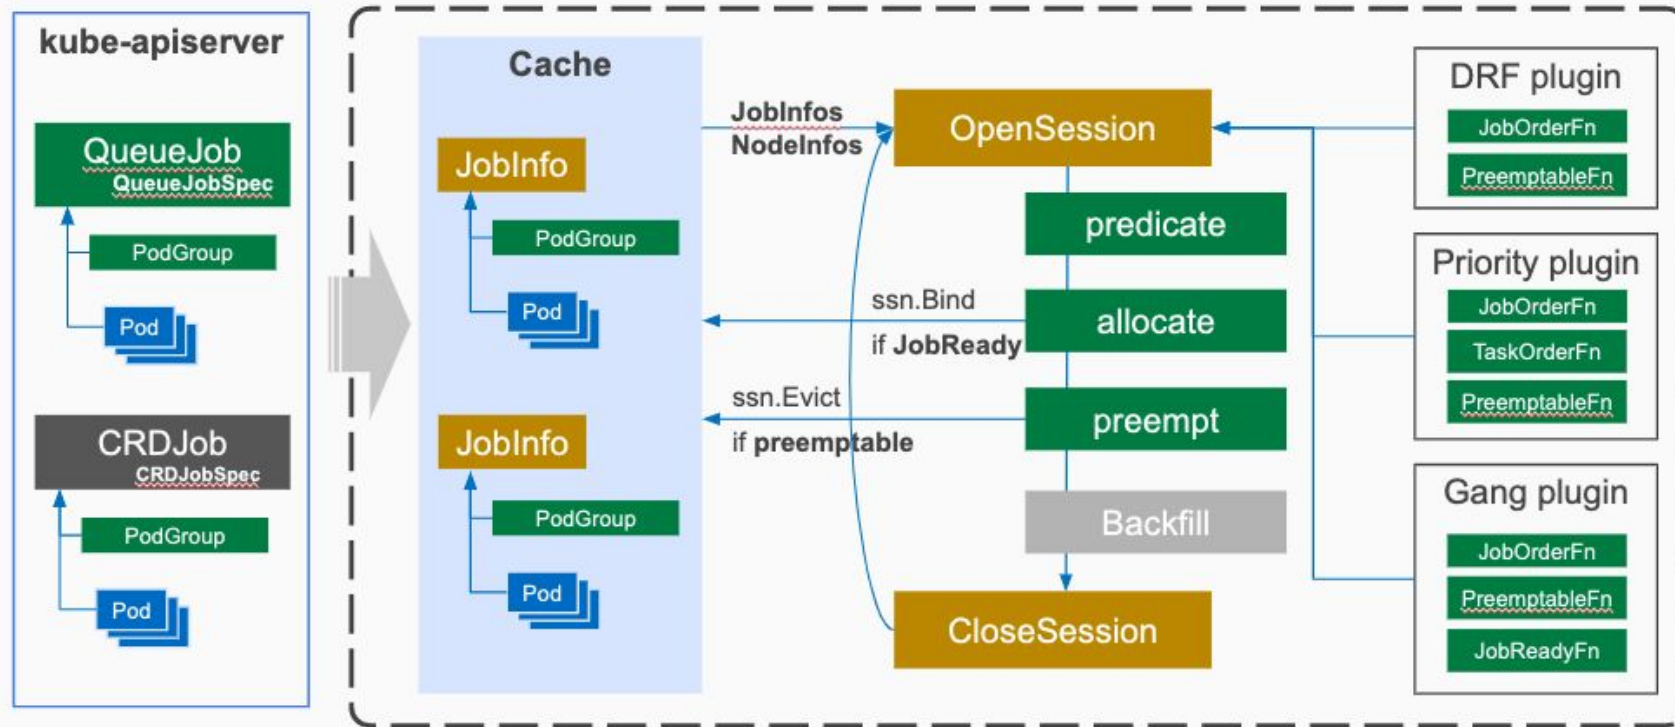
kube-batch **NOT** support:

- Data Management
- Accelerator (Kubelet), e.g. GPU
- Isolation for multi-tenant
- Job Management
- New container runtime, e.g. Singularity, Charis Cloud



kubernetes

# Overview of kube-batch (cont.)



Re-construct JobInfo in Cache by  
PodGroup

Predicate, allocate, preempt are  
Actions, and they're **pluggable**

Plugins on demand

ates

# Features of kube-batch



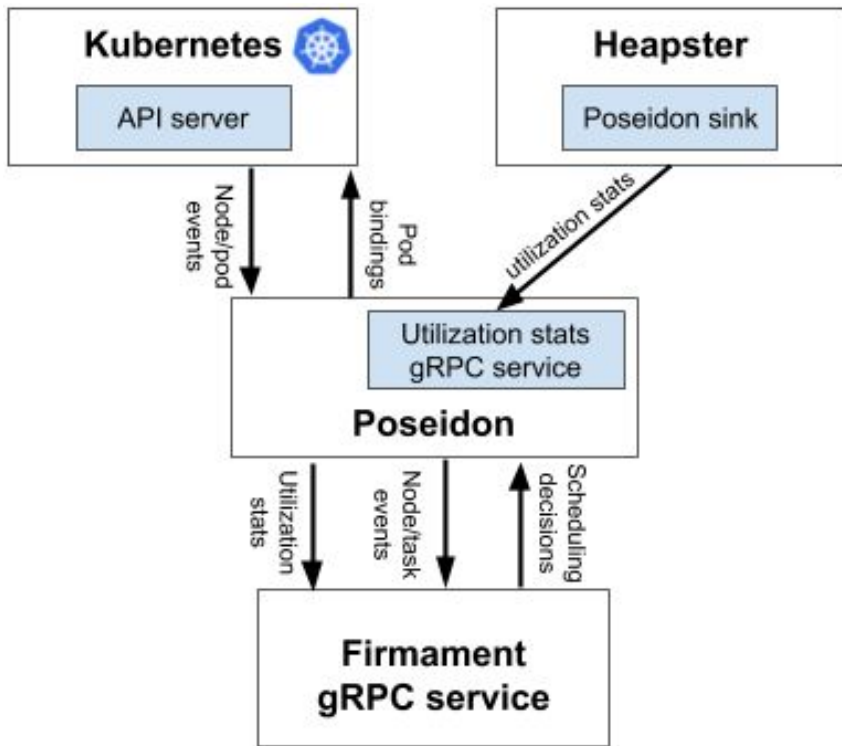
- Co-scheduling
- “Fair-sharing” (job/queue)
- Preemption/Reclaim
- Task Priority within Job
- Predicates
- Queue
- Backfill (partially)
- Dynamic configuration

**Batch Capability into Kubernetes (#68357)**



kubernetes

# Poseidon



Poseidon/Firmament scheduler augments the current Kubernetes scheduling capabilities by incorporating a new novel flow network graph based scheduling capabilities alongside the default Kubernetes Scheduler.

Firmament models workloads on a cluster as flow networks and runs min-cost flow optimizations over these networks to make scheduling decisions.



# Features of Poseidon



1. Node level Affinity and Anti-Affinity
2. Pod level Affinity and Anti-Affinity
3. Taints & Tolerations
4. Gang Scheduling






How those schedulers  
**work together ???**




**Sorry, I don-t know :(**

# Multi-Schedulers



Option 1




Option 2

# Multi-Schedulers



**Option 3**



**Option 4**



## Trigger Of Pod Movement/Migration

**Eviction** -> Creation -> Re-schedule



# User Cases of Descheduler



- Some nodes are under or over utilized.
- The original scheduling decision does not hold true any more, as taints or labels are added to or removed from nodes, pod/node affinity requirements are not satisfied any more.
- Some nodes failed and their pods moved to other nodes.
- New nodes are added to clusters.



# Policy & Strategy



- RemoveDuplicates
- LowNodeUtilization
- RemovePodsViolatingInterPodAntiAffinity
- RemovePodsViolatingNodeAffinity



# Pod Eviction Restriction



- Critical pods (with annotations `scheduler.alpha.kubernetes.io/critical-pod`) are never evicted.
- Pods (static or mirrored pods or stand alone pods) not part of an RC, RS, Deployment or Jobs are never evicted because these pods won't be recreated.
- Pods associated with DaemonSets are never evicted.
- Pods with local storage are never evicted.
- Best efforts pods are evicted before Burstable and Guaranteed pods.
- Pod are never evicted If violates its PDB



# Contact Us



# Contact Us



## Chairs

- @bsalamat
- @k82cn

**Home page:** <https://github.com/kubernetes/community/tree/master/sig-scheduling>

**Slack channel:** <https://kubernetes.slack.com/messages/sig-scheduling>

**Mail list:** <https://kubernetes.slack.com/messages/sig-scheduling>

## Google doc:

[https://docs.google.com/document/d/13mwye7nvrnV11q9\\_Eg77z-1w3X7Q1GTbslpml4J7F3A/view](https://docs.google.com/document/d/13mwye7nvrnV11q9_Eg77z-1w3X7Q1GTbslpml4J7F3A/view)

Thanks!

Q & A

