



# Falco Deep Dive: Extending Event Sources for Kubernetes Audit

---

Kaizhe Huang, Sysdig

# About me.

**Kaizhe Huang**

Security Researcher, Sysdig

 @kaizhe

 @kaizhe



# Join the community.

## Website

- <https://falco.org>

## Public Slack

- <http://slack.sysdig.com/>
- <https://sysdig.slack.com/messages/falco>

## Blog

- <https://sysdig.com/blog/tag/falco/>

## Workshop

- <https://github.com/falcosecurity/falco-security-workshop>

## Github

- <https://github.com/falcosecurity/falco>

## Documentation

- <https://github.com/falcosecurity/falco/wiki>

## Docker Hub

- <https://hub.docker.com/r/falcosecurity/falco/>

# Falco Team



@leodido

Leonardo



@fntlnz

Lorenzo




@mstemm

Mark



@ldegio

Loris



@mfdii

Ducky

 @mfdii





# Help Wanted

- **Profiles Library**
  - Expand predefined rule sets for common use cases and applications
- **Prometheus Metrics Exporter**
  - Expose metric counts for rules fired, etc
- **GRPC Based Outputs**
  - Streaming based outputs for alerts
- **Helm Chart**
  - General maintenance
- **Kubernetes Operator**
  - Add features to improve operator functionality, developer defined rules, etc
- **New Event Sources**
  - Container runtime events, application level events, more



# Falco.

## **A behavioral activity monitor**

- Detects suspicious activity defined by a set of rules
- Uses sysdig's flexible and powerful filtering expressions

## **With full support for containers/orchestration**

- Utilizes sysdig's container & orchestrator support

## **And flexible notification methods**

- Alert to files, standard output, syslog, programs

## **Open Source**

- Anyone can contribute rules or improvements

# Falco: a CNCF sandbox project.

## Runtime security for cloud native platforms

- Detect abnormal behavior in applications, containers, and hosts.
- Audit orchestrator activity.

## Cloud Native Computing Foundation (CNCF)

- Sandbox level project
- [sysdig.com/blog/falco-cncf-sandbox](https://sysdig.com/blog/falco-cncf-sandbox)



# Example.

```
sysdig:~ $ sudo falco
Sat Jan 13 07:11:15 2018: Falco initialized with configuration file /etc/falco/falco.yaml
Sat Jan 13 07:11:15 2018: Parsed rules from file /etc/falco/falco_rules.yaml
Sat Jan 13 07:11:15 2018: Parsed rules from file /etc/falco/falco_rules.local.yaml
```

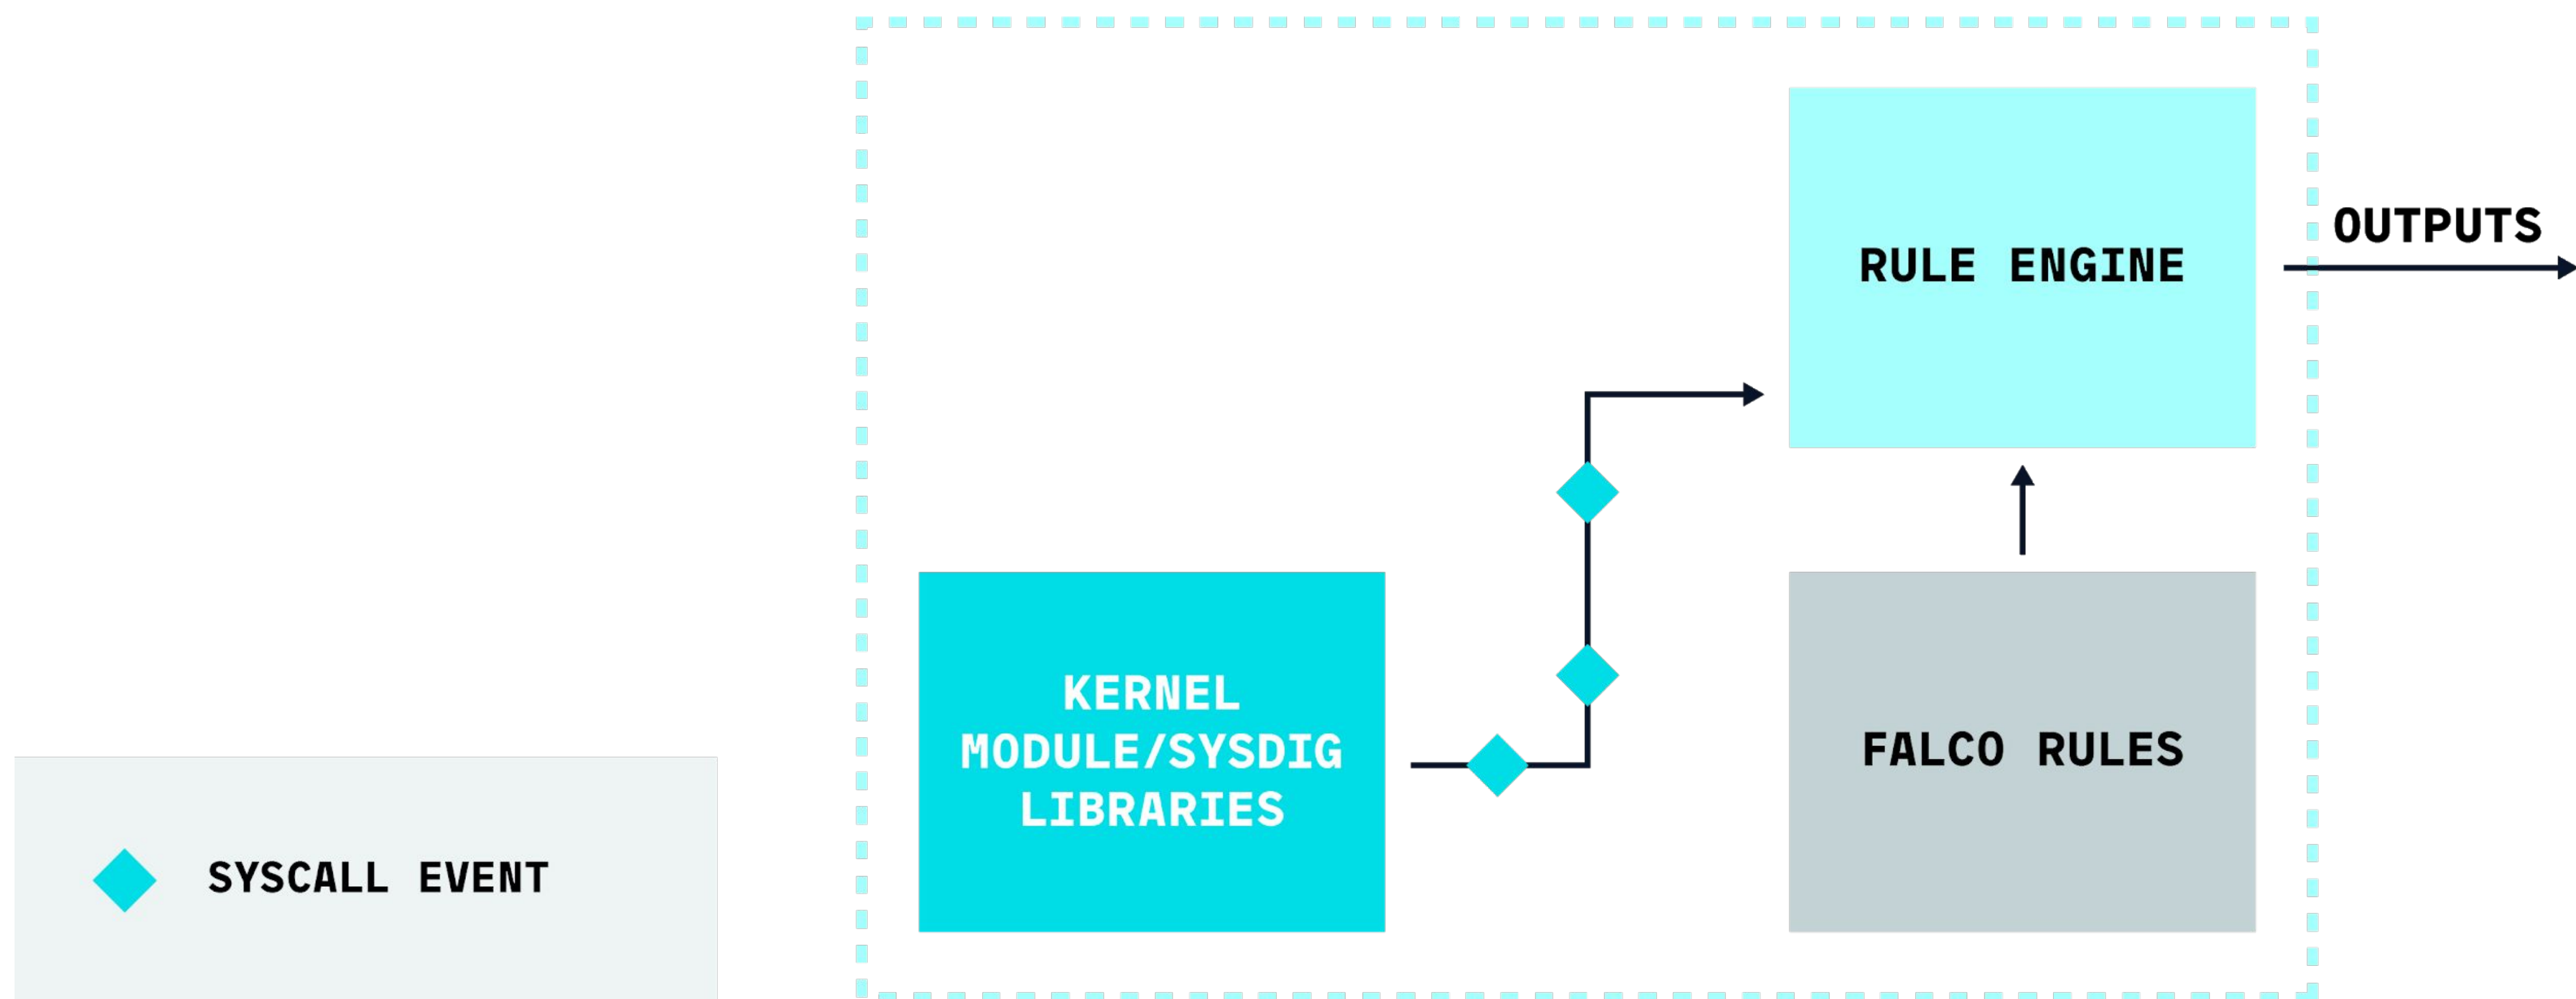
0 Falco

```
sysdig:~ $
```

1 Falco



# Falco internal architecture.



# Additional event sources.

Falco extracts values from events using *filter fields*


```
- rule: Write Below Binary Dir
  desc: An attempt to write to any file below a set of binary directories
  condition: evt.type=open and evt.is_open_write=true and fd.directory in
(/bin, /usr/bin)
  output: >
    File below a known binary directory opened for writing (user=%user.name
    command=%proc.cmdline file=%fd.name)
  priority: WARNING
```



# Additional event sources.

Falco extracts values from events using *filter fields*

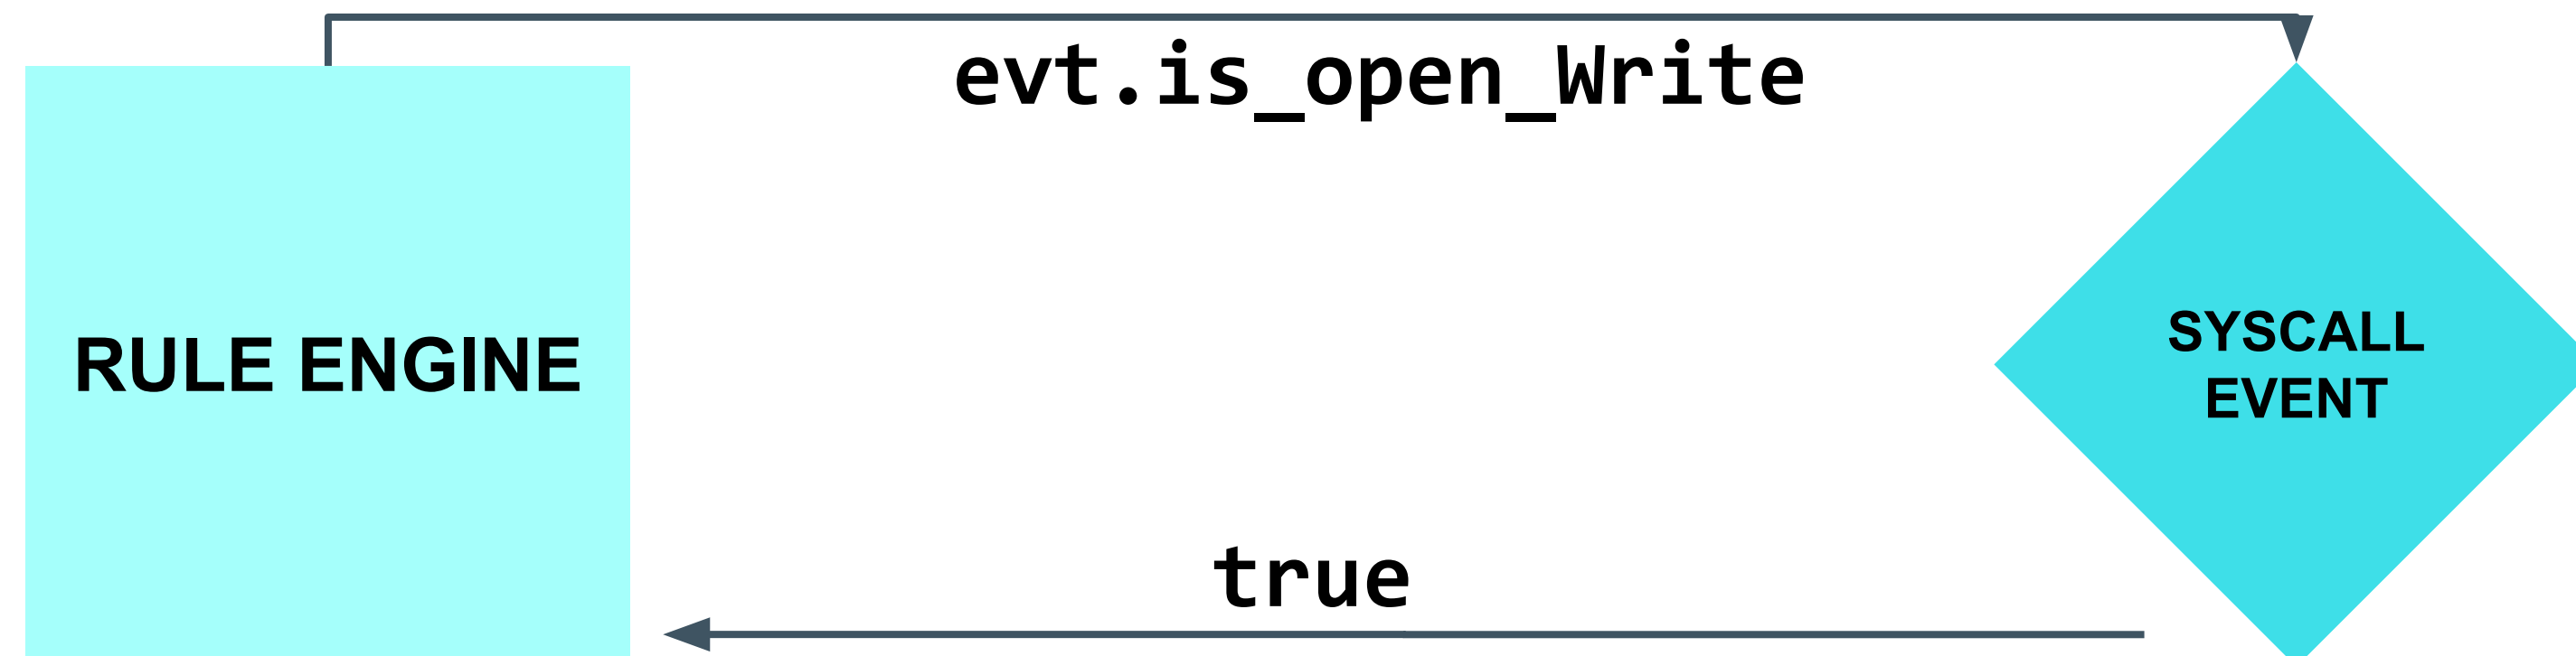
```
- rule: Write Below Binary Dir
  desc: An attempt to write to any file below a set of binary directories
  condition: evt.type=open and evt.is_open_write=true and fd.directory in
(/bin, /usr/bin)
  output: >
    File below a known binary directory opened for writing (user=%user.name
    command=%proc.cmdline file=%fd.name)
  priority: WARNING
```



# Additional event sources.

Falco extracts values from events using *filter fields*

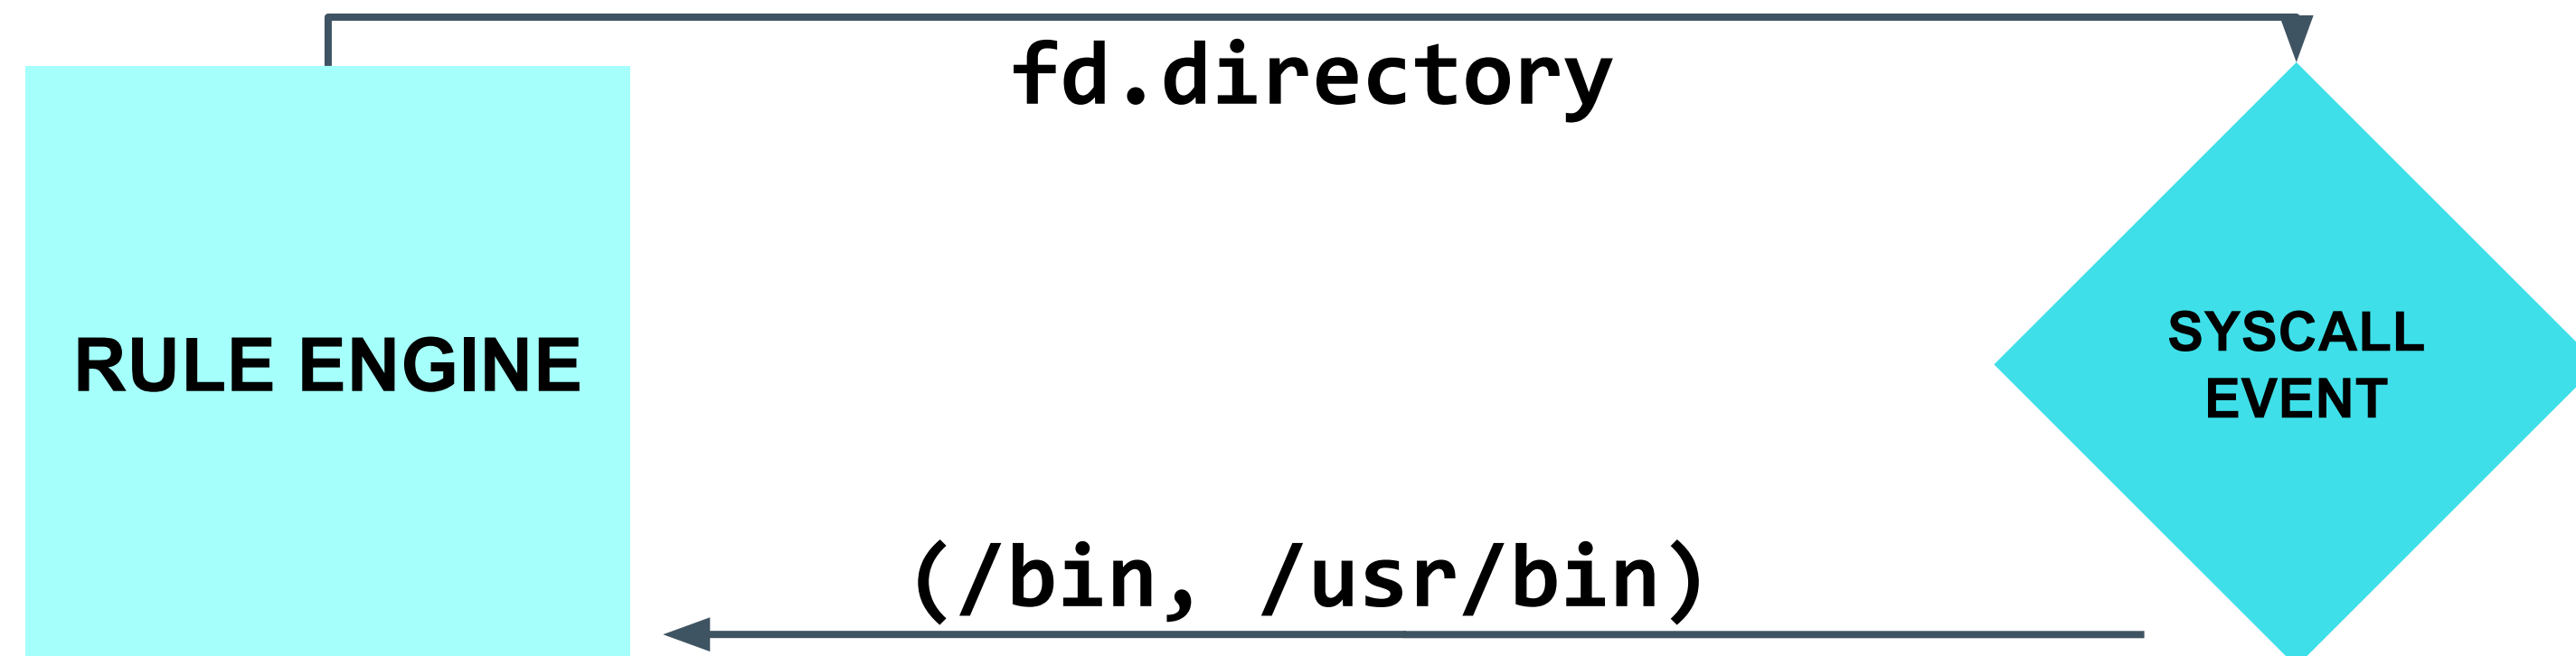
```
- rule: Write Below Binary Dir
  desc: An attempt to write to any file below a set of binary directories
  condition: evt.type=open and evt.is_open_write=true and fd.directory in
  (/bin, /usr/bin)
  output: >
    File below a known binary directory opened for writing (user=%user.name
    command=%proc.cmdline file=%fd.name)
  priority: WARNING
```



# Additional event sources.

Falco extracts values from events using *filter fields*

```
- rule: Write Below Binary Dir
  desc: An attempt to write to any file below a set of binary directories
  condition: evt.type=open and evt.is_open_write=true and fd.directory in
(/bin, /usr/bin)
  output: >
    File below a known binary directory opened for writing (user=%user.name
    command=%proc.cmdline file=%fd.name)
  priority: WARNING
```






# Additional event sources.

Falco extracts values from events using *filter fields*

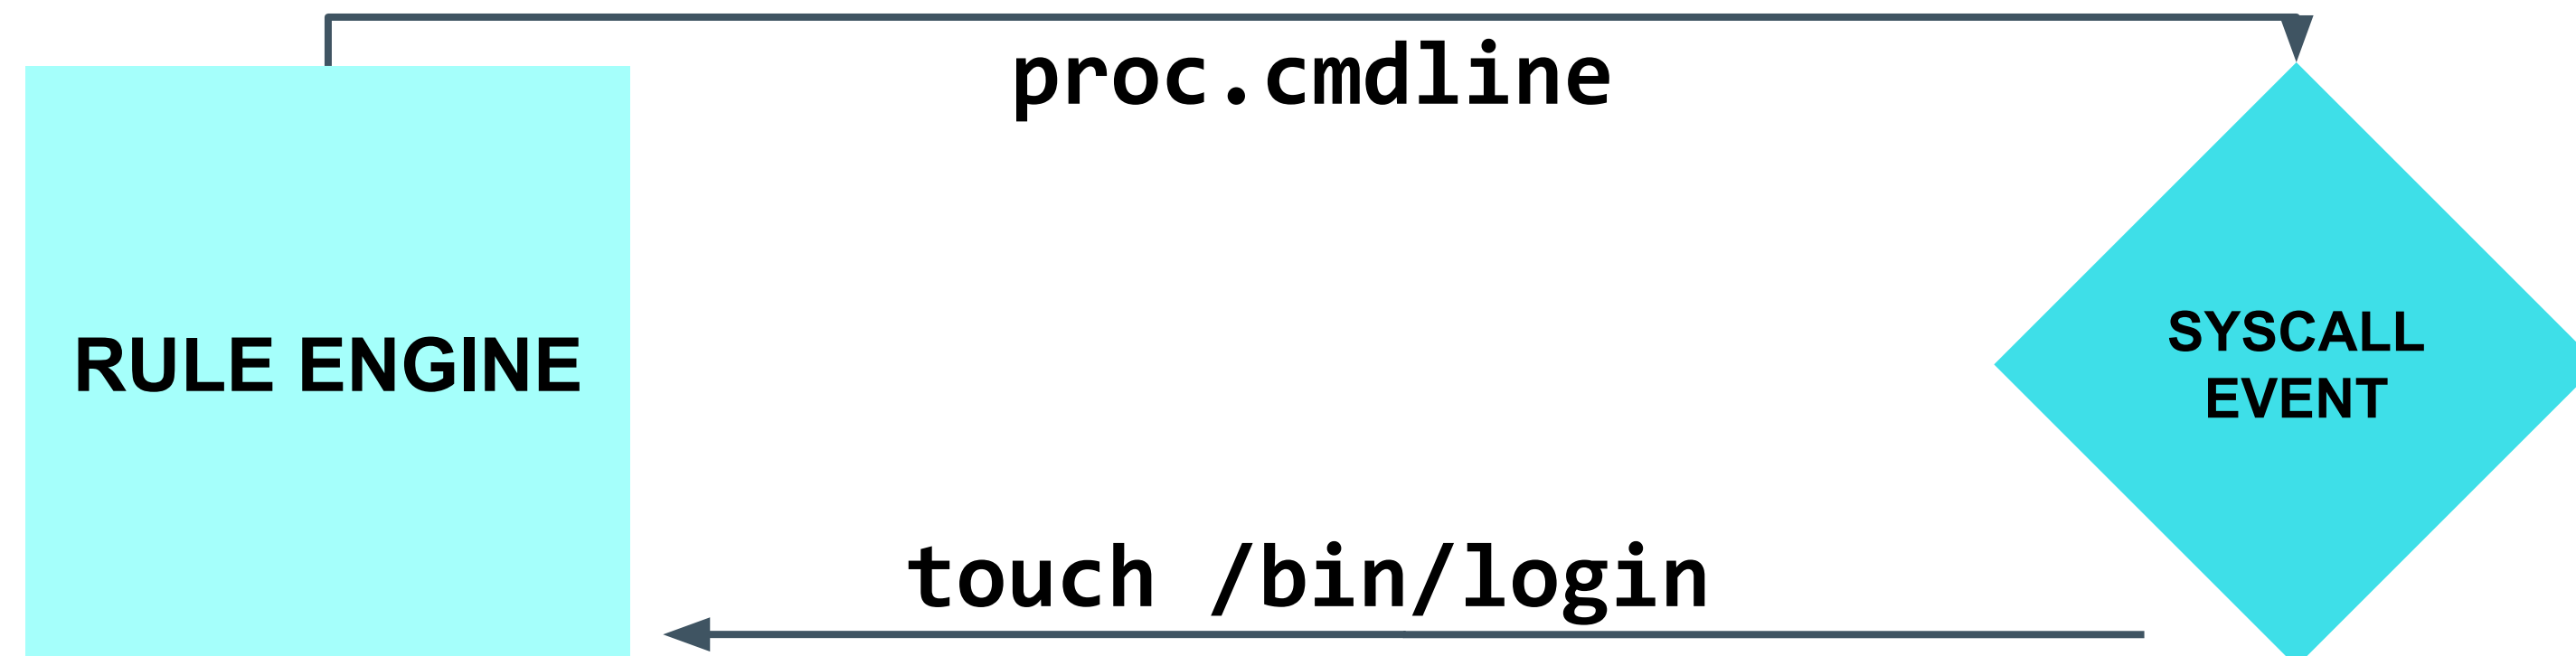
```
- rule: Write Below Binary Dir
  desc: An attempt to write to any file below a set of binary directories
  condition: evt.type=open and evt.is_open_write=true and fd.directory in
(/bin, /usr/bin)
  output: >
    File below a known binary directory opened for writing (user=%user.name
    command=%proc.cmdline file=%fd.name)
  priority: WARNING
```



# Additional event sources.

Falco extracts values from events using *filter fields*

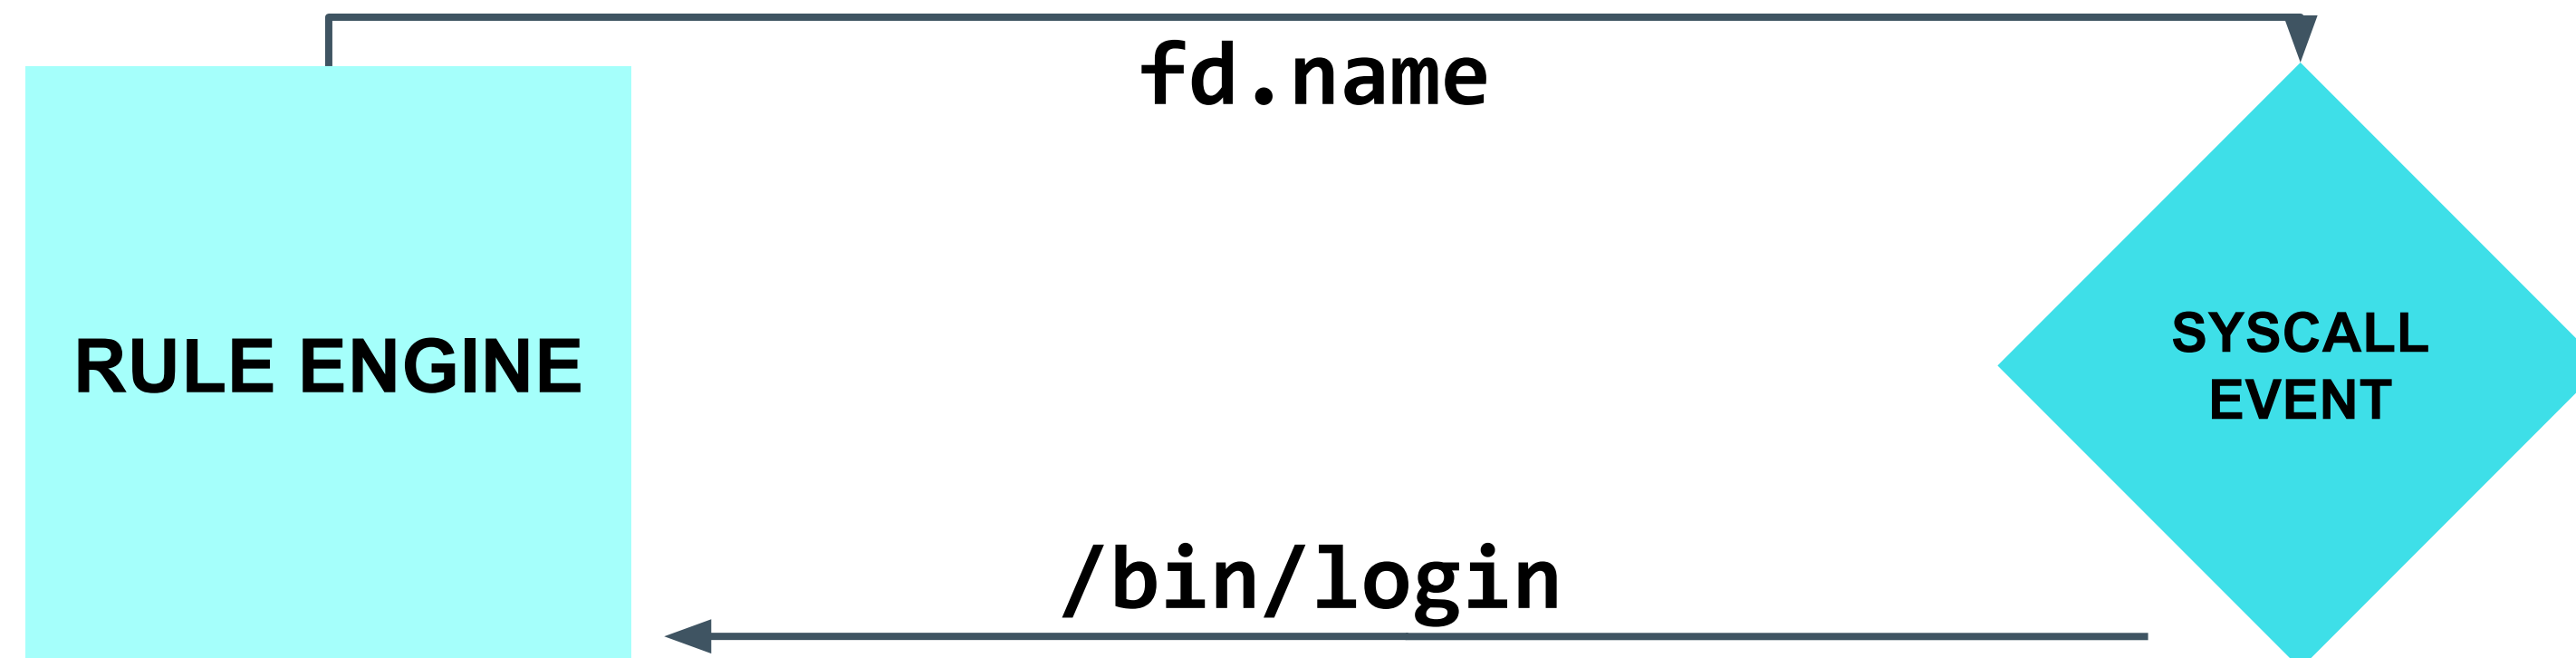
```
- rule: Write Below Binary Dir
  desc: An attempt to write to any file below a set of binary directories
  condition: evt.type=open and evt.is_open_write=true and fd.directory in
(/bin, /usr/bin)
  output: >
    File below a known binary directory opened for writing (user=%user.name
    command=%proc.cmdline file=%fd.name)
  priority: WARNING
```



# Additional event sources.

Falco extracts values from events using *filter fields*

```
- rule: Write Below Binary Dir
  desc: An attempt to write to any file below a set of binary directories
  condition: evt.type=open and evt.is_open_write=true and fd.directory in
(/bin, /usr/bin)
  output: >
    File below a known binary directory opened for writing (user=%user.name
    command=%proc.cmdline file=%fd.name)
  priority: WARNING
```



# Additional event sources.

```
- rule: Write Below Binary Dir
  desc: An attempt to write to any file below a set of binary directories
  condition: evt.type=open and and evt.is_open_write=true and fd.directory in
  (/bin, /usr/bin)
  output: >
    File below a known binary directory opened for writing (user=%user.name
    command=%proc.cmdline file=%fd.name)
  priority: WARNING
```

Falco can be extended to work  
on other “kinds” of events



# Kubernetes audit events.

- New in K8s v1.11
- Provides chronological set of records documenting changes to cluster
- Each record is a JSON object
- Audit policy controls which events are included in event log
- Log backend controls where events are sent
  - Log file
  - Webhook
  - AuditSink (alpha as of 1.13)

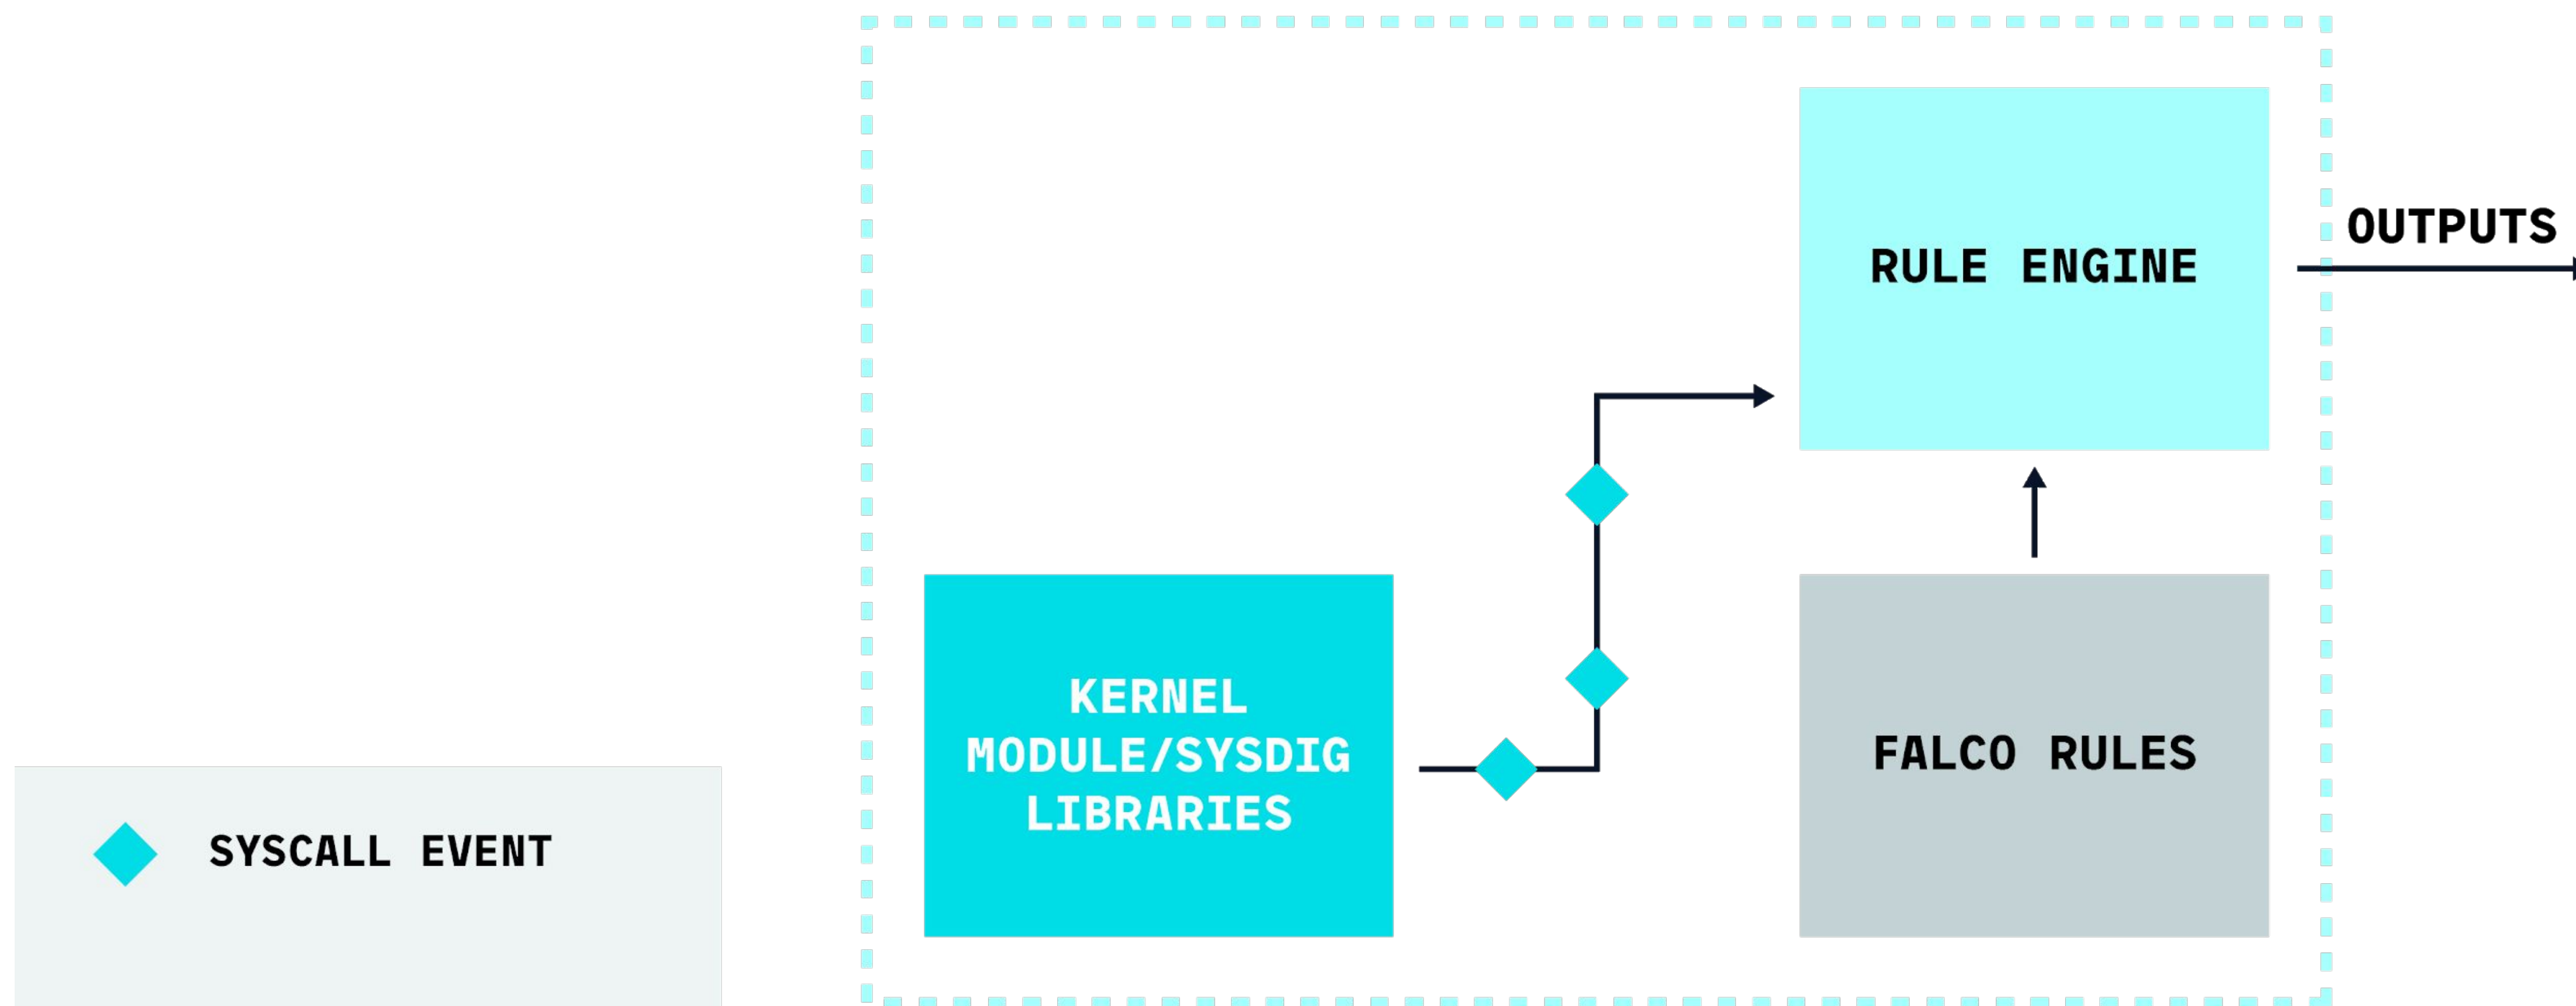


# K8s audit events.

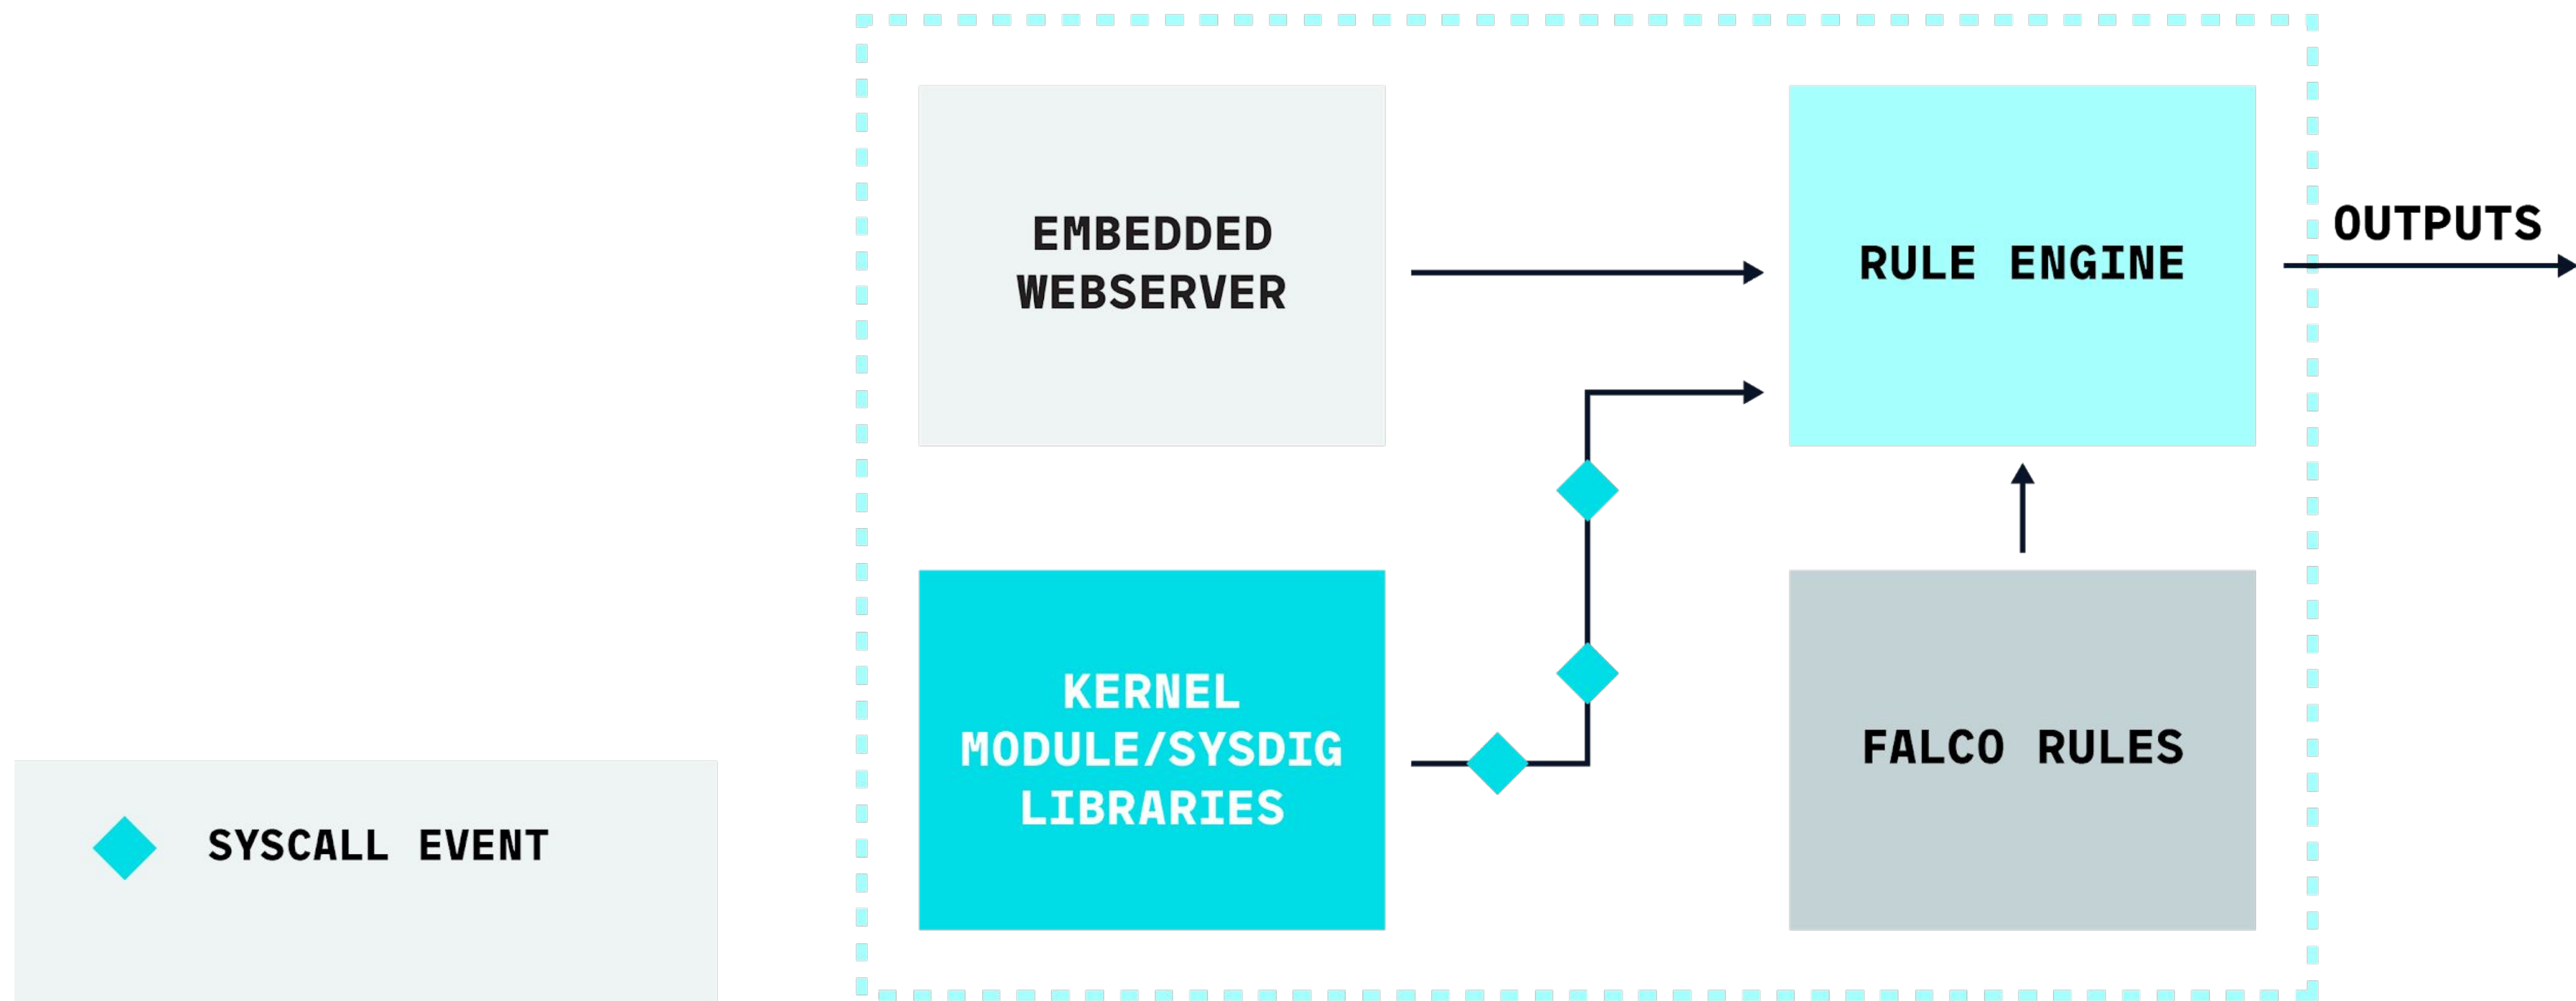
```
{
  "kind": "Event",
  "timestamp": "2018-10-26T13:00:25Z",
  "stage": "ResponseComplete",
  "verb": "delete",
  "requestURI": "/api/v1/namespaces/foo",
  "user": { "username": "minikube-user" },
  "responseStatus": { "code": 200 },
  "objectRef": { "resource": "namespaces", "namespace": "foo" },
  "level": "Request",
  "auditID": "693f4726-2430-450a-83e1-123c050fde98",
  "annotations": { "authorization.k8s.io/decision": "allow" }
}
```



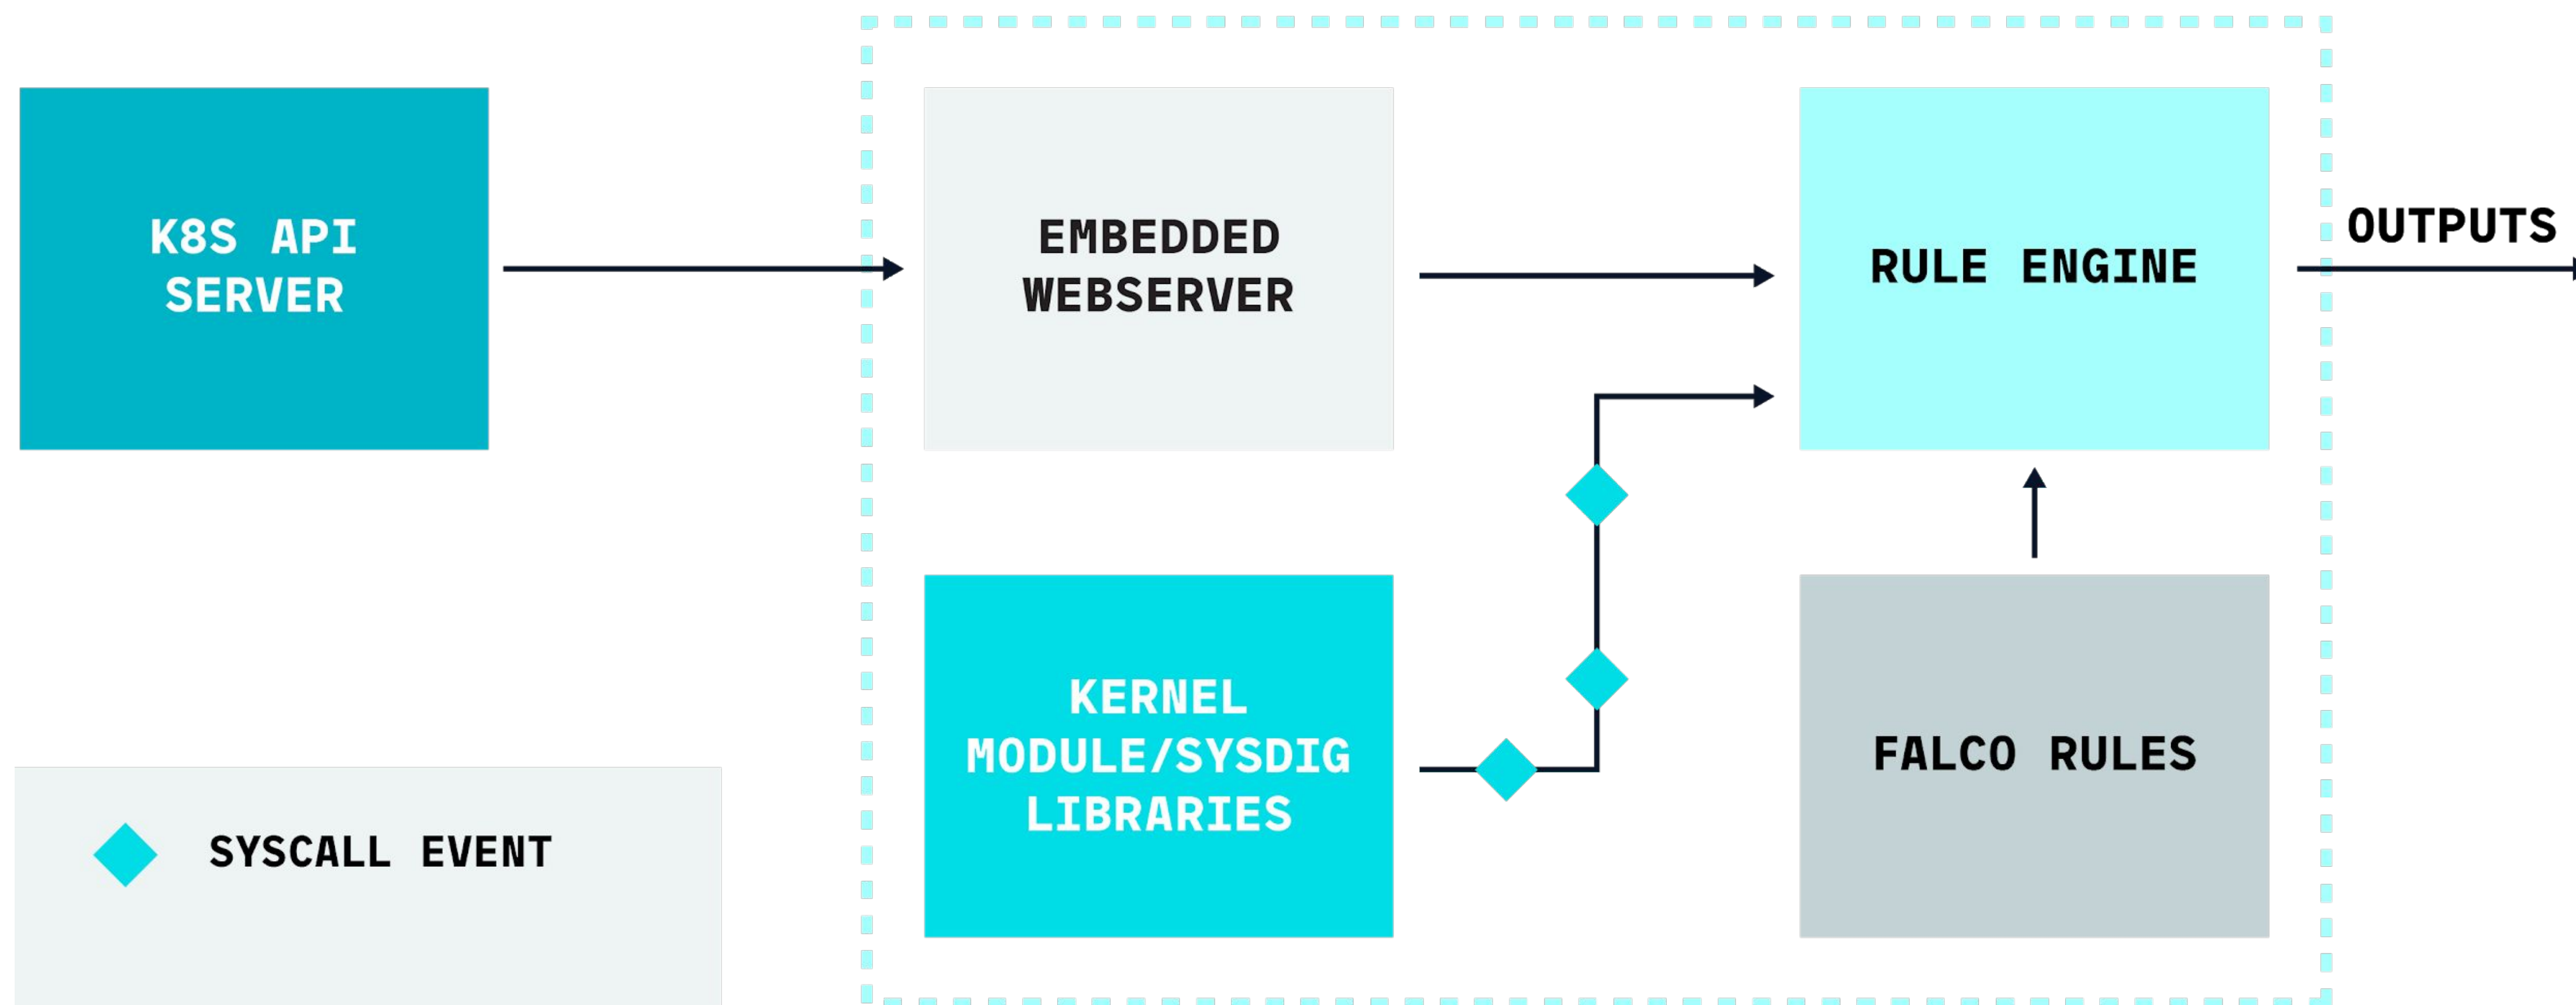
# Falco internal architecture.



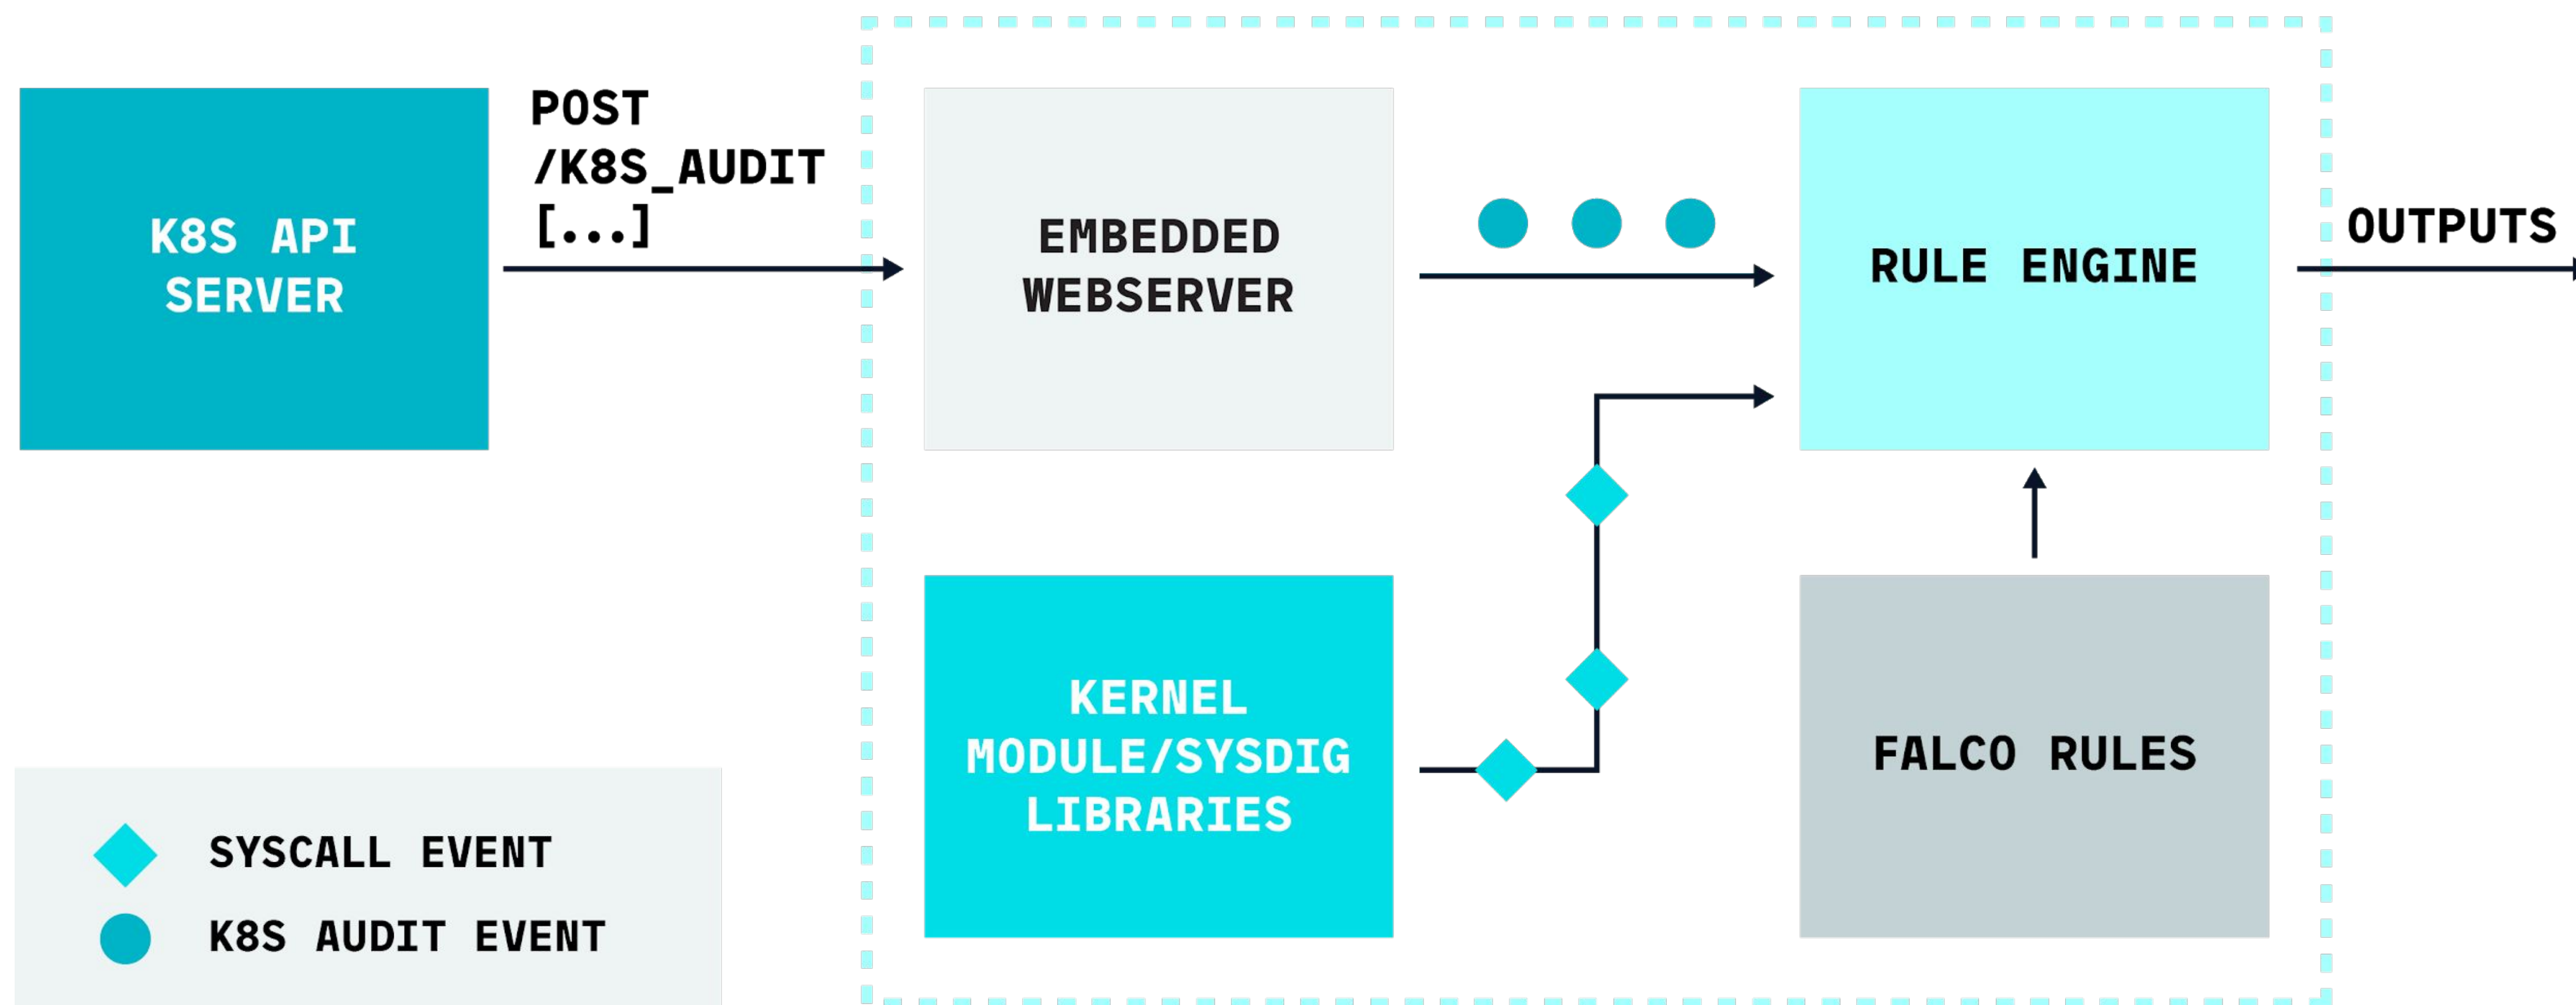
# Falco internal architecture.



# Falco internal architecture.



# Falco internal architecture.





# Supporting K8s audit events.

Create a new “Generic Event” interface

- Event time, ability to extract values using fields

Create a K8s Audit Event object

- Event data is json object, stored in event

Define new fields to extract values from K8s Audit Events

- Uses *Json Pointers* to extract values

Each Falco Rule now has a source

- Default “syscall”, “k8s\_audit” for K8s Audit Events

# JSON pointers.

```
{  
  "foo": 1,  
  "bar": {  
    "baz": "tee"  
  },  
  "arr": [1, 3, 4]  
}
```

- `"/foo"` -> 1
- `"/bar/baz"` -> "tee"
- `"/arr/0"` -> 1
- `"/arr/2"` -> 4

# Accessing K8s audit events.

```
{
  "verb": "delete",
  "responseStatus": { "code": 200 },
  "objectRef": {
    "resource": "namespaces",
    "namespace": "foo"
  },
  "level": "Request"
}
```

- “/verb” -> “delete”
- “/responseStatus/code” -> 200
- “/objectRef/resource” -> “namespaces”

# K8s audit event fields.

`jevt.value[<json_pointer>]`

- Access any field from json object

`jevt.time`

- Access event timestamp

`ka.verb, ka.uri, ka.user.name, ka.target.resource, ...`

- Access specific values from object

- Implemented as macros:

- `ka.verb -> jevt.value[/verb]`

- `ka.target.resource -> jevt.value[/objectRef/resource]`

- Full list: `falco —list=k8s_audit`

# K8s audit rule example.

```
- macro: contains_private_credentials
  condition: >
    (ka.req.configmap.obj contains "aws_access_key_id" or
     ka.req.configmap.obj contains "aws_s3_access_key_id" or
     ka.req.configmap.obj contains "password")

- macro: configmap
  condition: ka.target.resource=configmaps

- macro: modify
  condition: (ka.verb in (create,update,patch))

- rule: Create/Modify Configmap With Private Credentials
  desc: Detect creating/modifying a configmap containing a private
  credential
    (aws key, password, etc.)
  condition: configmap and modify and contains_private_credentials
  output: K8s configmap with private credential (user=%ka.user.name
    verb=%ka.verb name=%ka.req.configmap.name
    configmap=%ka.req.configmap.name config=%ka.req.configmap.obj)
  priority: WARNING
  source: k8s_audit
  tags: [k8s]
```



Demo.



# Join the community.

## Website

- <https://falco.org>

## Public Slack

- <http://slack.sysdig.com/>
- <https://sysdig.slack.com/messages/falco>

## Blog

- <https://sysdig.com/blog/tag/falco/>

## Workshop

- <https://github.com/falcosecurity/falco-security-workshop>

## Github

- <https://github.com/falcosecurity/falco>

## Documentation

- <https://github.com/falcosecurity/falco/wiki>

## Docker Hub

- <https://hub.docker.com/r/falcosecurity/falco/>



**Thank You!**

