To CRD, or Not to CRD, *That* Is the Question

KubeCon | CloudNativeCon

OPEN SOURCE SUMMIT

China 2019

Intros

Sam Gunaratne

Pivotal, London

@sam_gun

Ed King

Pivotal, London

@edking2

This Talk

- Motivations
- Where we've come from:
  - The k8s API, CRDs & Operators
- Where we **may** be heading:
  - Building on top of the k8s API
  - Pros, cons and considerations
- Closing remarks and Questions

Motivations

K8s is API-centric

The k8s API

Declarative

- What vs How
- World can be built from a set of config
- Easy to record changes

State separation

- Desired state
- Observed state

Level-based

- Does not rely on individual changes
- Less prone to error

Transparent

- One control plane
- Facilities composability

Example

k8s api

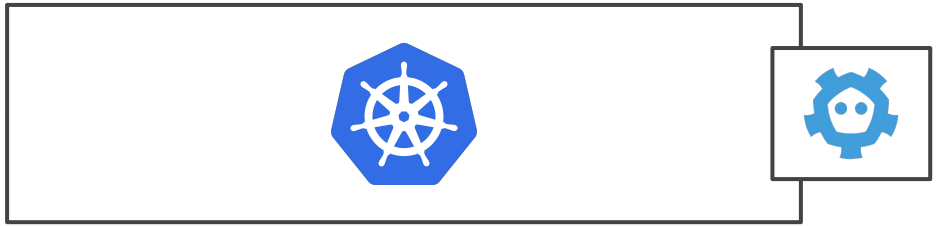

“I want 1 pod to be running”

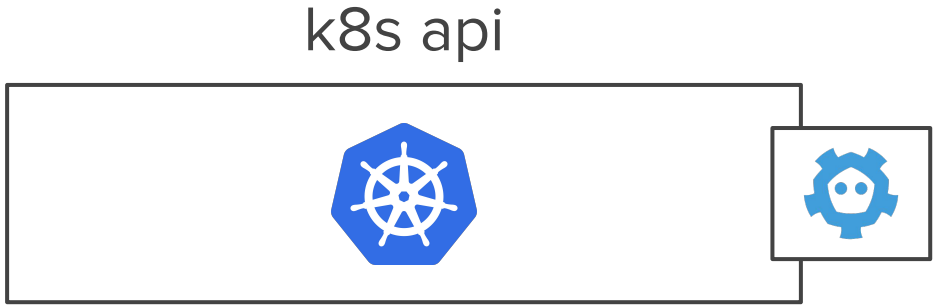

desired

“I want 1 pod to be running”

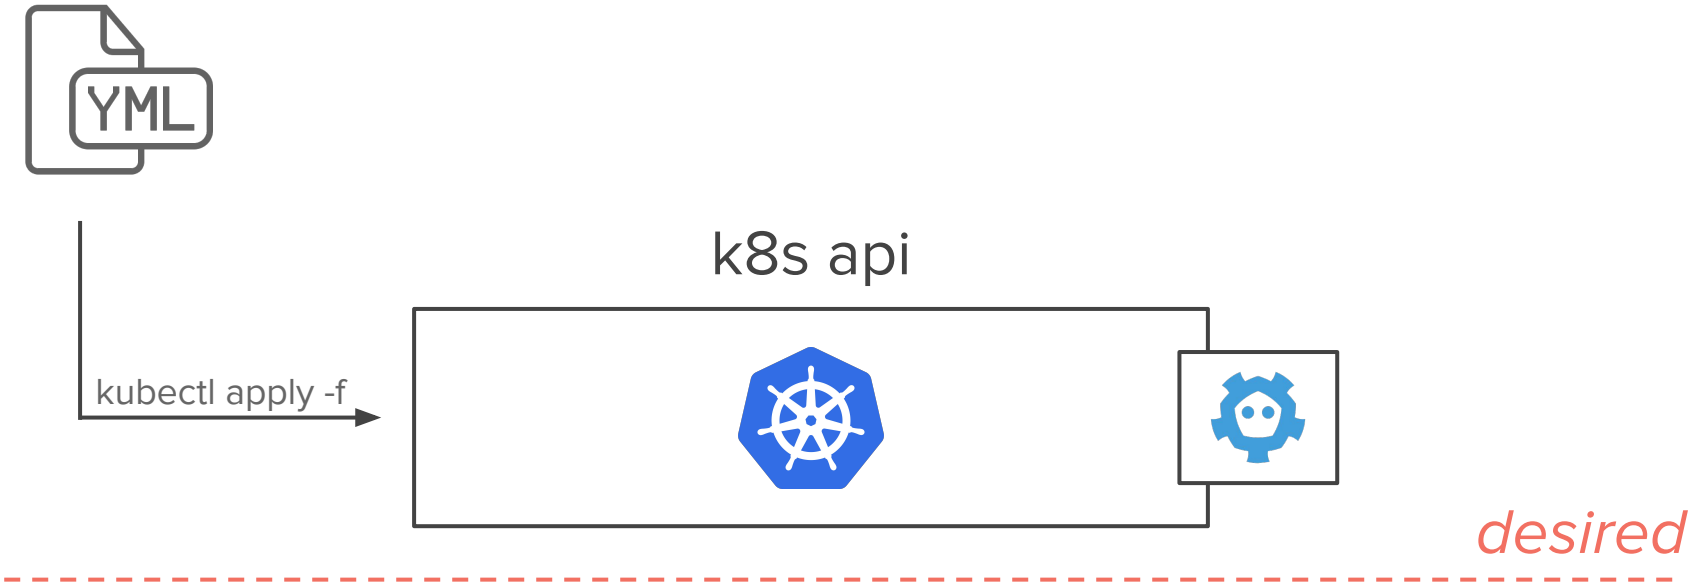


```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-1
  namespace: default
spec:
  containers:
    name: my-ctr-1
    image: reg.io/my-img:v1.0.0
```

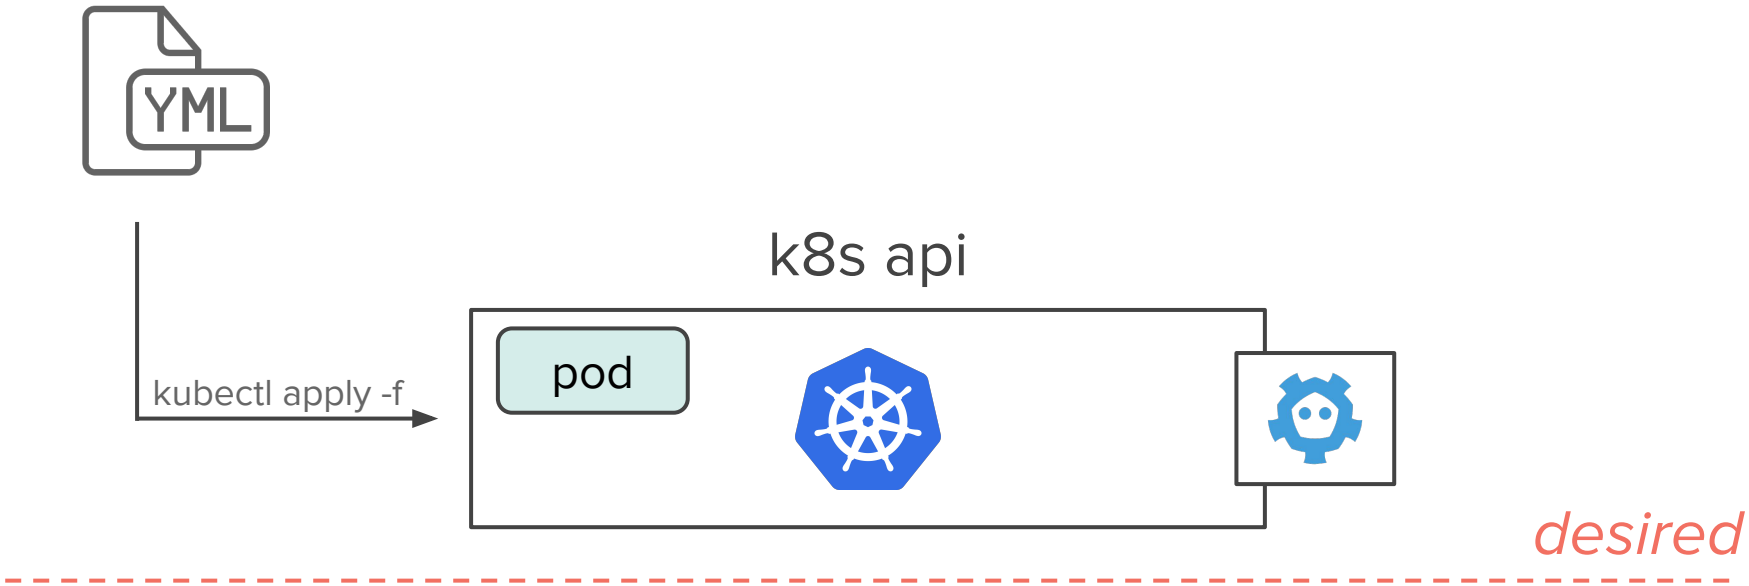
YML
kubectl apply -f
k8s api
pod
desired

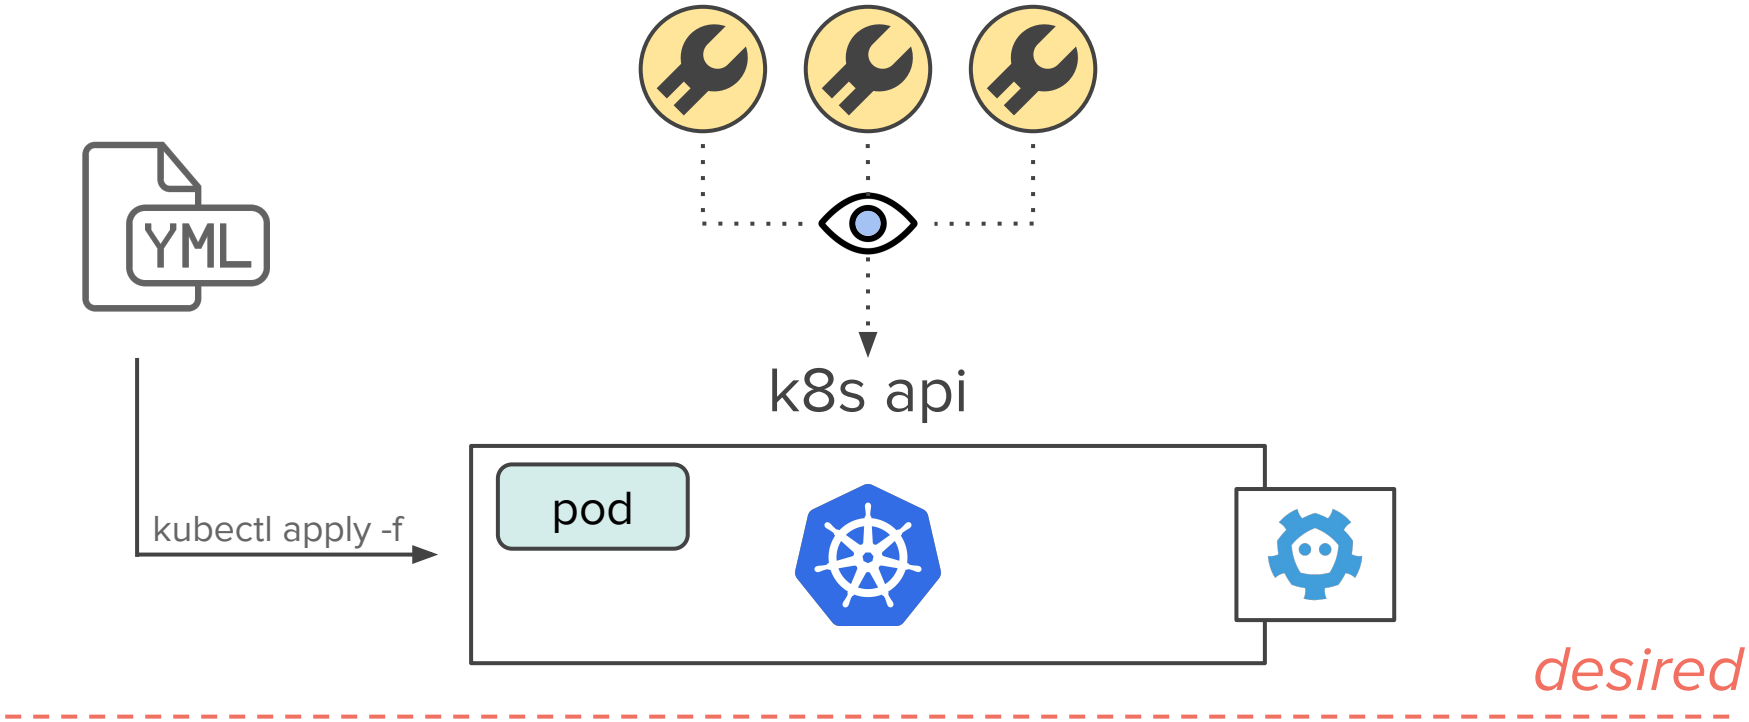
YML
kubectl apply -f
k8s api
pod
desired

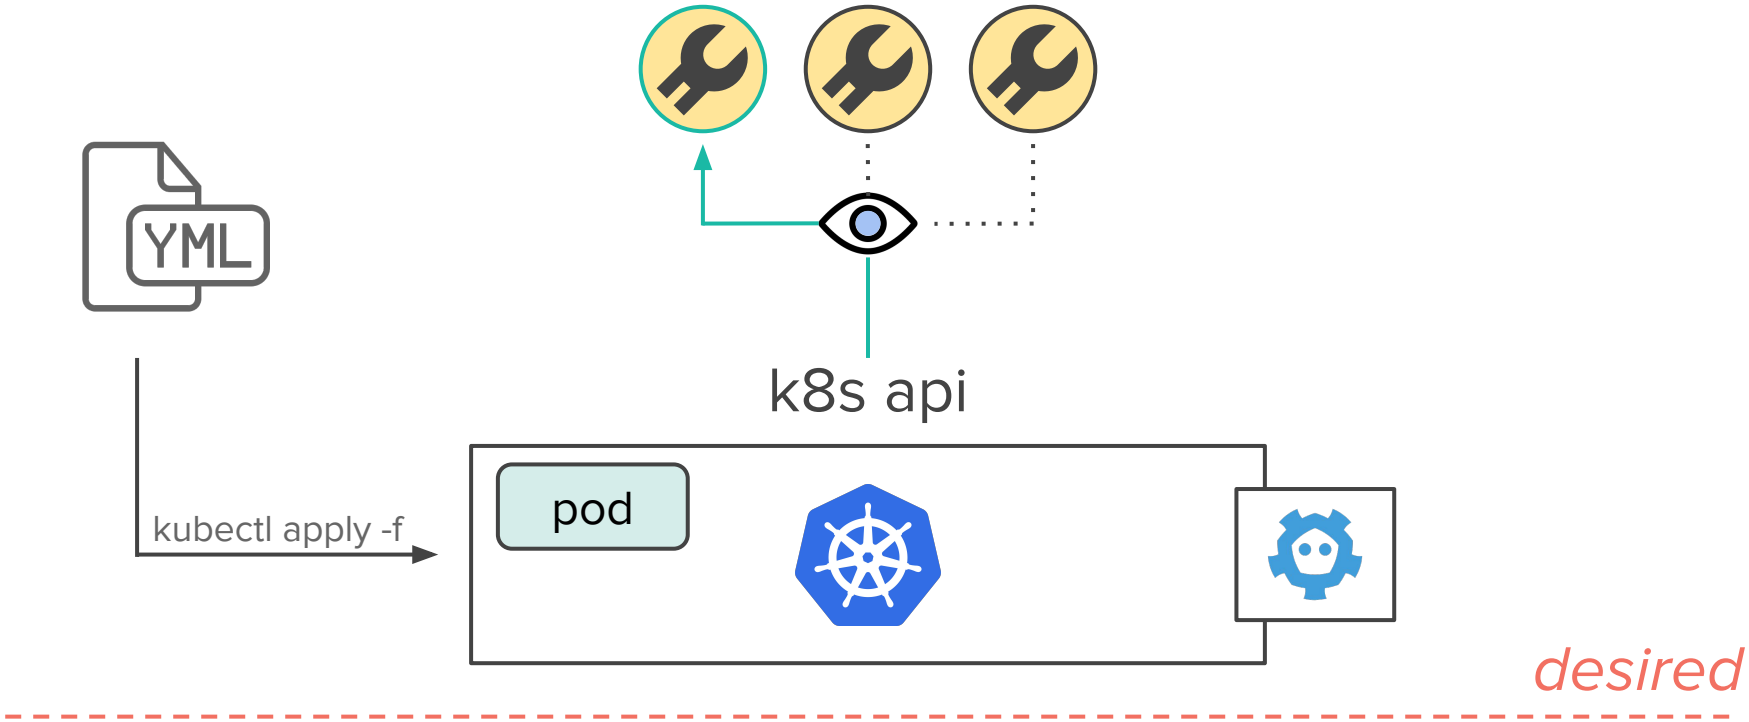
YML
k8s api
kubectl apply -f
pod
desired

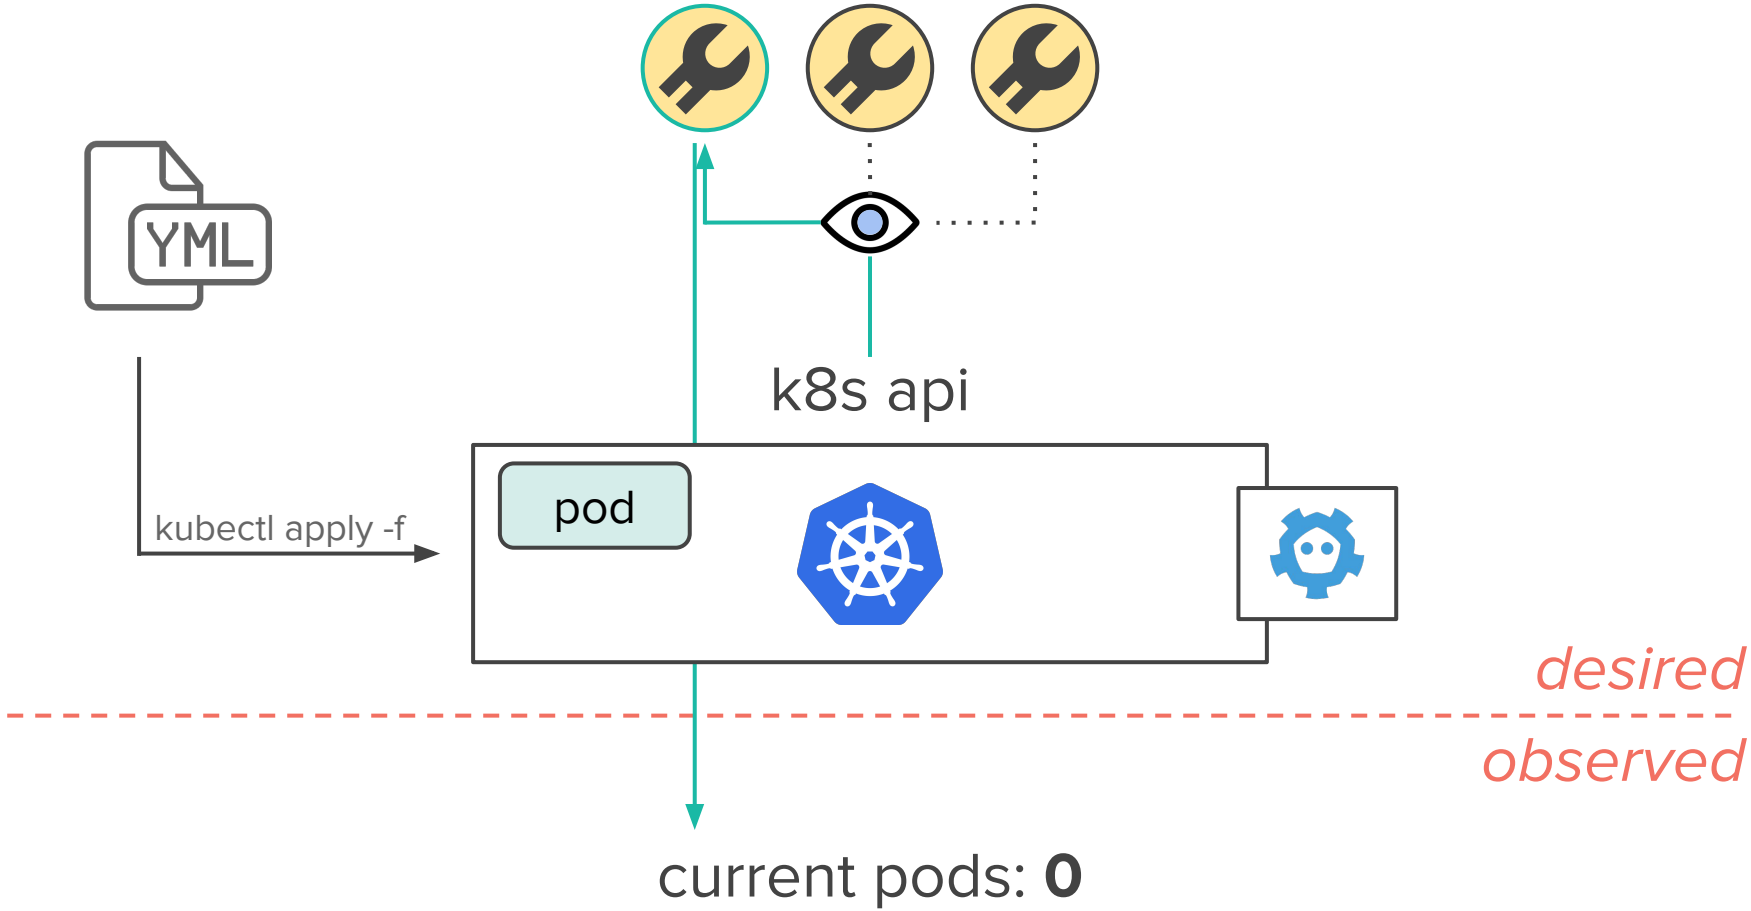
YML
kubectl apply -f
k8s api
pod
desired
observed
current pods: 0

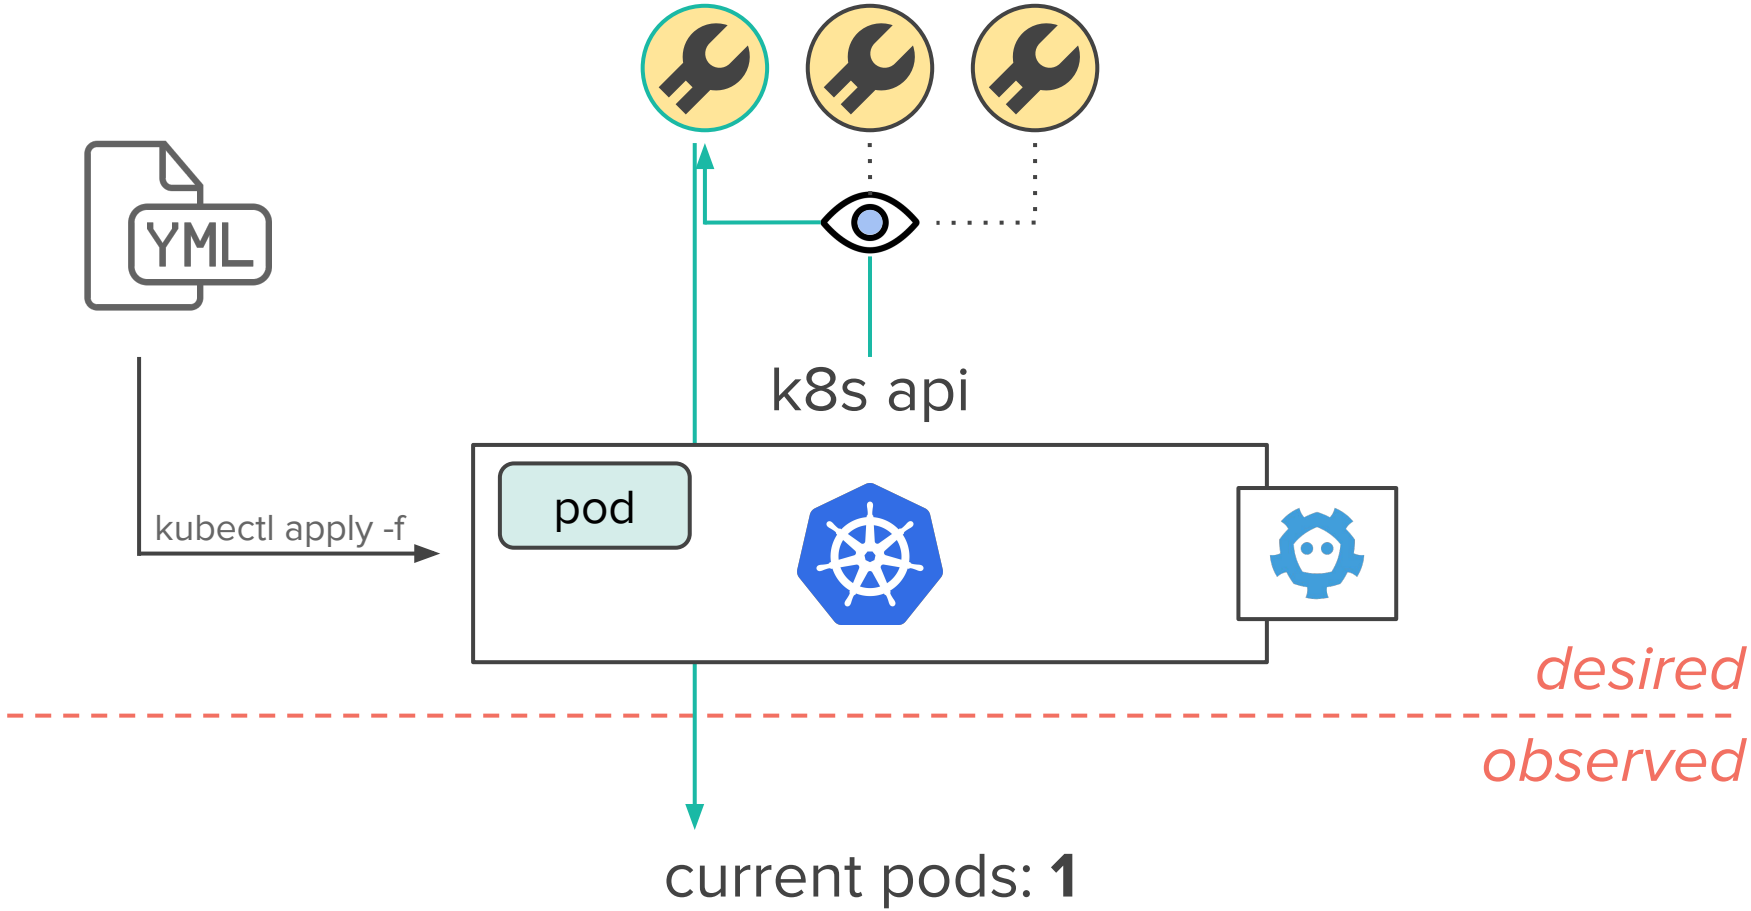
YML
kubectl apply -f
k8s api
pod
desired
observed
current pods: 1

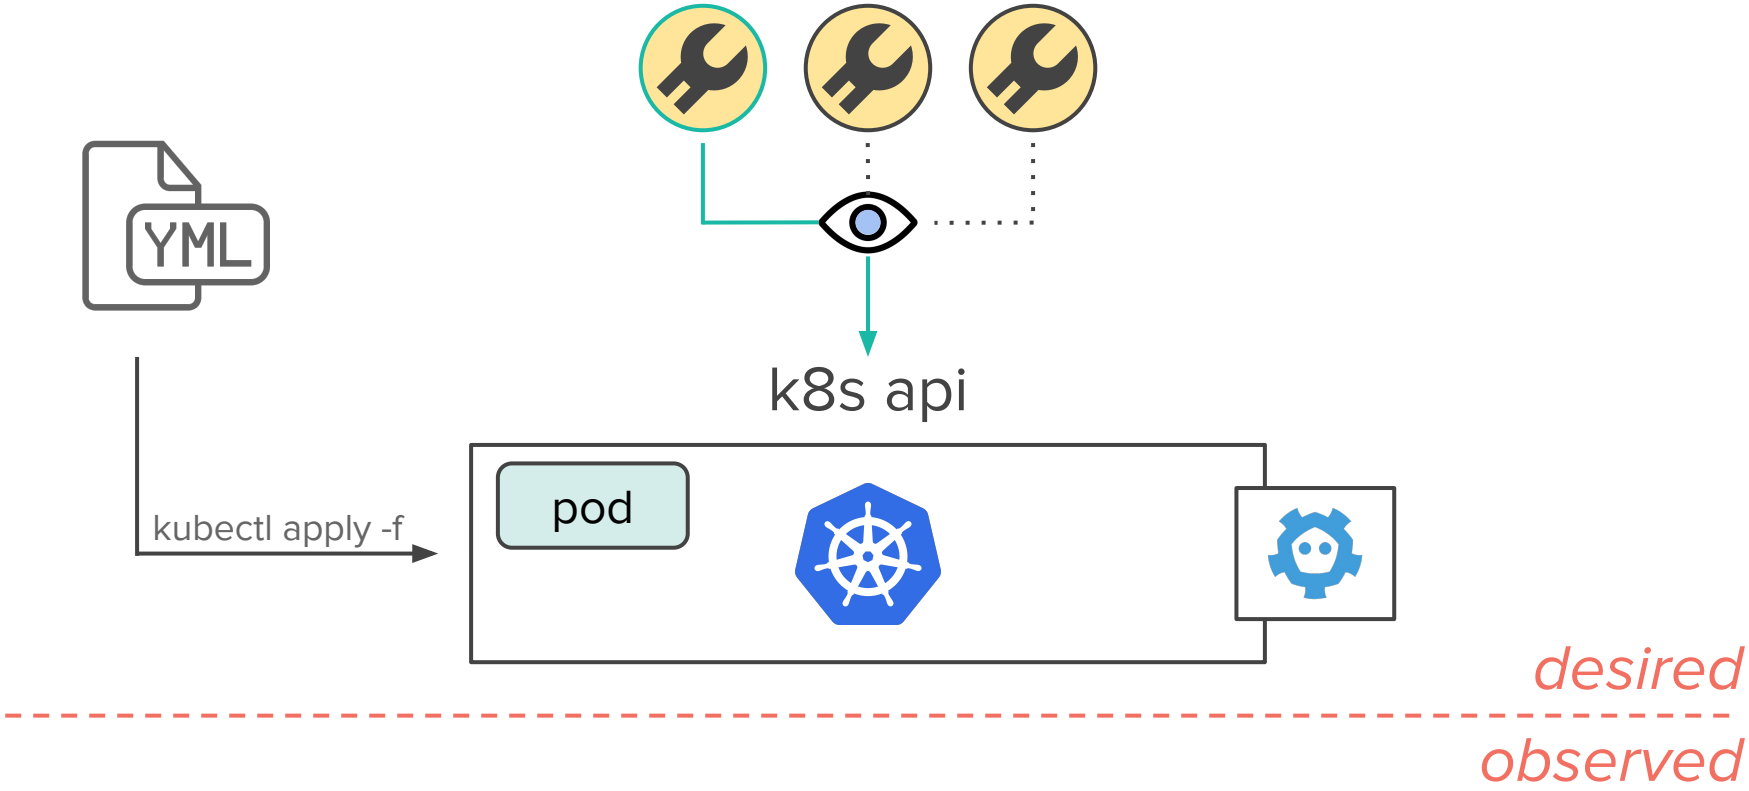

current pods: **1**


```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-1
  namespace: default
spec:
  containers:
    name: my-ctr-1
    image: reg.io/my-img:v1.0.0
status:
  state: running
  ready: true
```

Custom Resource Definitions


```
apiVersion: v1
kind: Cake
metadata:
  name: my-delicious-cake-1
  namespace: default
spec:
  name: victoria-sponge
  ingredients:
  - sugar
  - eggs
  - flour
  bakeTime: 30m
```

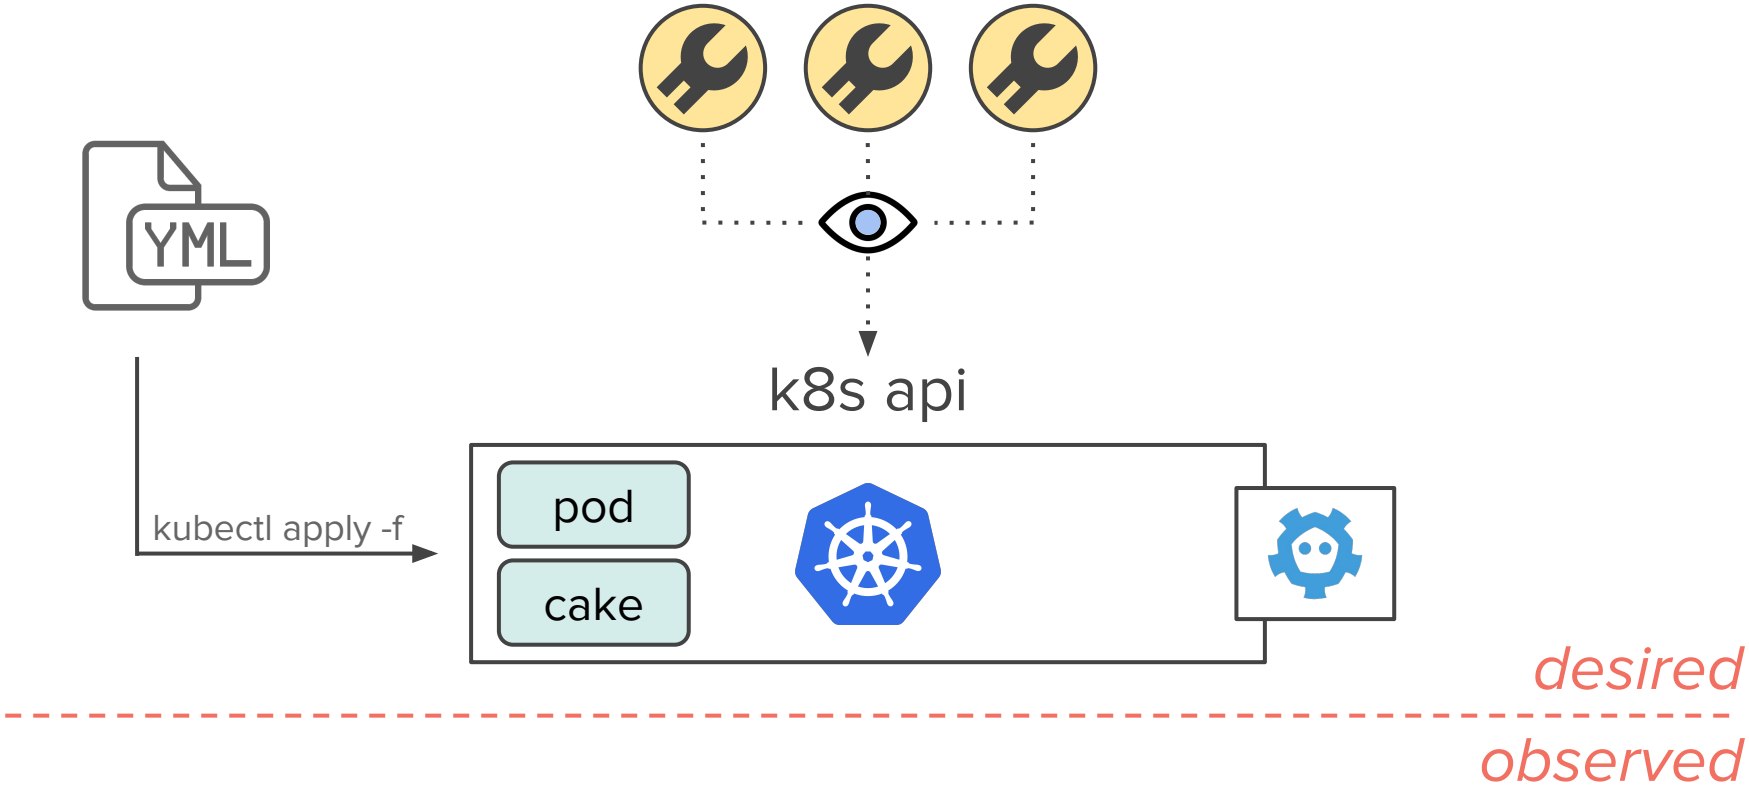
YML
kubectl apply -f
k8s api
pod
cake
desired
observed

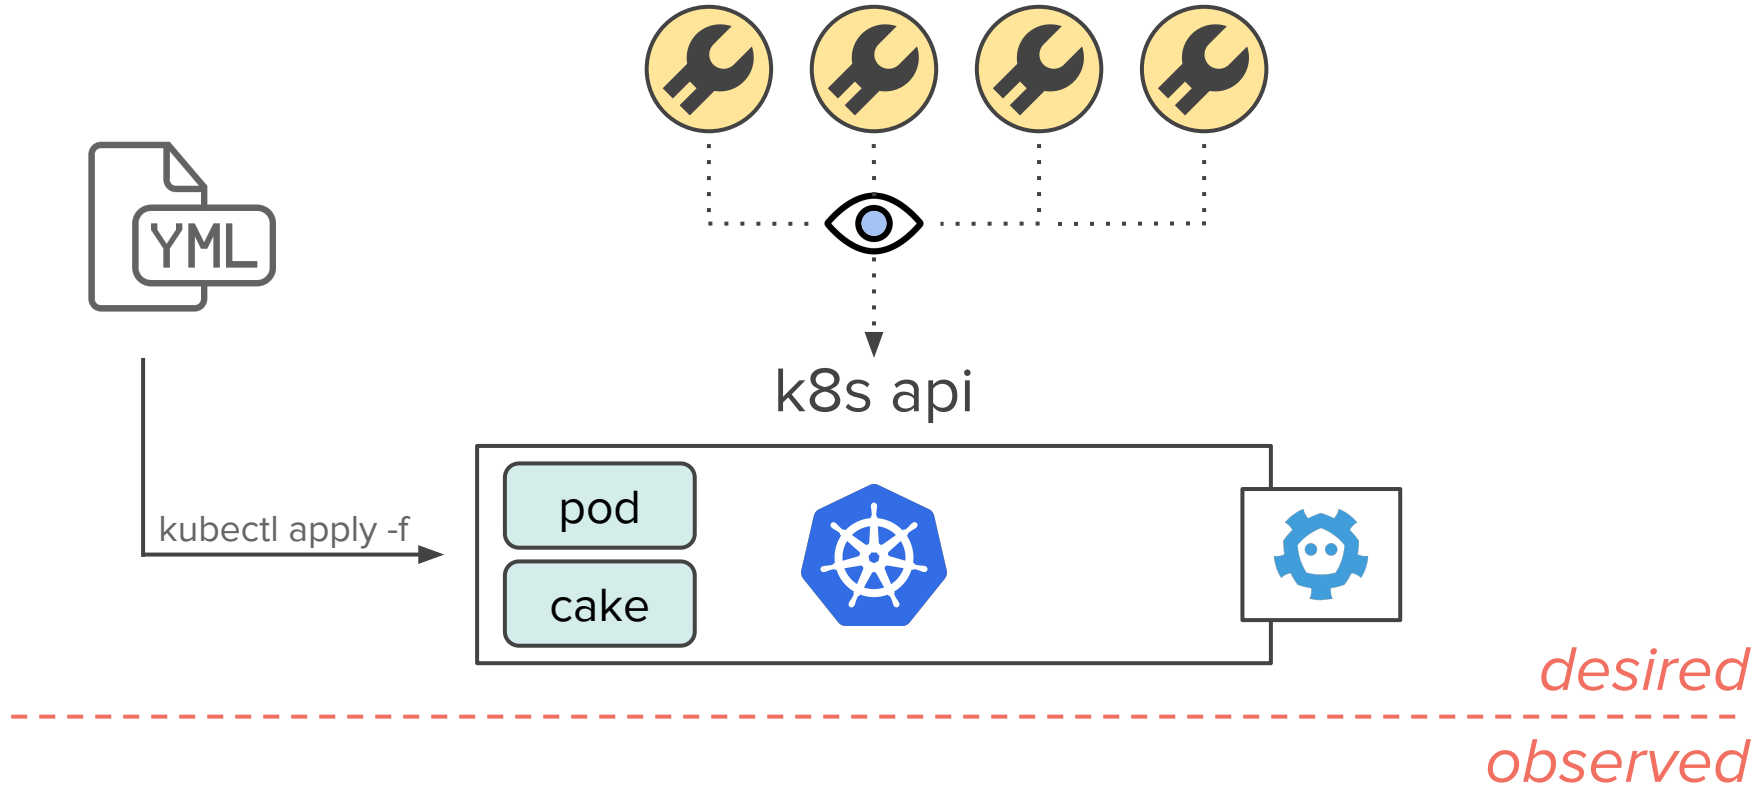
YML
kubectl apply -f
k8s api
pod
cake
desired
observed

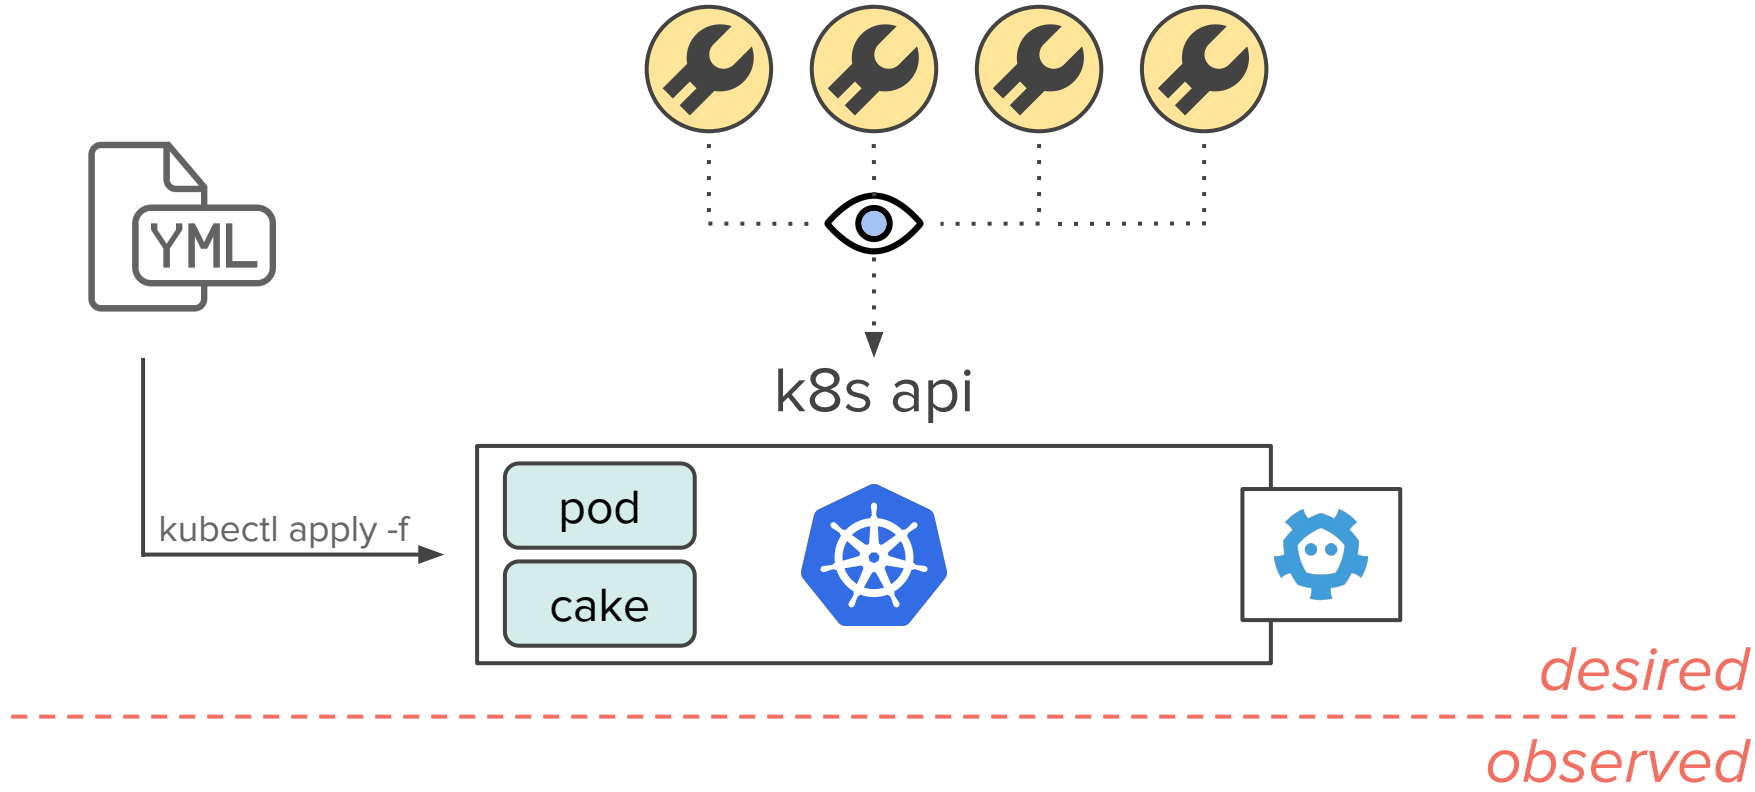
YML
kubectl apply -f
k8s api
pod
cake
desired
observed

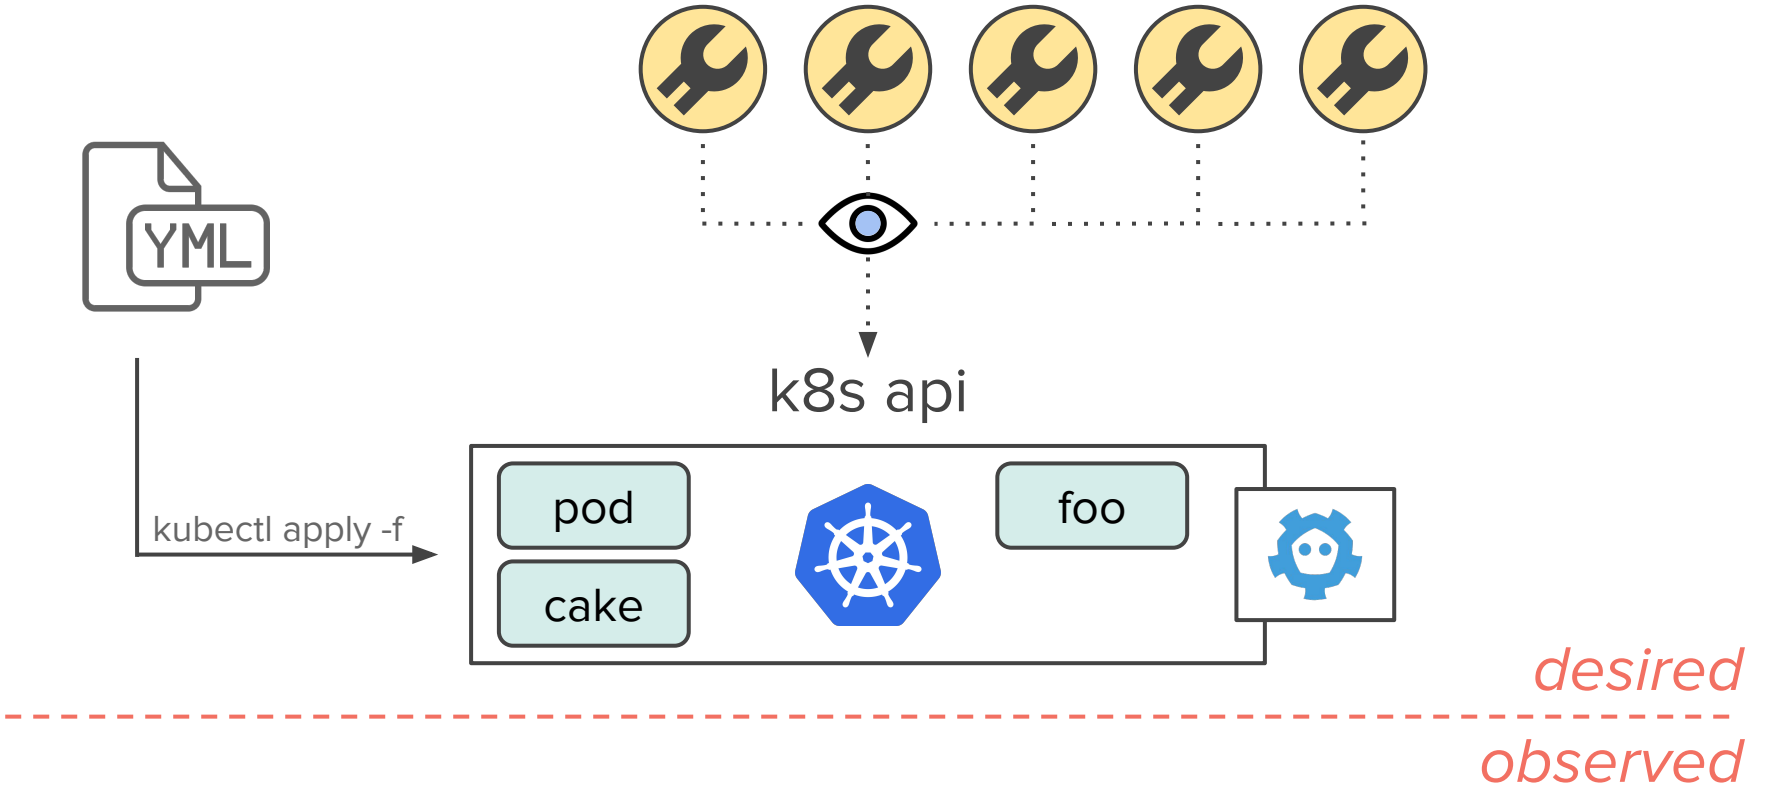
YML
kubectl apply -f
k8s api
pod
cake
foo
desired
observed

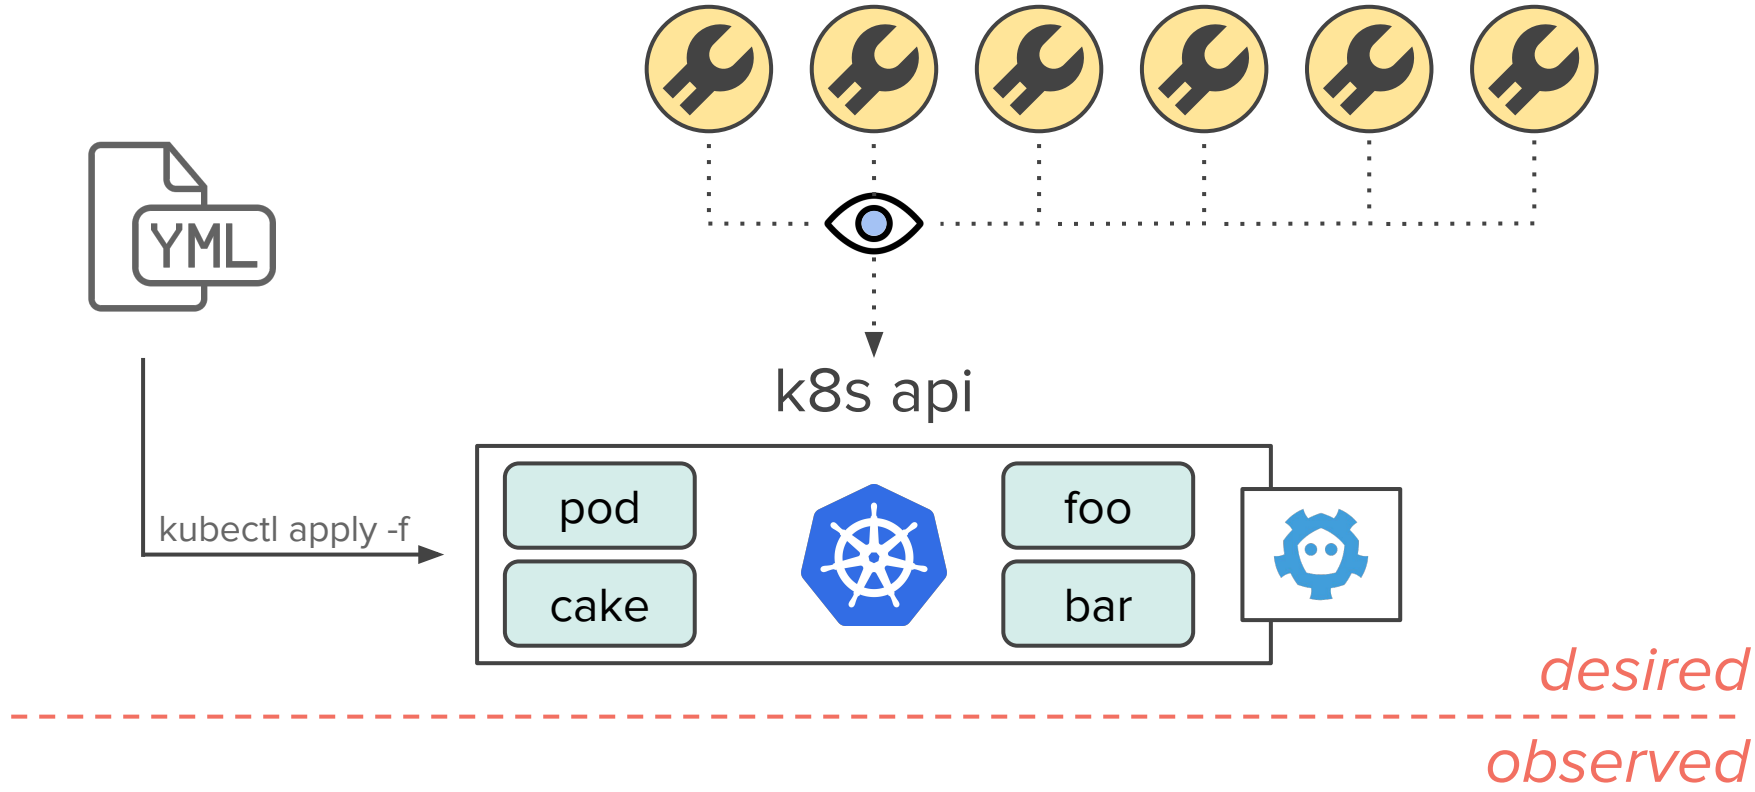
YML
kubectl apply -f
k8s api
pod
cake
foo
bar
desired
observed

Timeline

March 2016

Third party resources

March 2016

Third party resources

Nov 2016

Operator framework

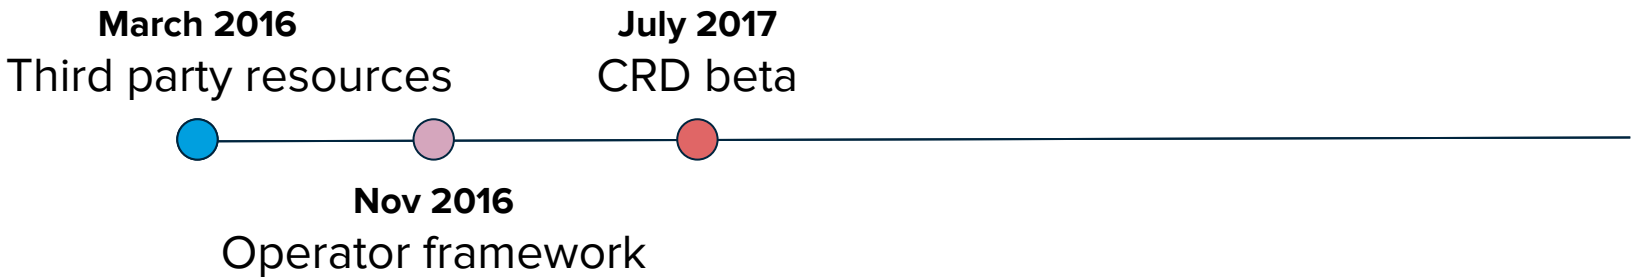
March 2016
Third party resources
Nov 2016
Operator framework
July 2017
CRD beta

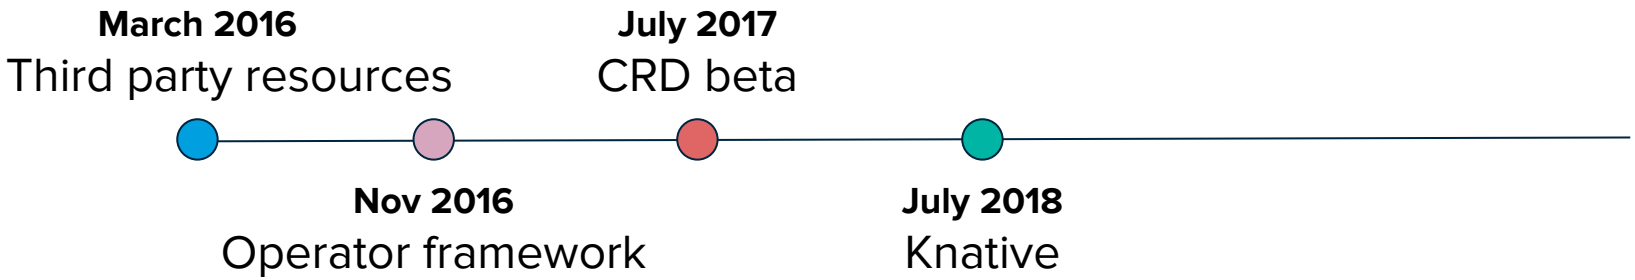
March 2016
Third party resources
Nov 2016
Operator framework
July 2017
CRD beta
July 2018
Knative

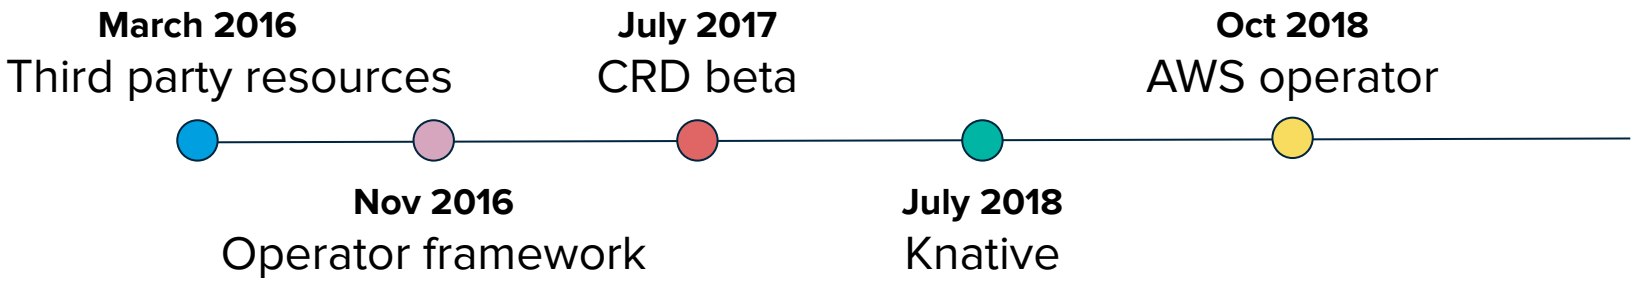
March 2016
Third party resources
Nov 2016
Operator framework
July 2017
CRD beta
July 2018
Knative
Oct 2018
AWS operator

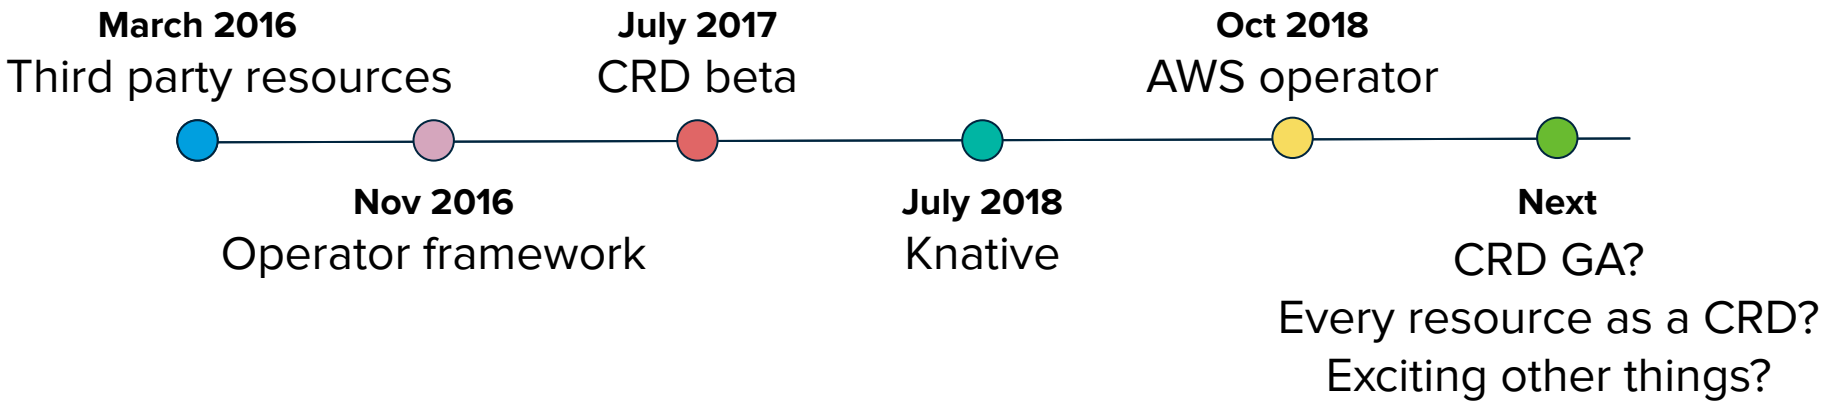
March 2016
Third party resources
Nov 2016
Operator framework
July 2017
CRD beta
July 2018
Knative
Oct 2018
AWS operator
Next
CRD GA?
Every resource as a CRD?
Exciting other things?


Let me paint you a picture

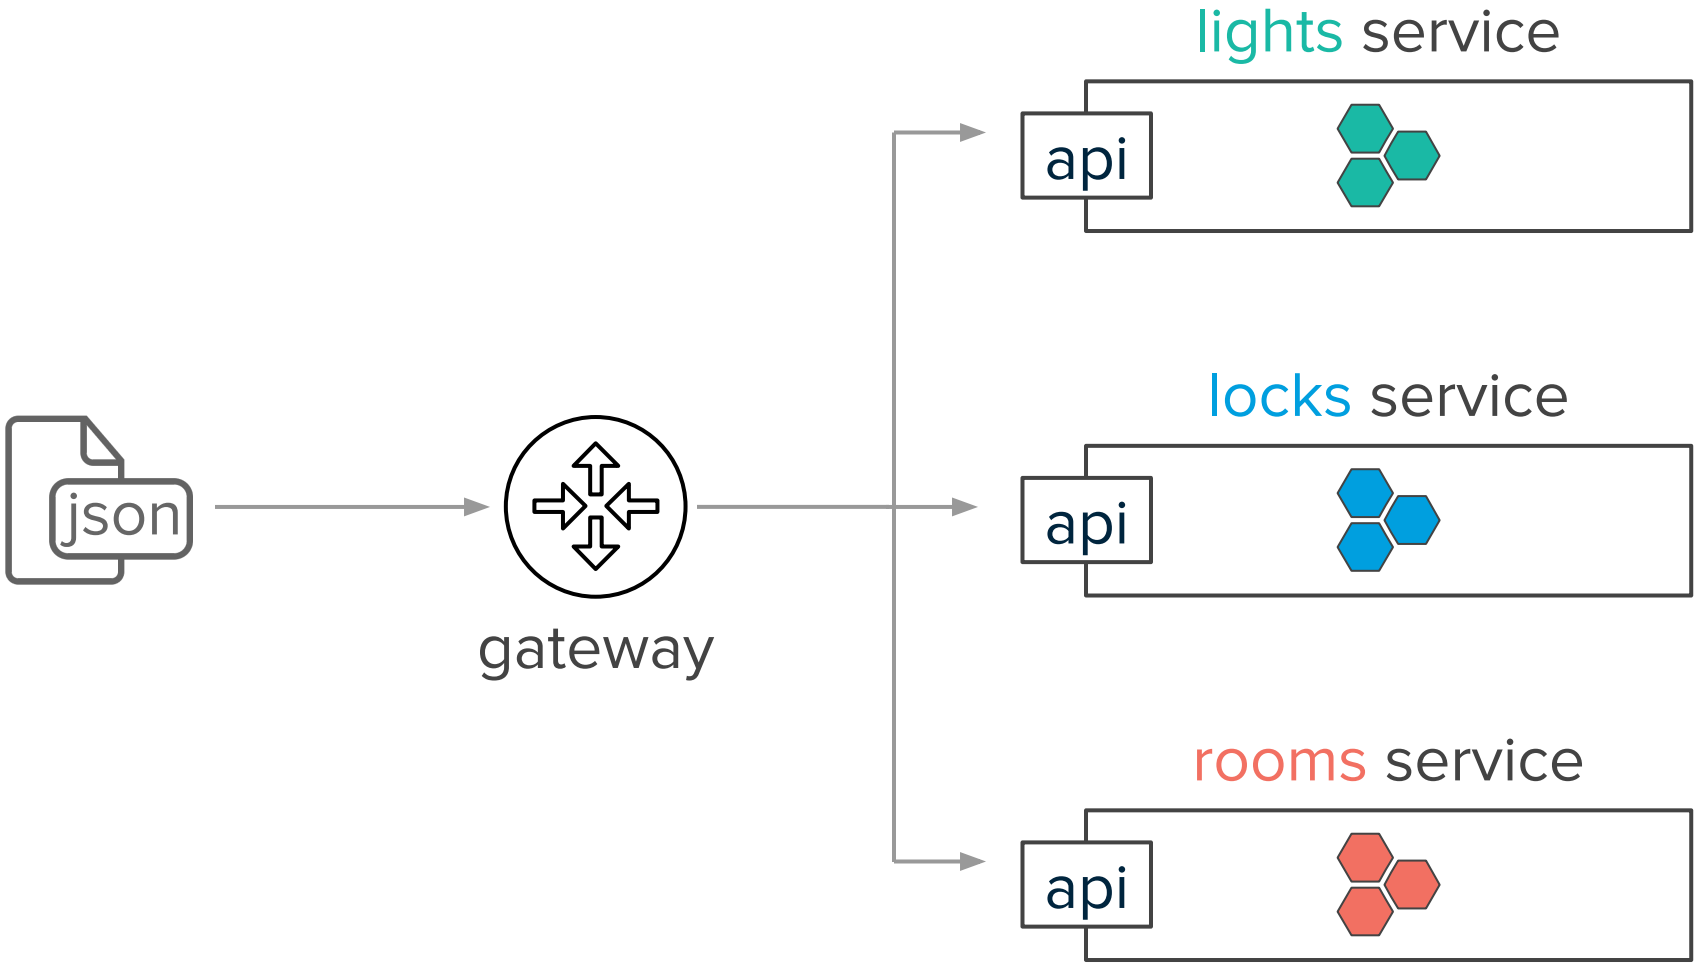
lights service
api
locks service
api
rooms service
api
json
gateway

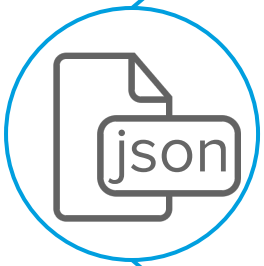


```
{
  “action”: “switch_on”,
  “lights”: [
    “lamp-1”,
    “lamp-2”
  ],
  “room”: “kitchen”
}
```

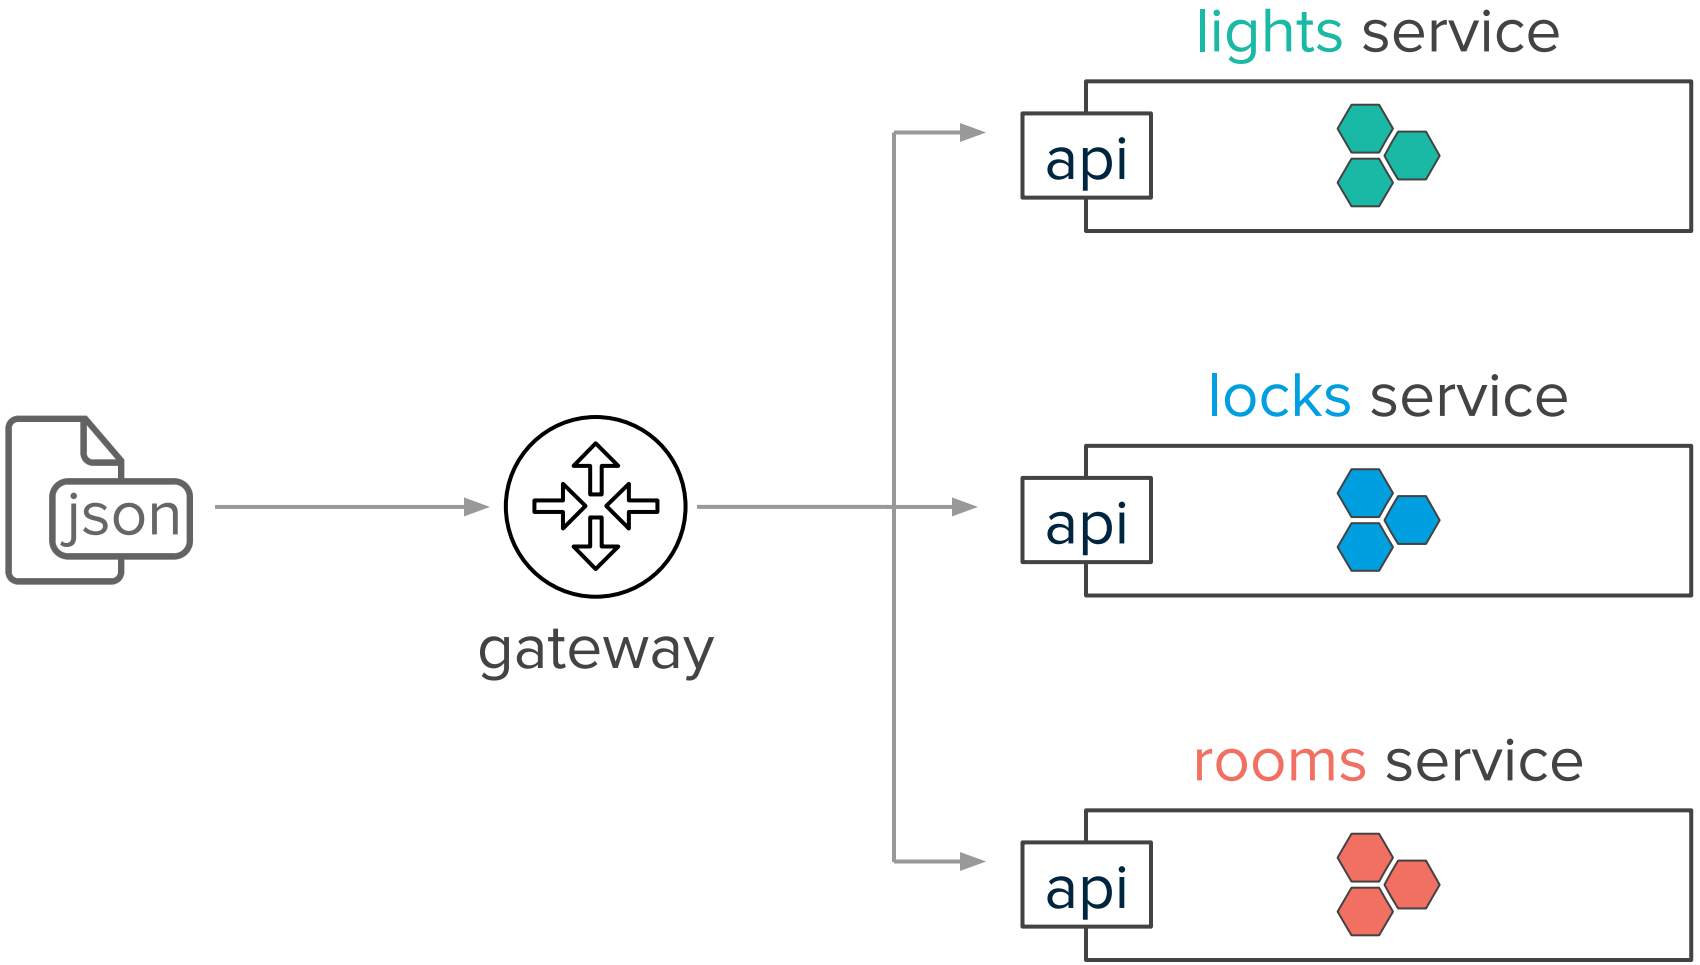
lights service
api
locks service
api
rooms service
api
json
gateway

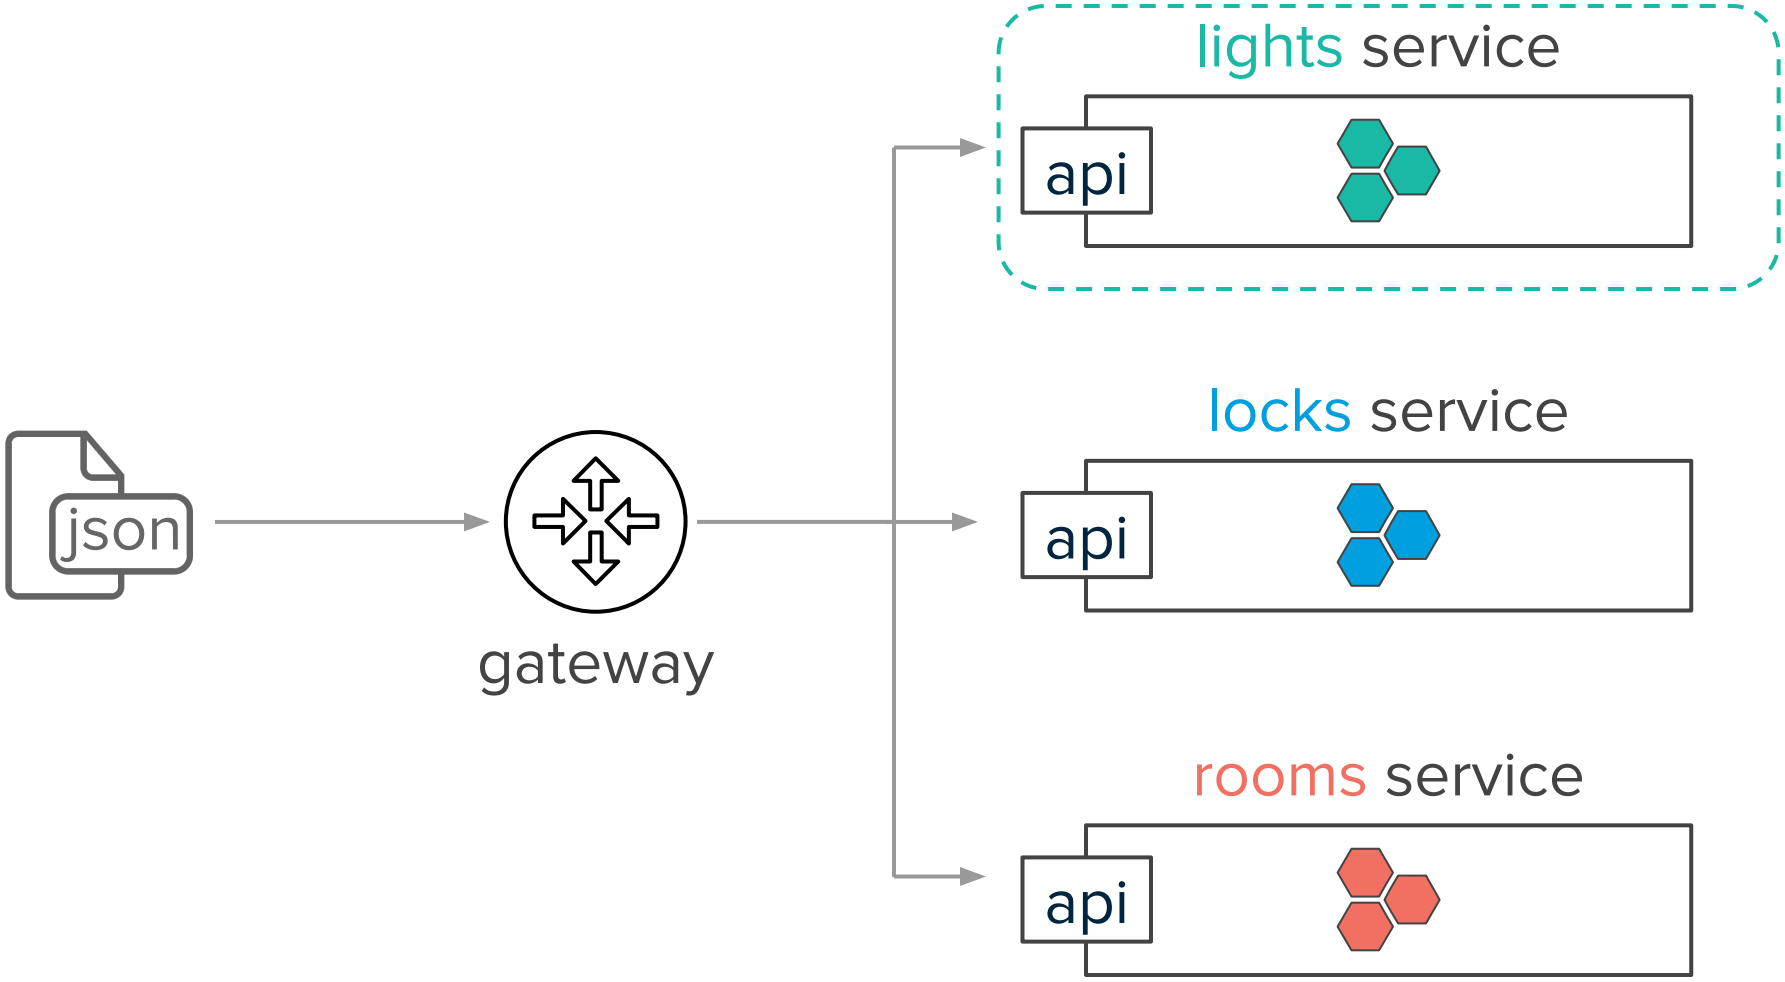
lights service
api
locks service
api
rooms service
api
json
gateway

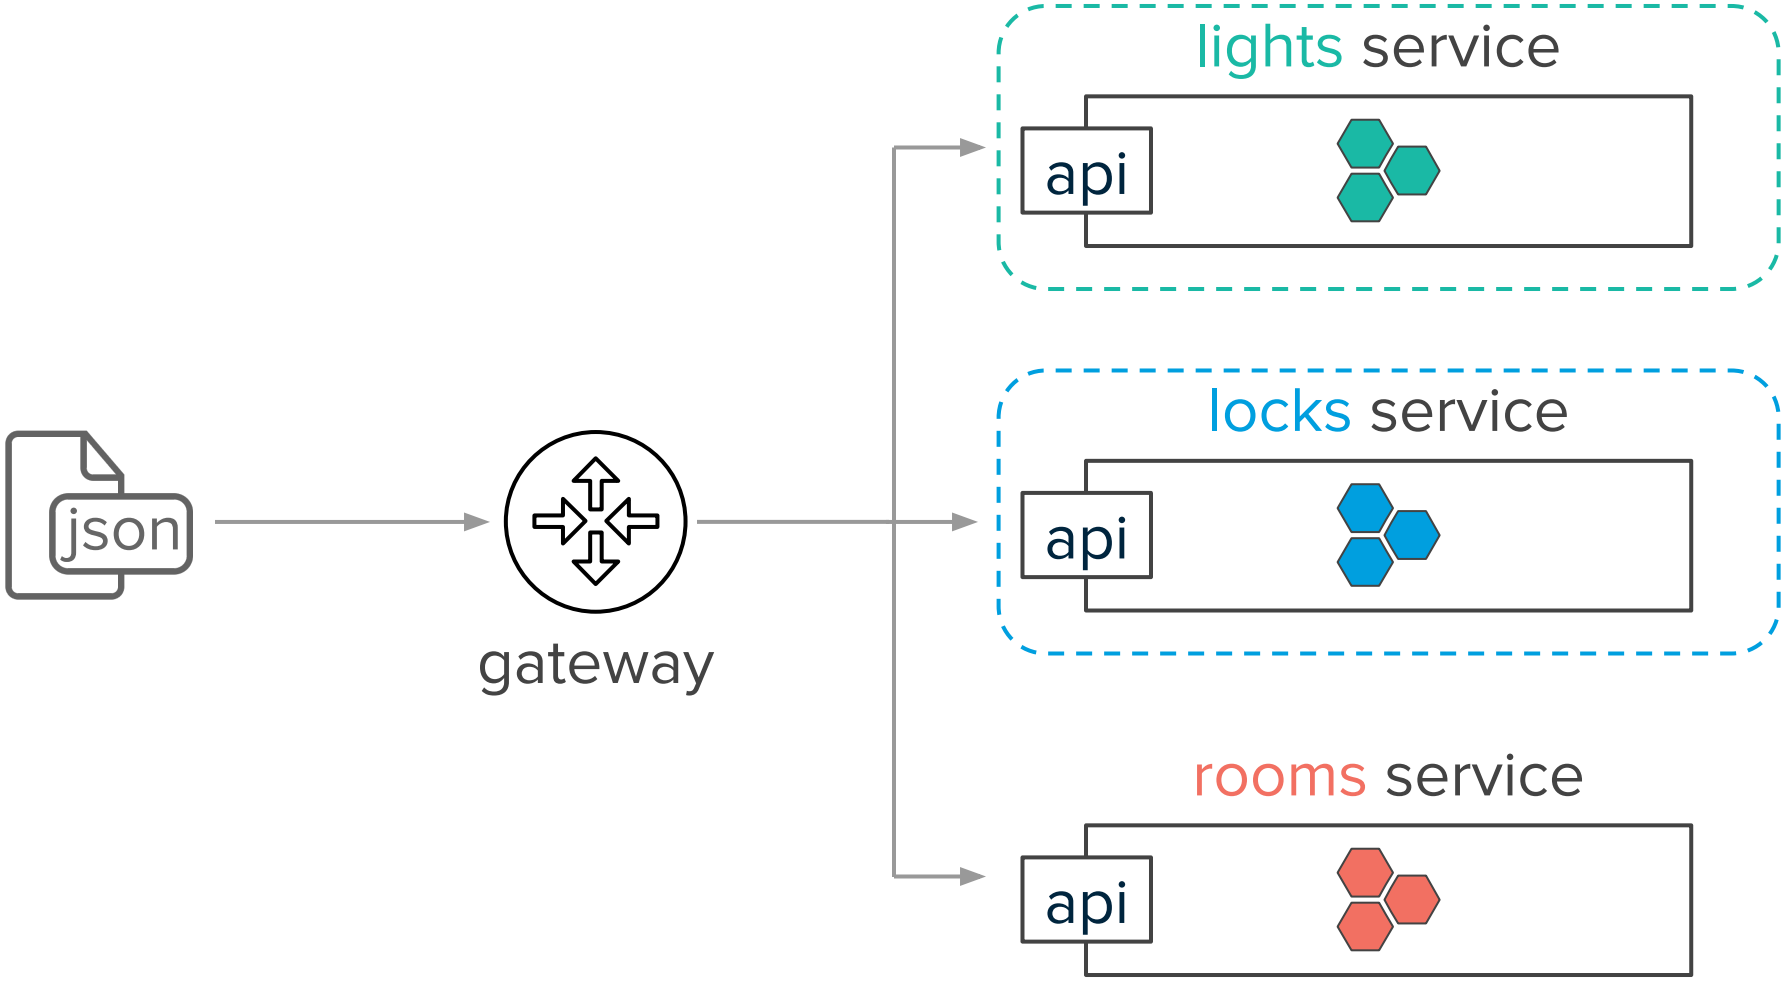
lights service
api
locks service
api
rooms service
api
json
gateway

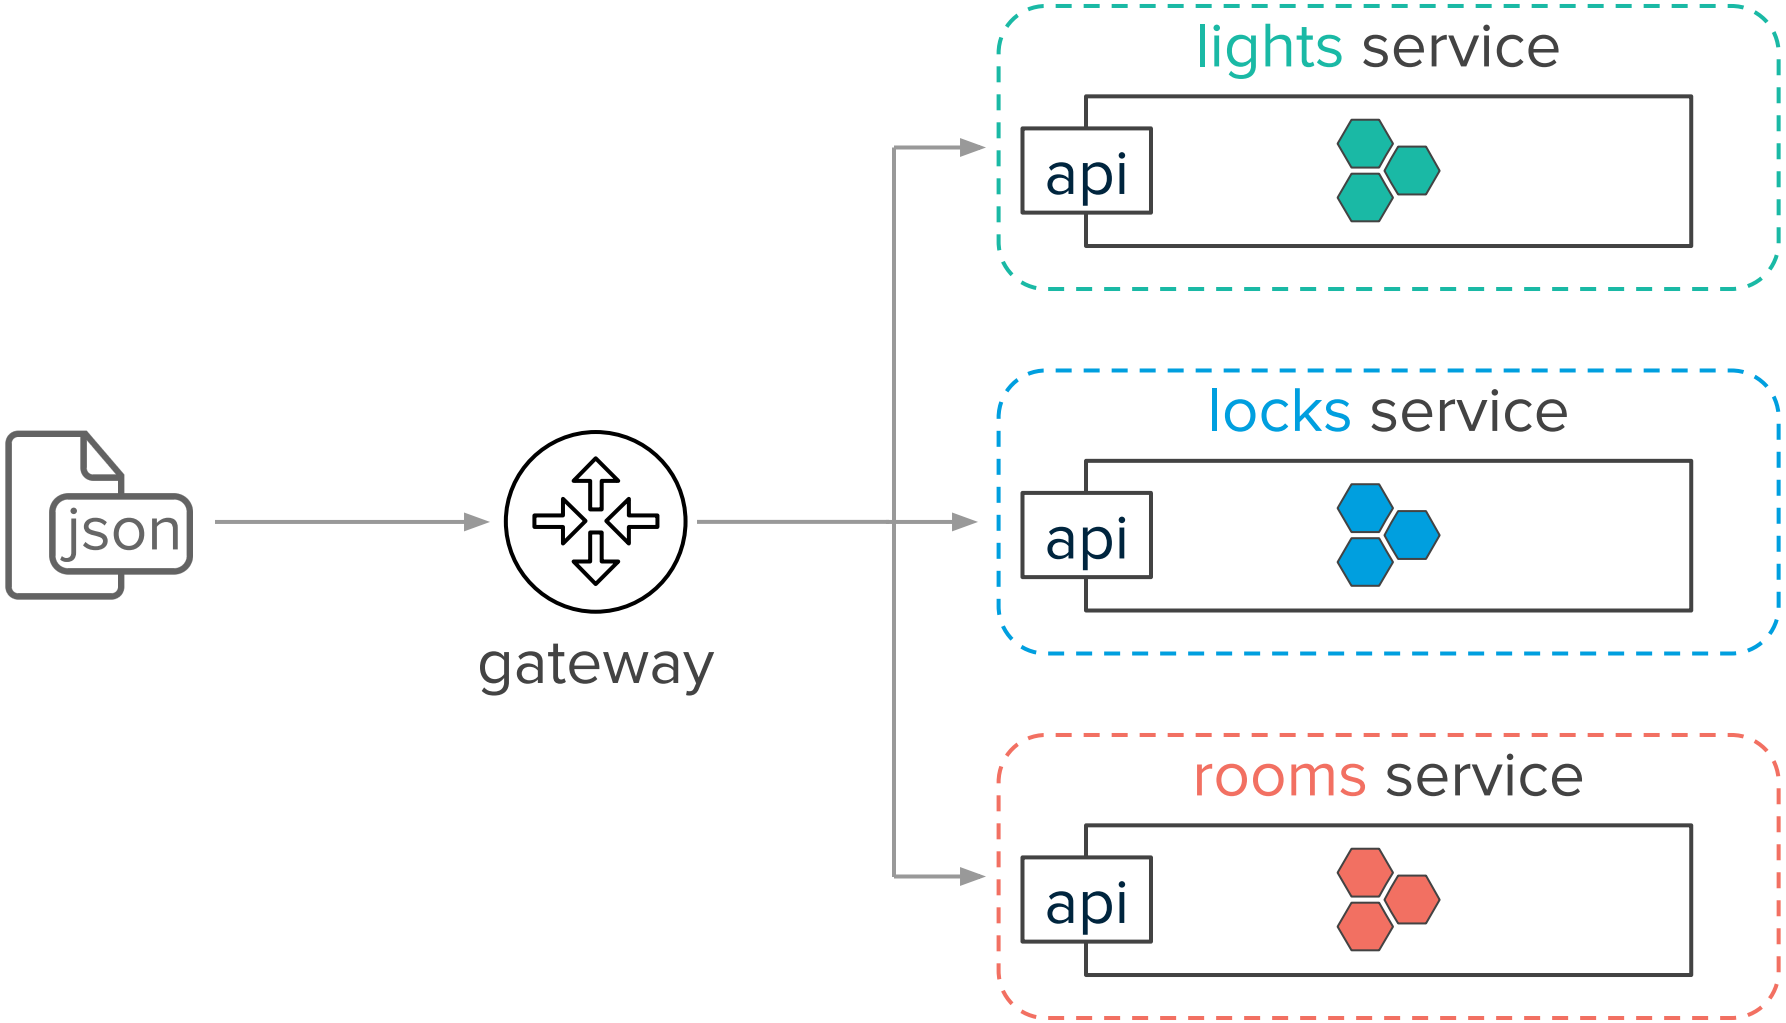
lights service
api
locks service
api
rooms service
api
json
gateway

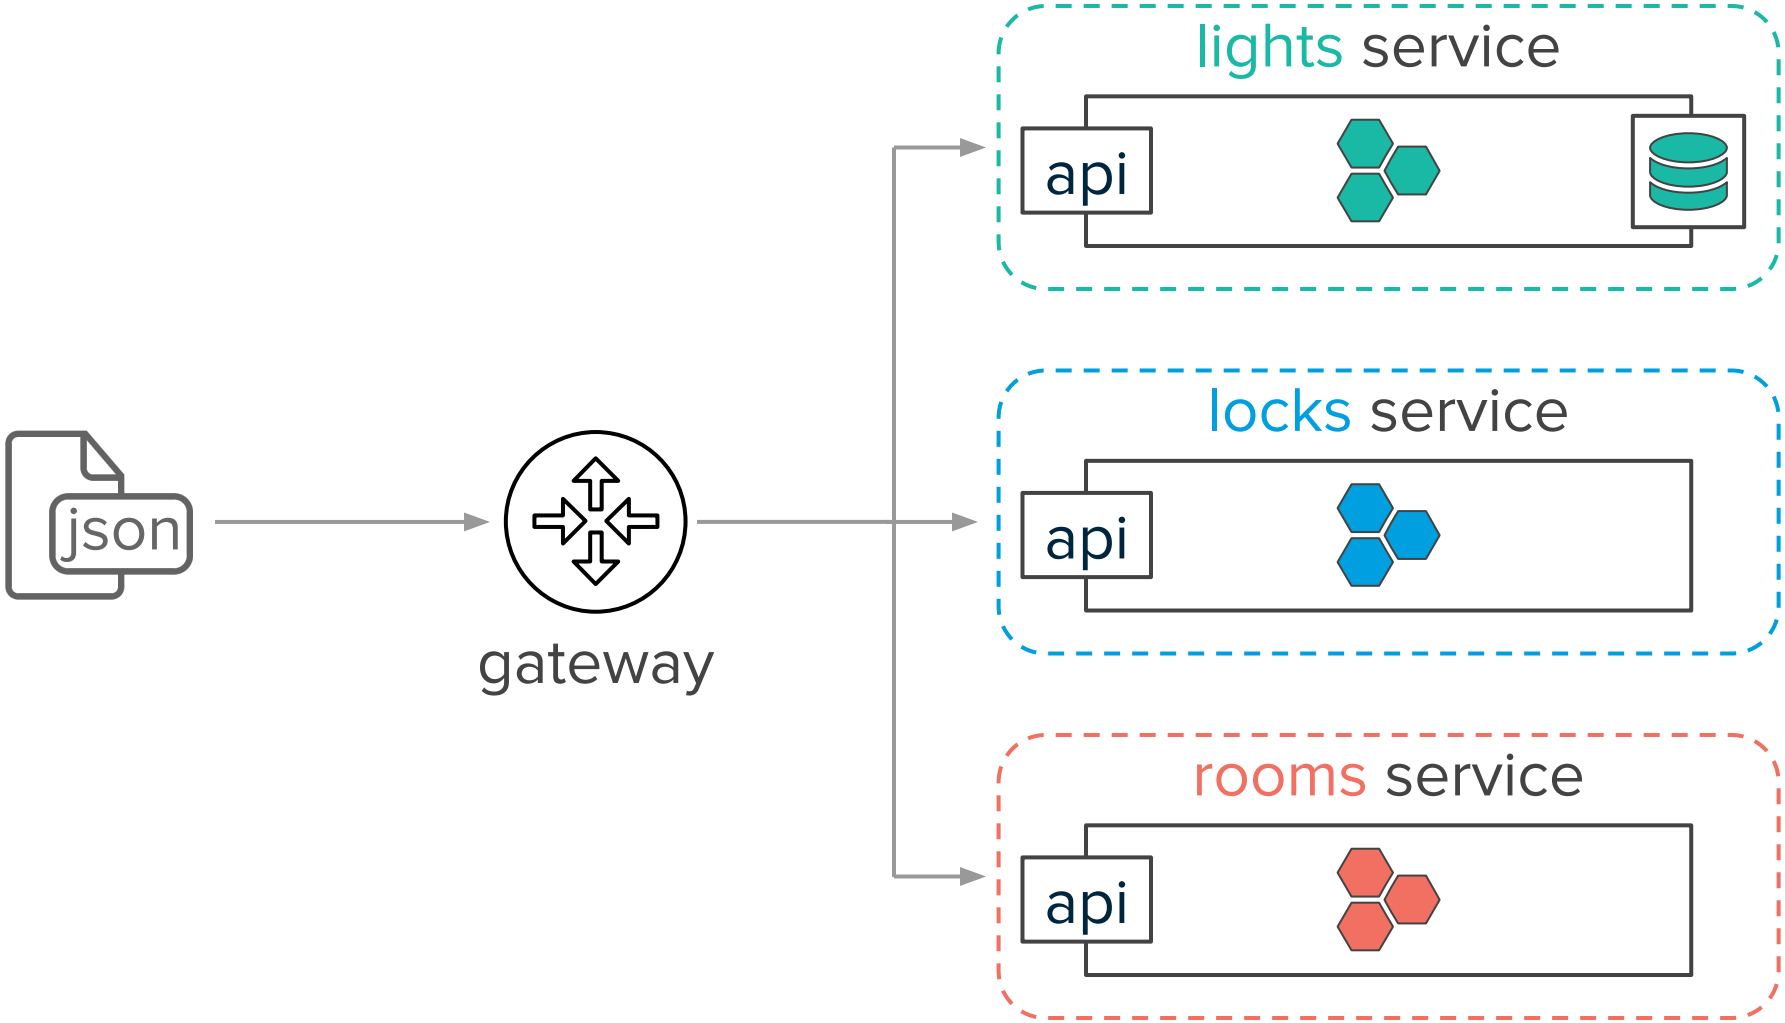
lights service
api
locks service
api
rooms service
api
json
gateway

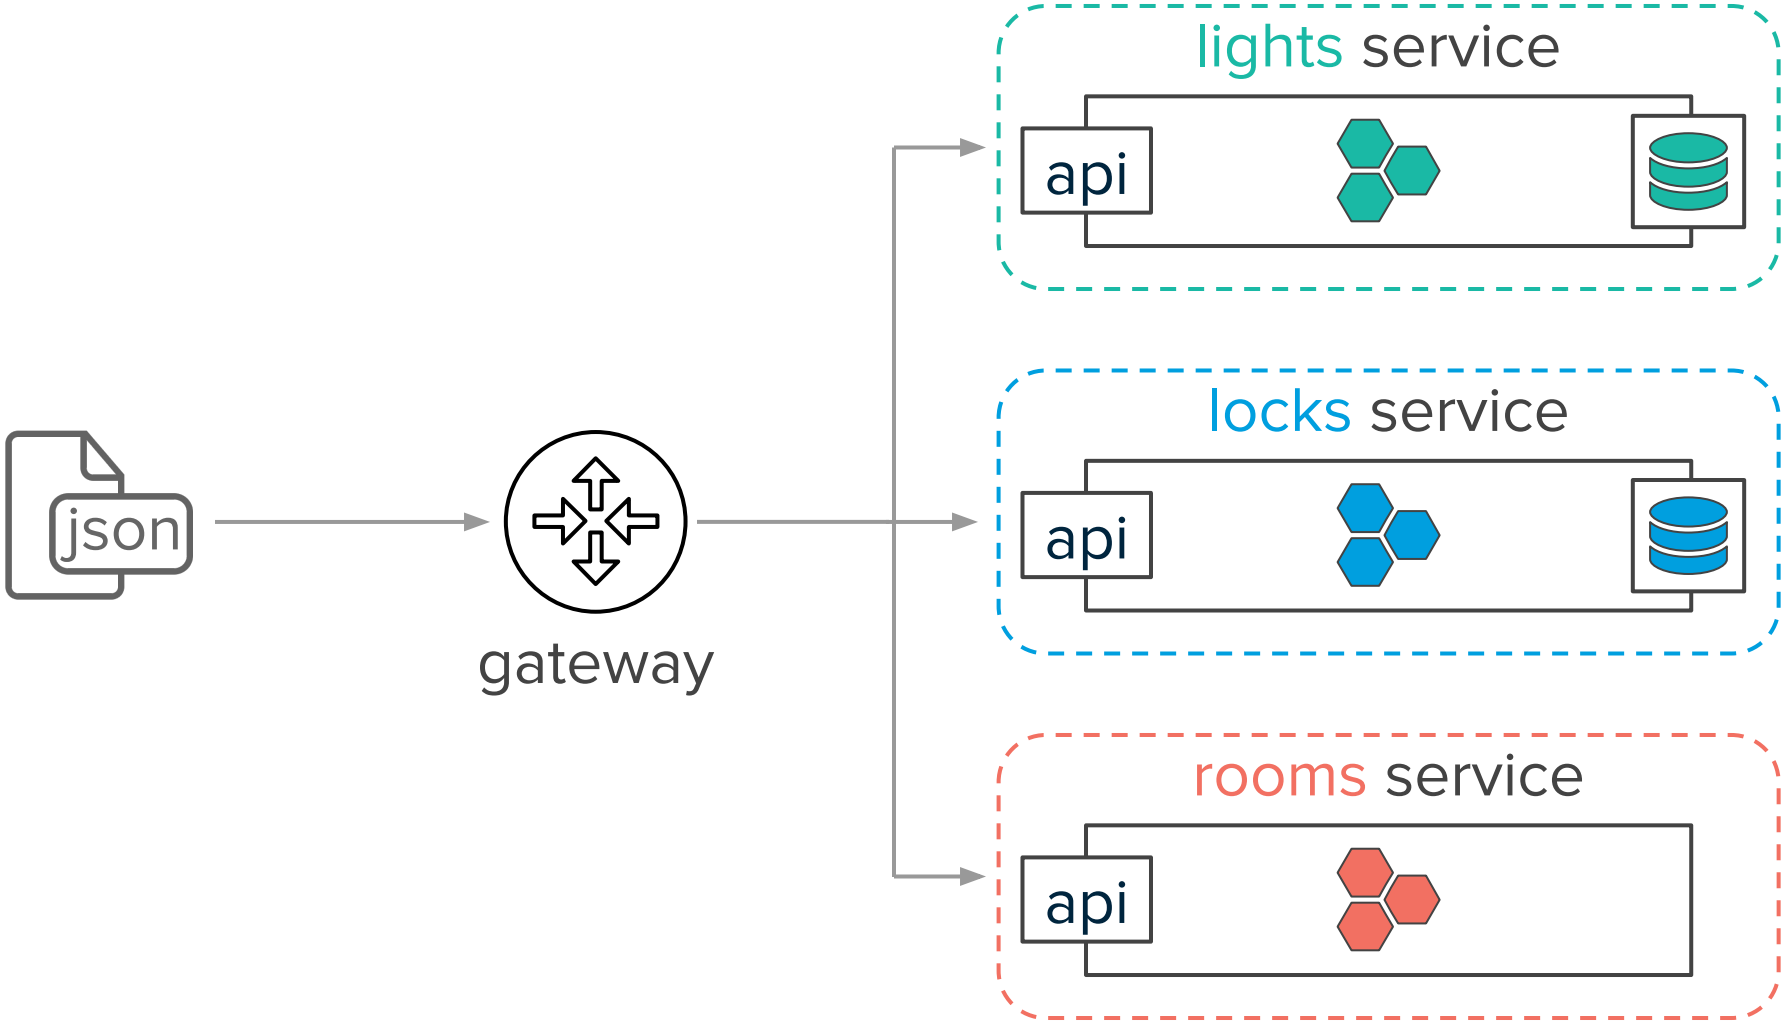
lights service
api
locks service
api
rooms service
api
json
gateway

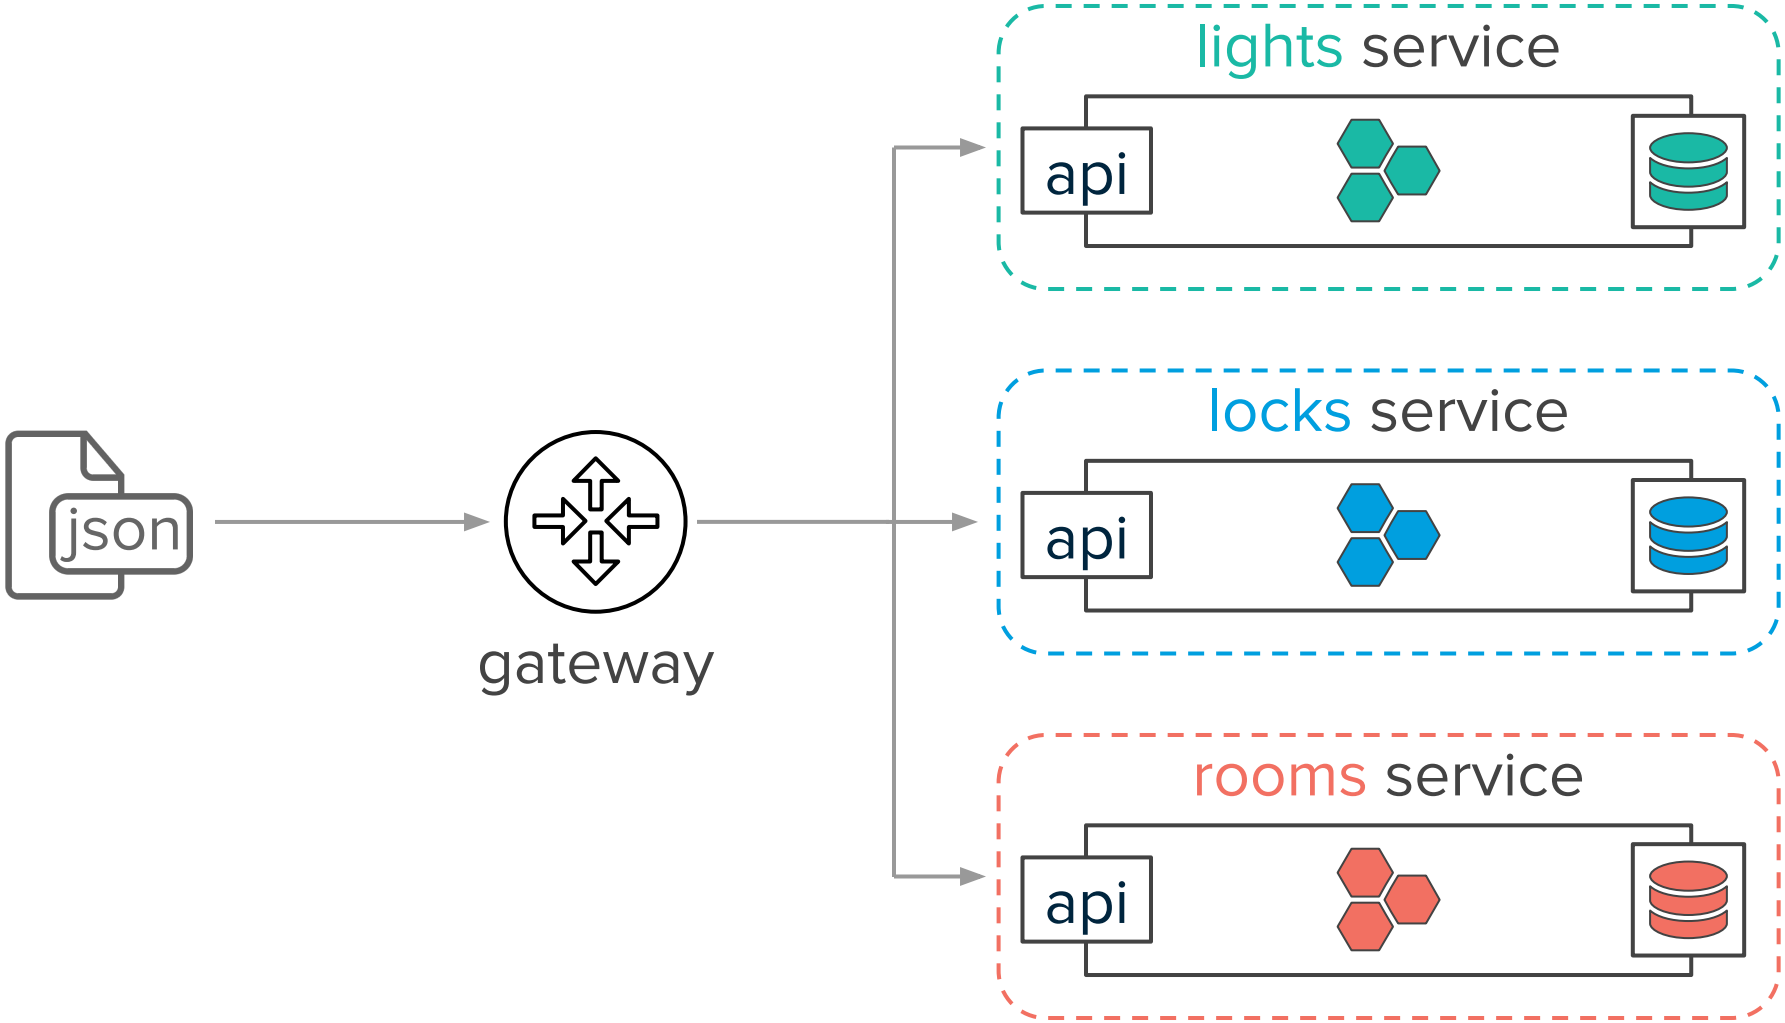
lights service
api
locks service
api
rooms service
api
json
gateway

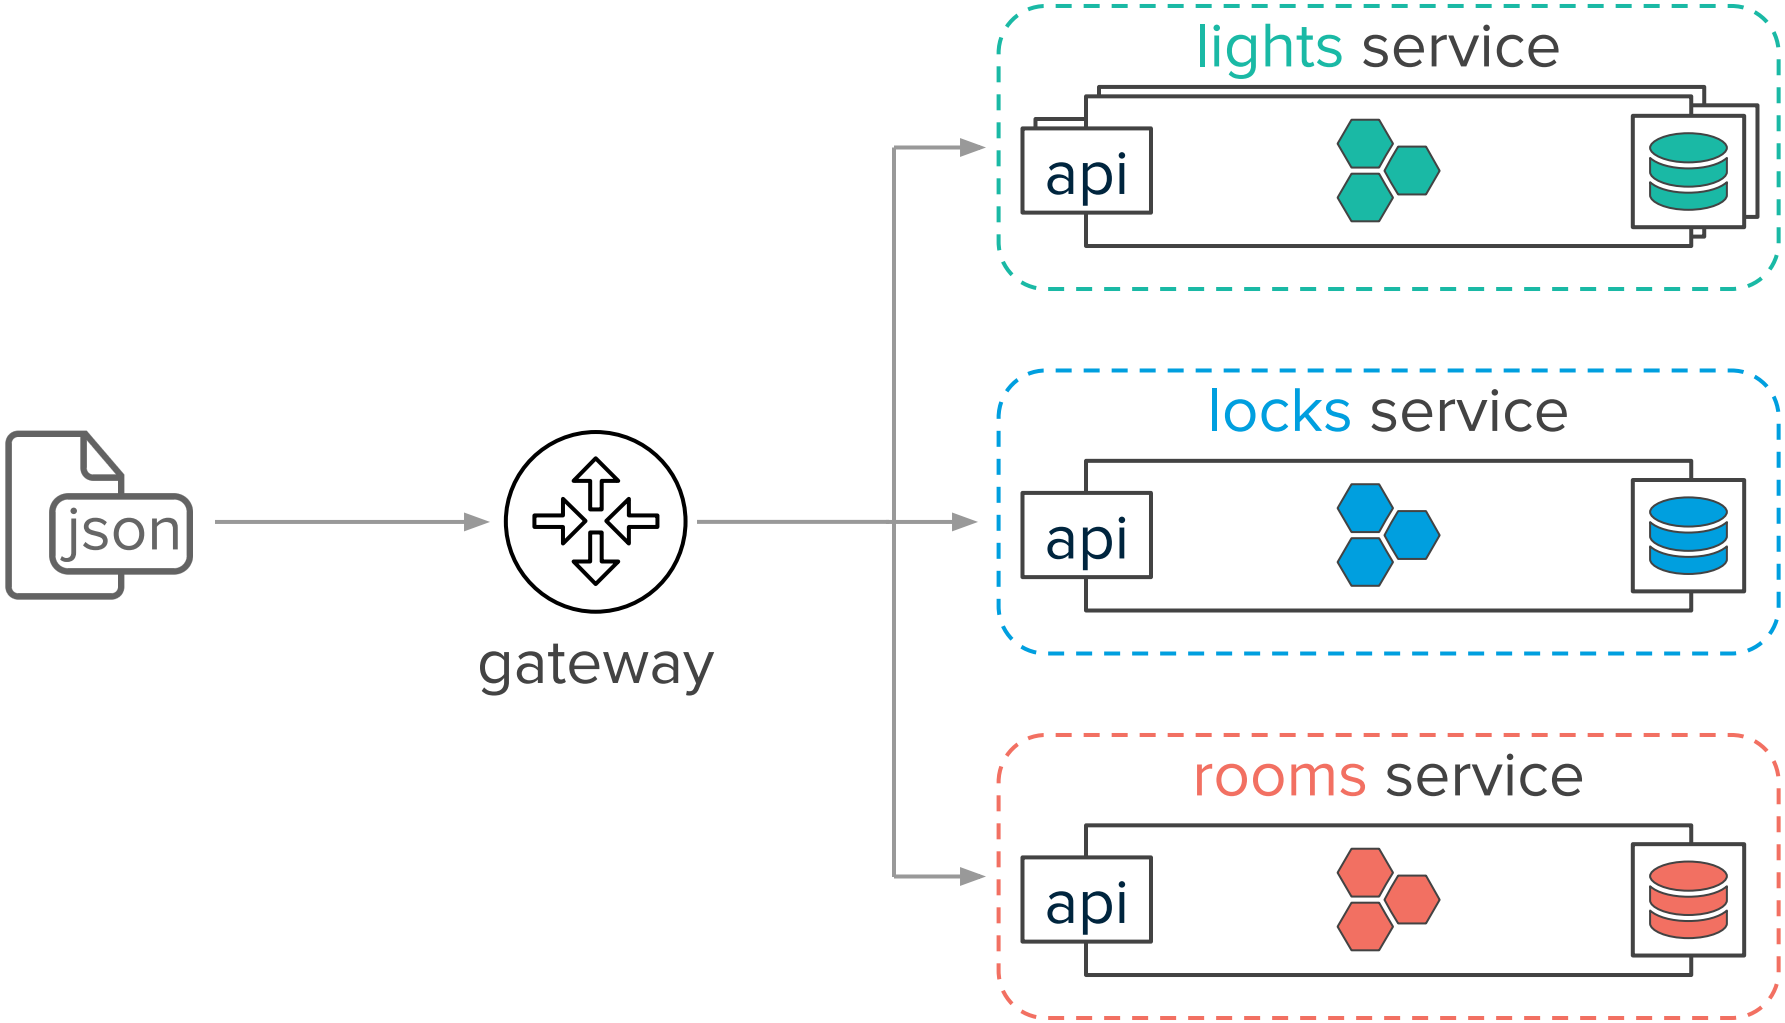
lights service
api
locks service
api
rooms service
api
json
gateway

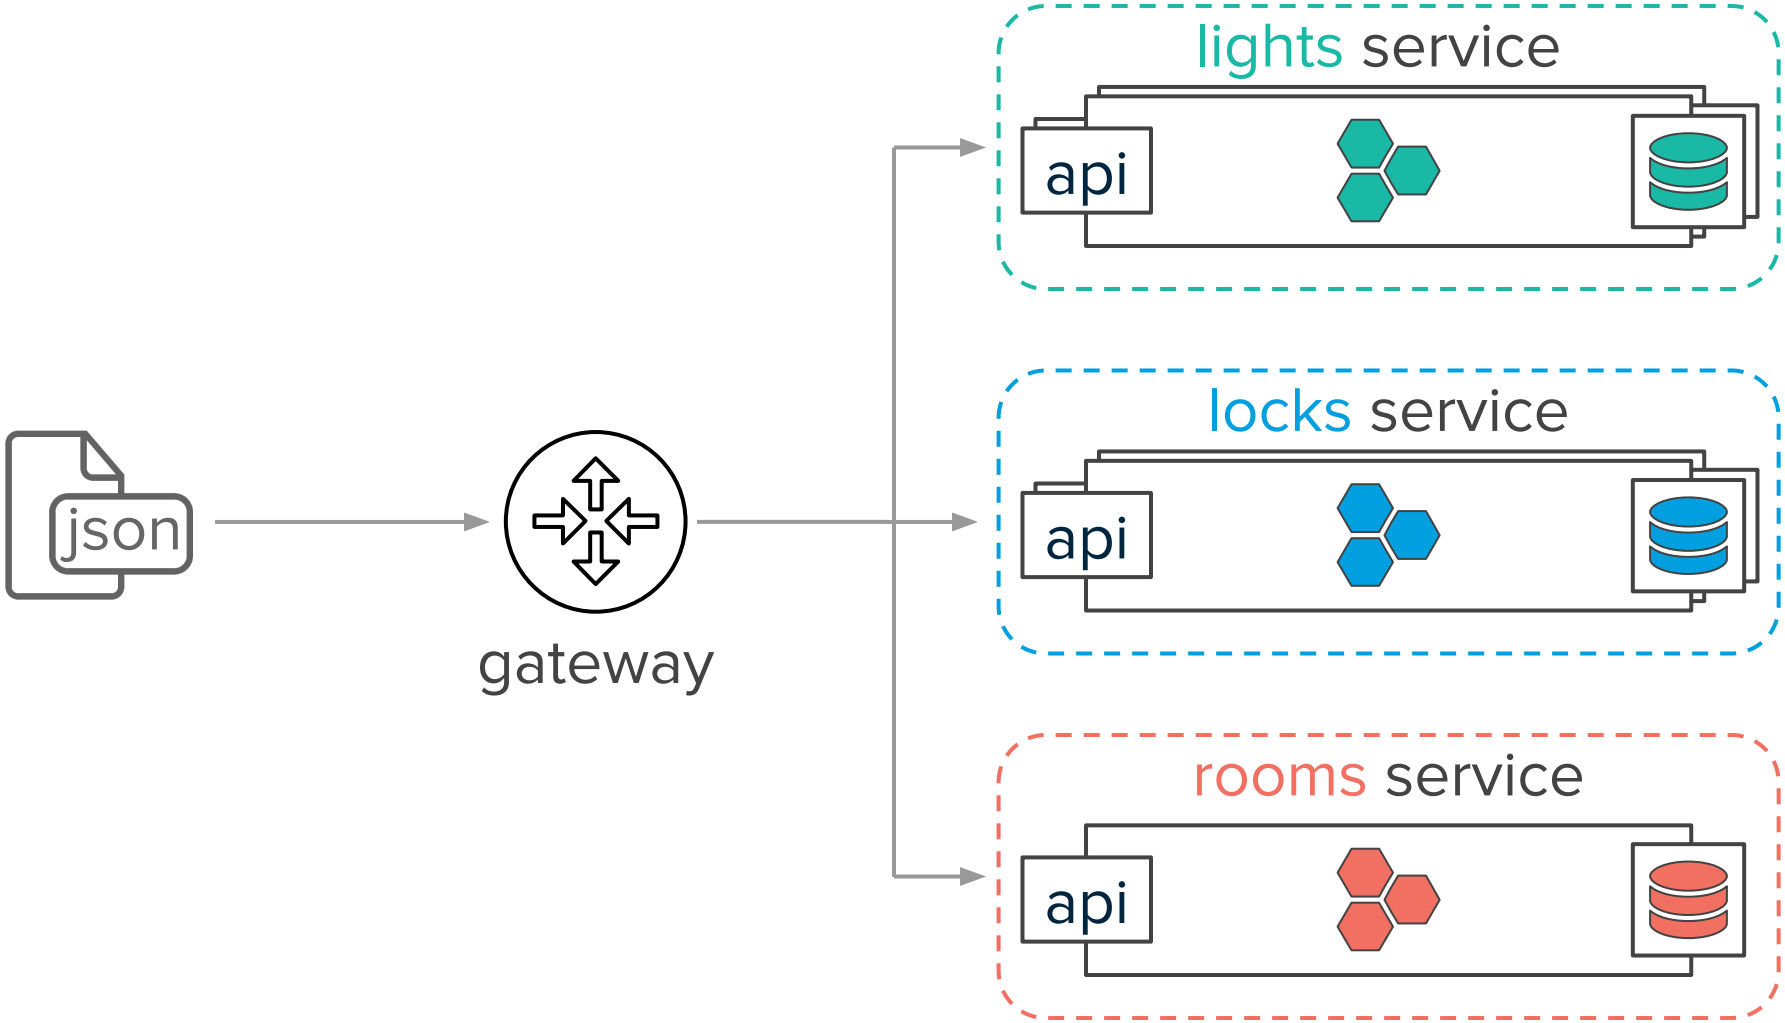
lights service
api
locks service
api
rooms service
api
json
gateway

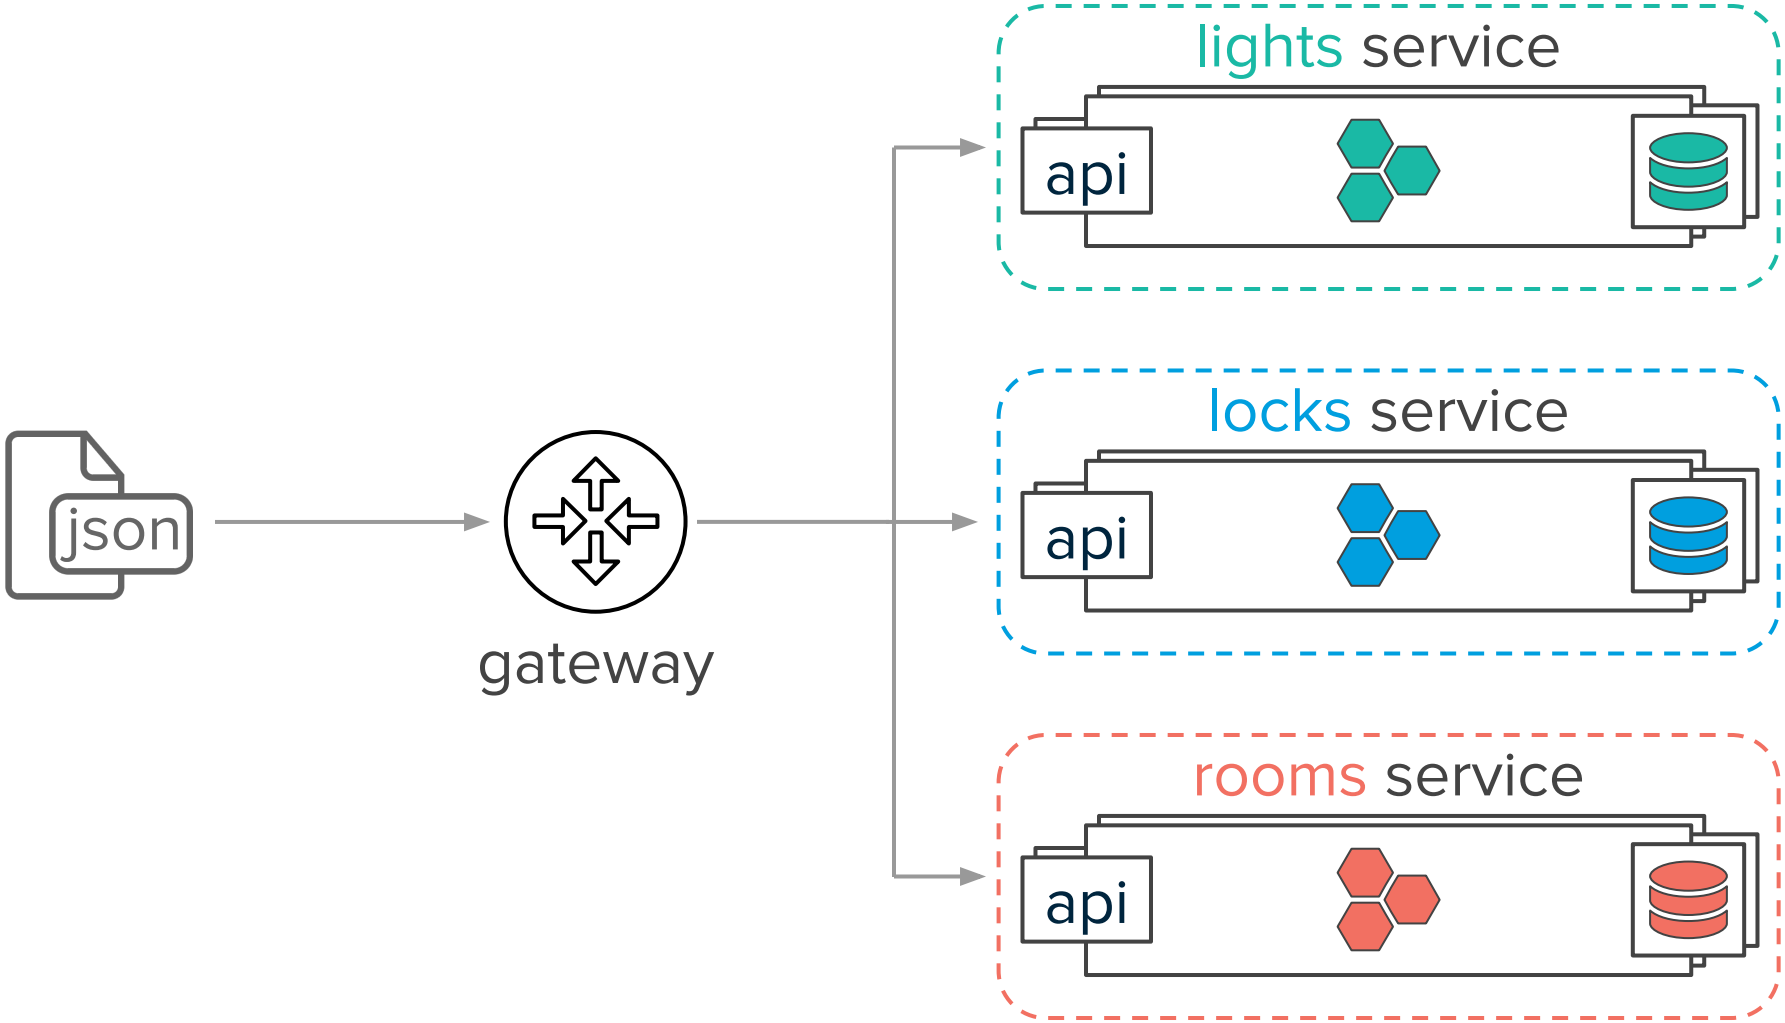
lights service
api
locks service
api
rooms service
api
json
gateway

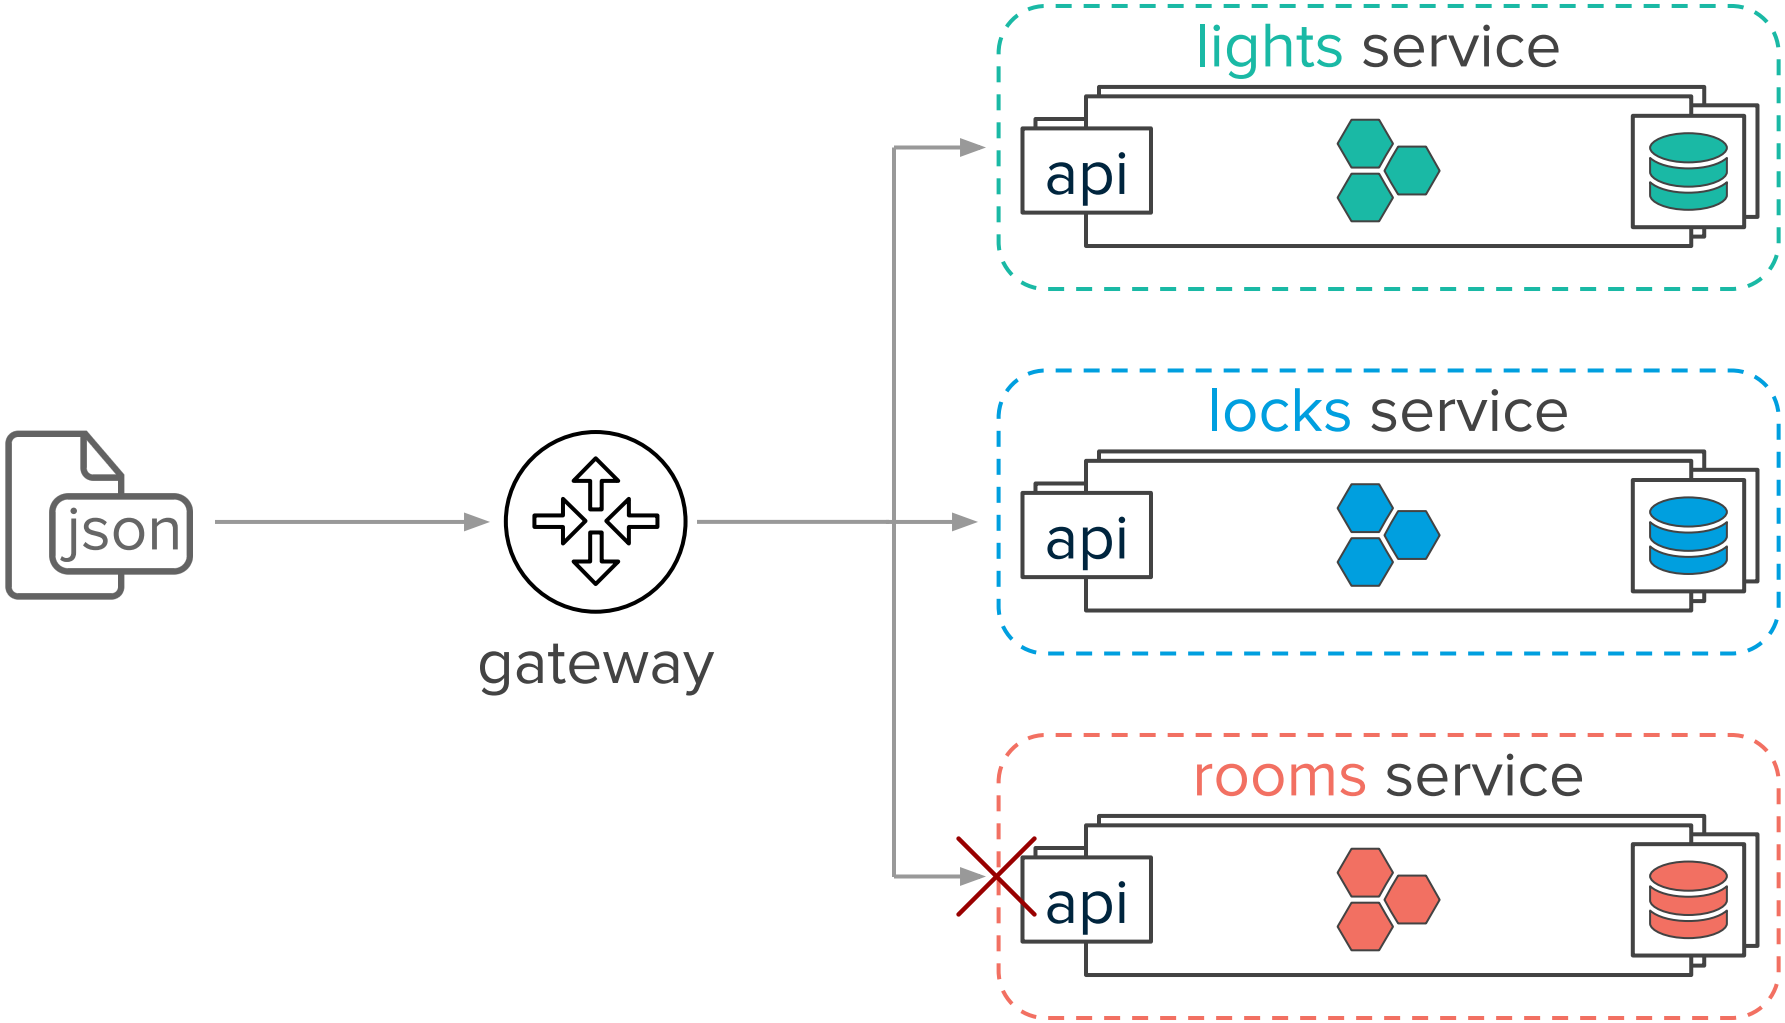
lights service
api
locks service
api
rooms service
api
json
gateway

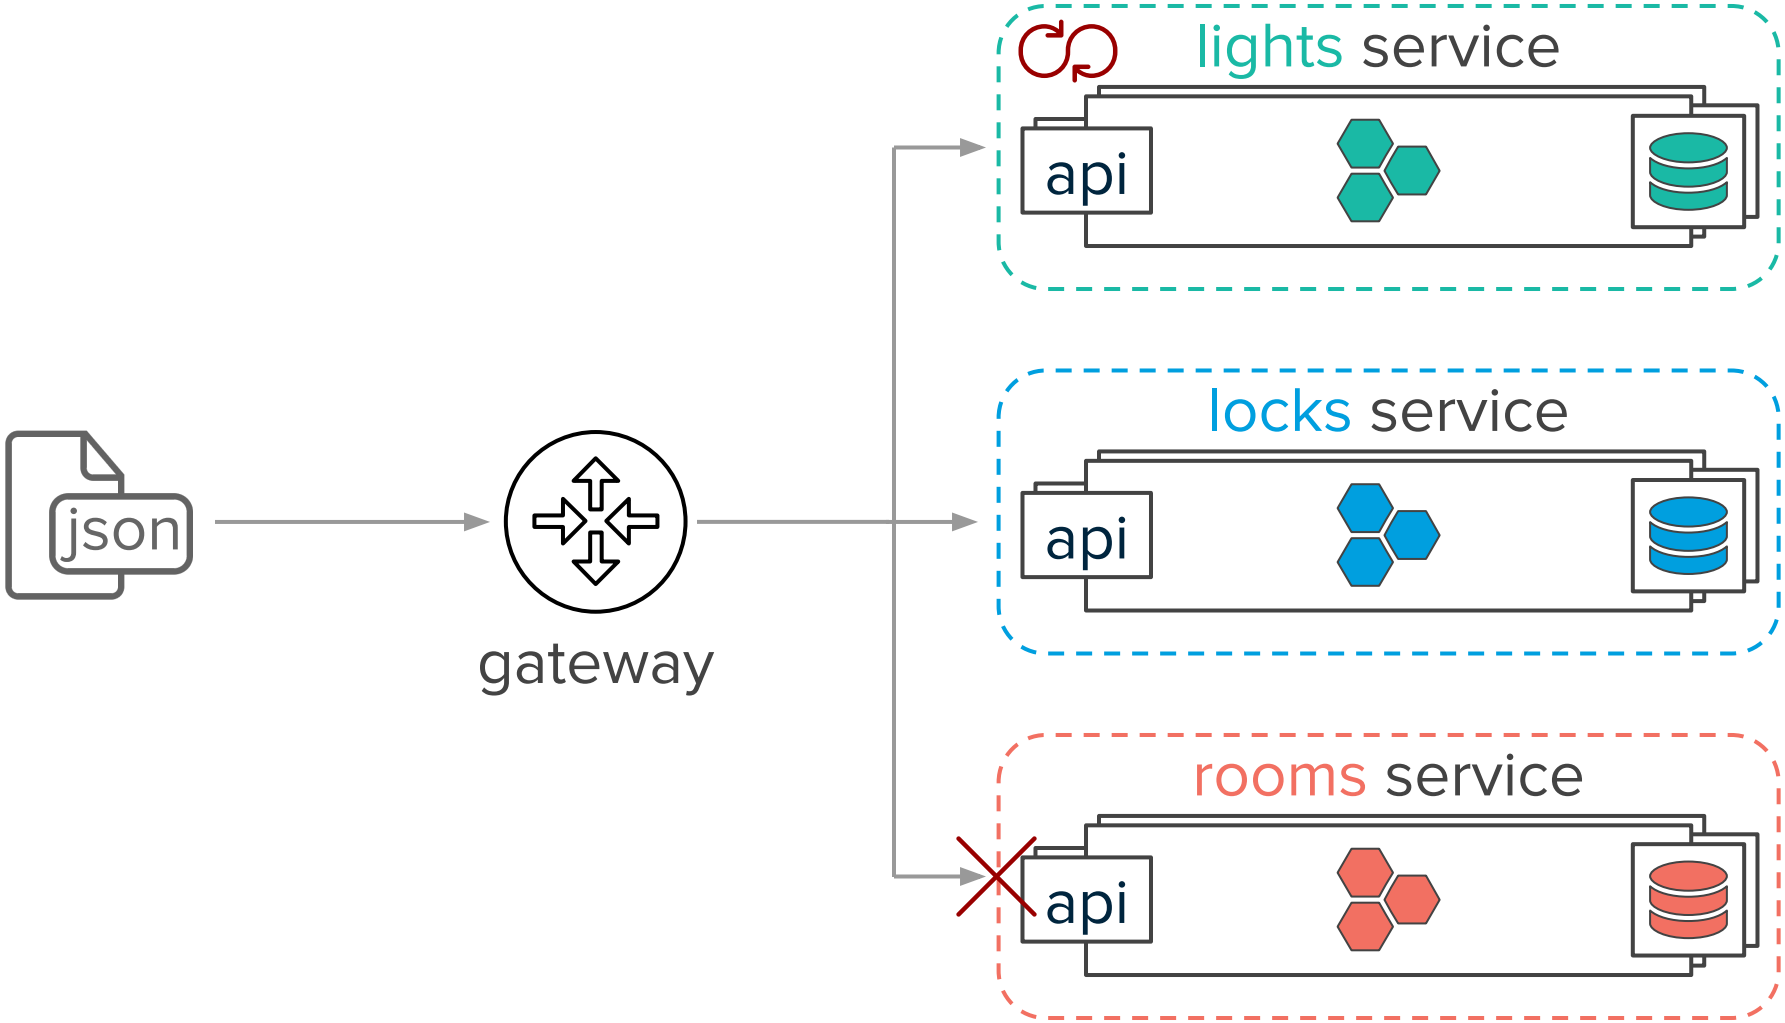
lights service
api
locks service
api
rooms service
api
json
gateway

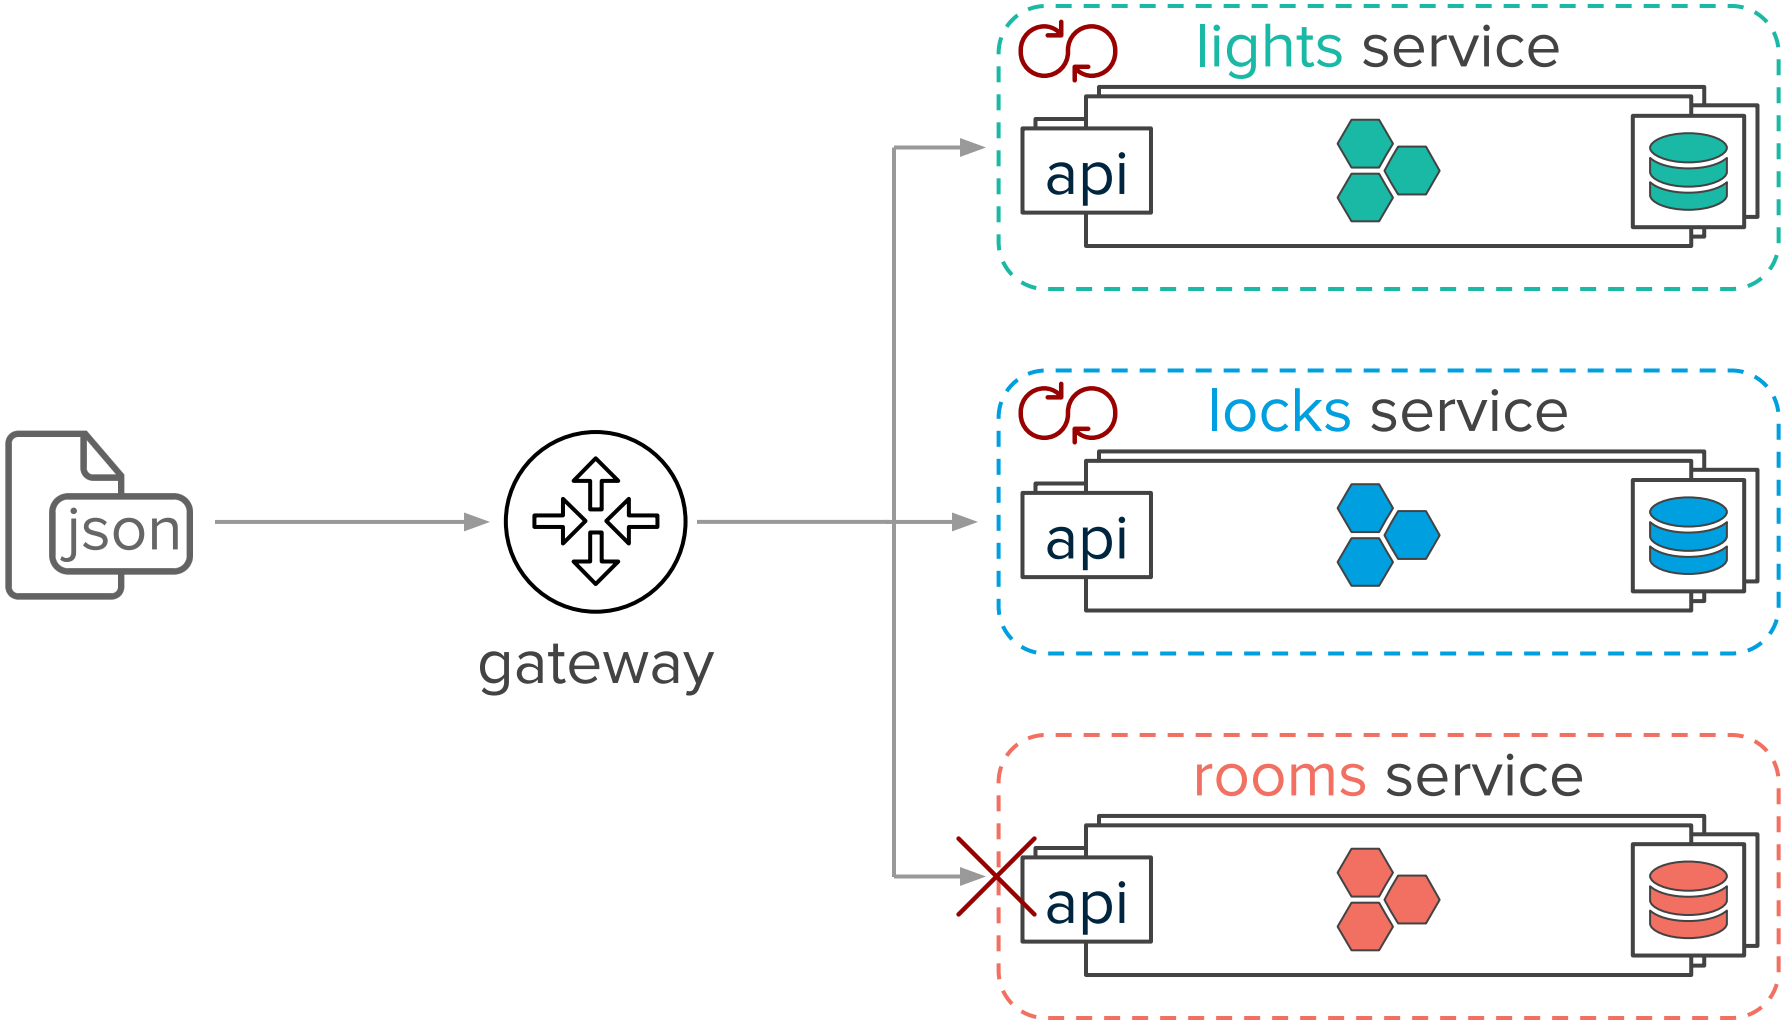
lights service
api
locks service
api
rooms service
api
json
gateway

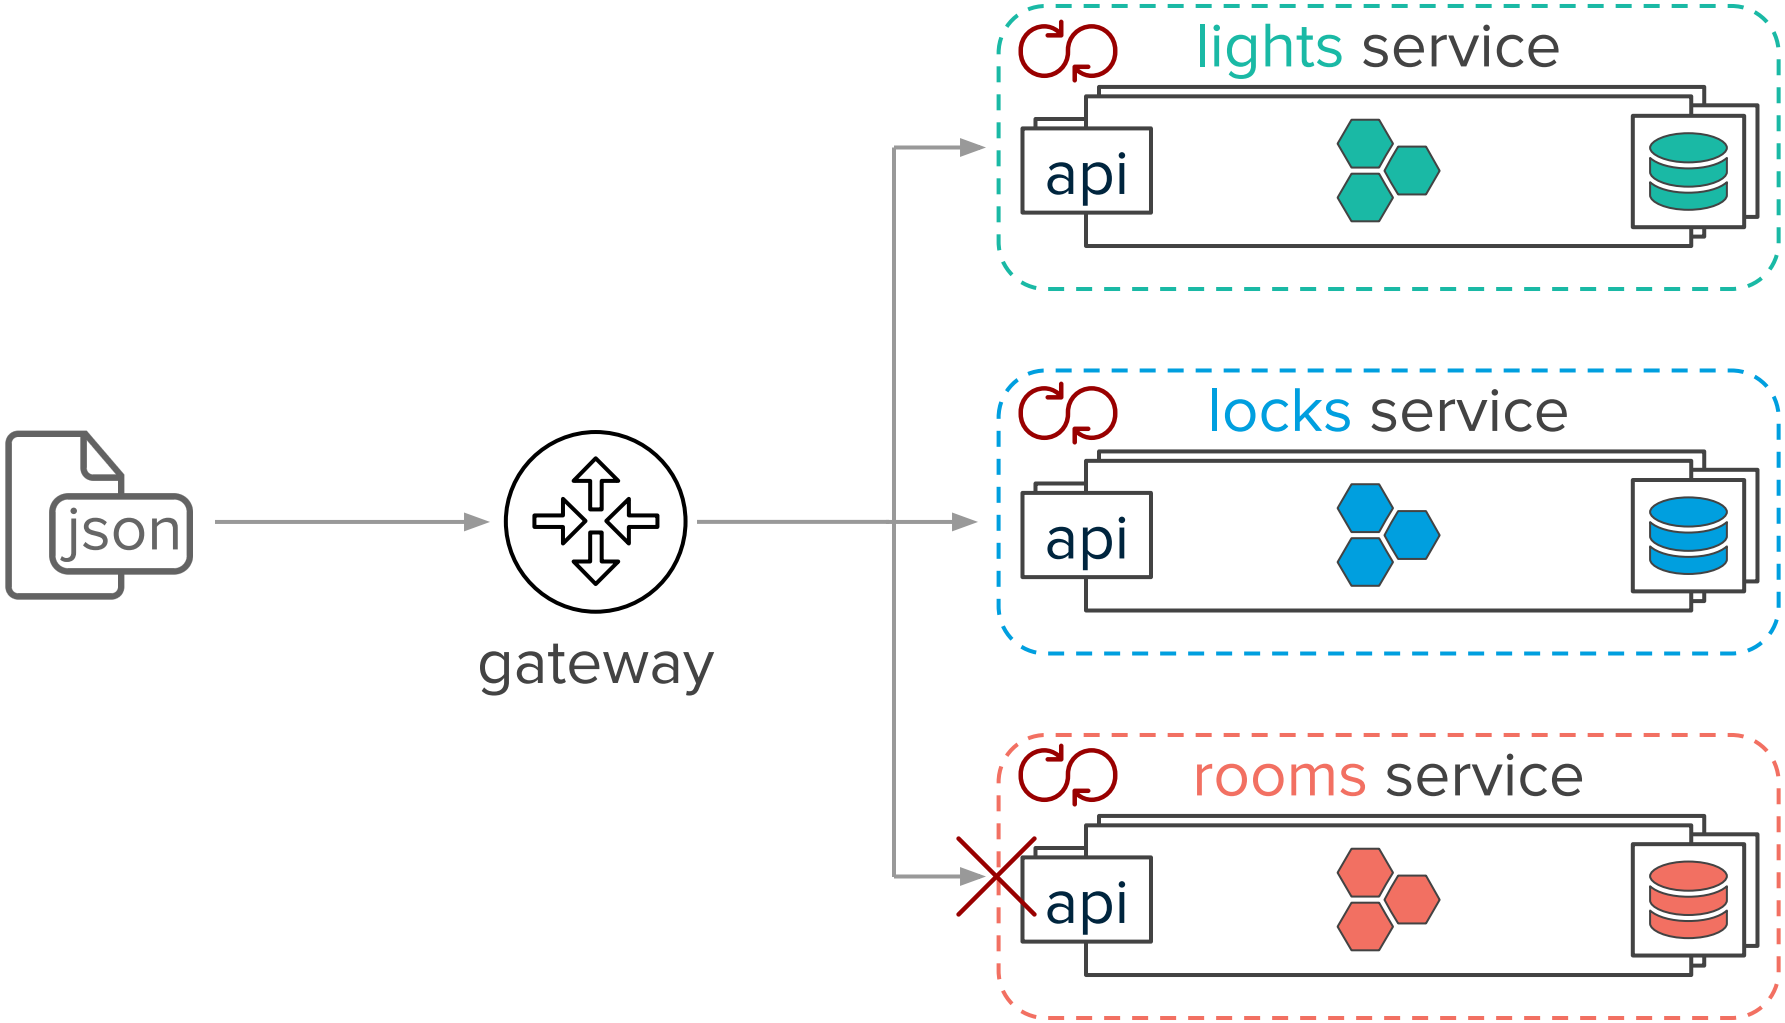
lights service
api
locks service
api
rooms service
api
json
gateway

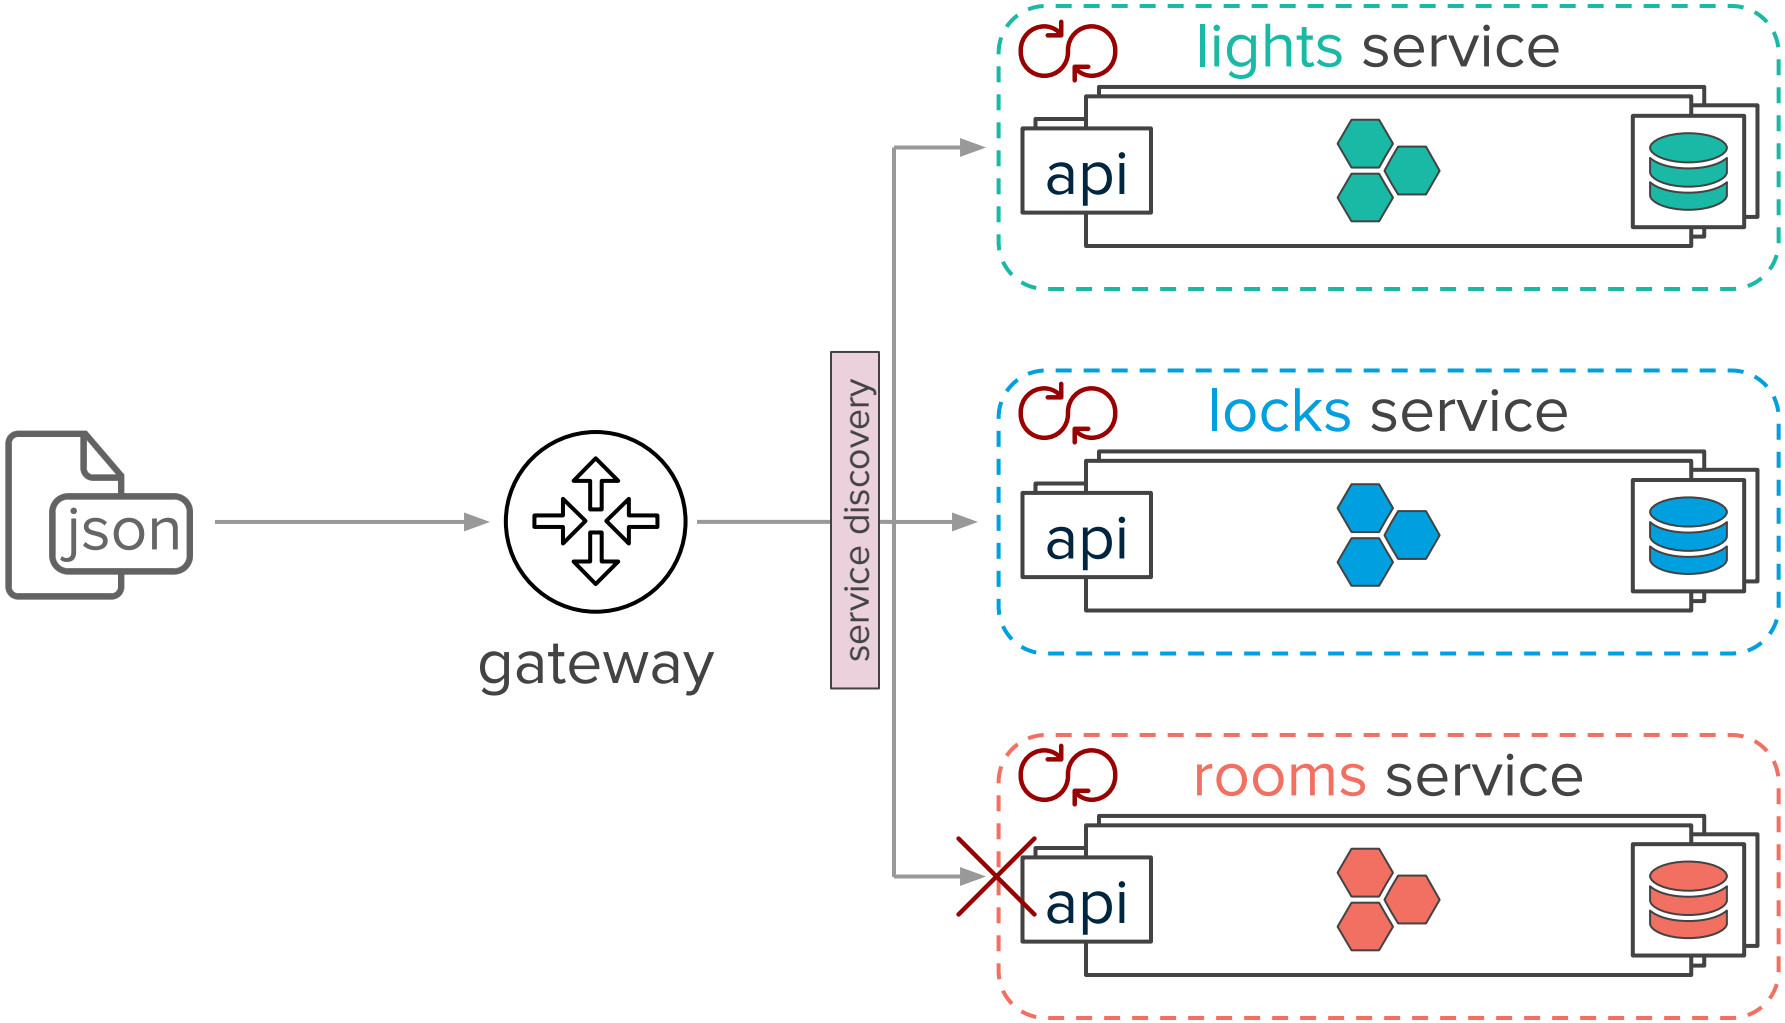
lights service
api
locks service
api
rooms service
api
service discovery
json
gateway

Distractions ...

- Storage
- High Availability
- Reliability
- The API contracts of each service (e.g. company standards)
- Team collaboration over APIs

These all take time and effort away from the main focus of each team, which is to provide the best service they can

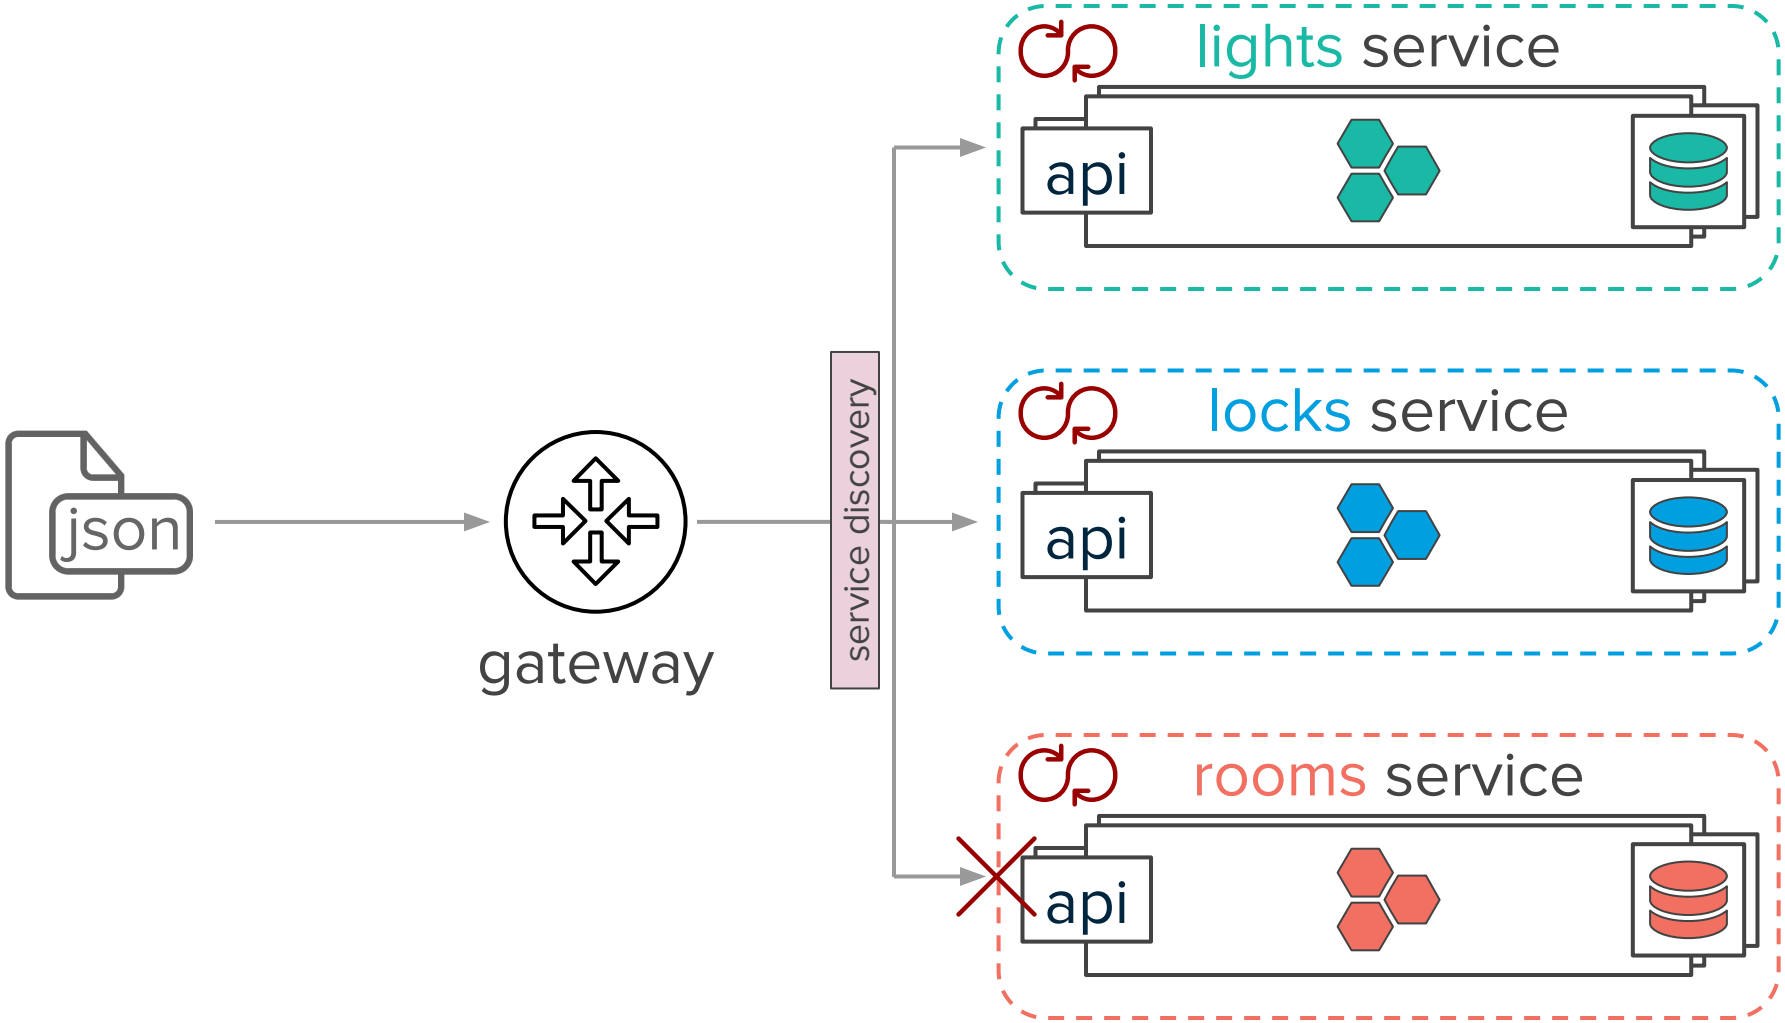
lights service
api
locks service
api
rooms service
api
service discovery
json
gateway

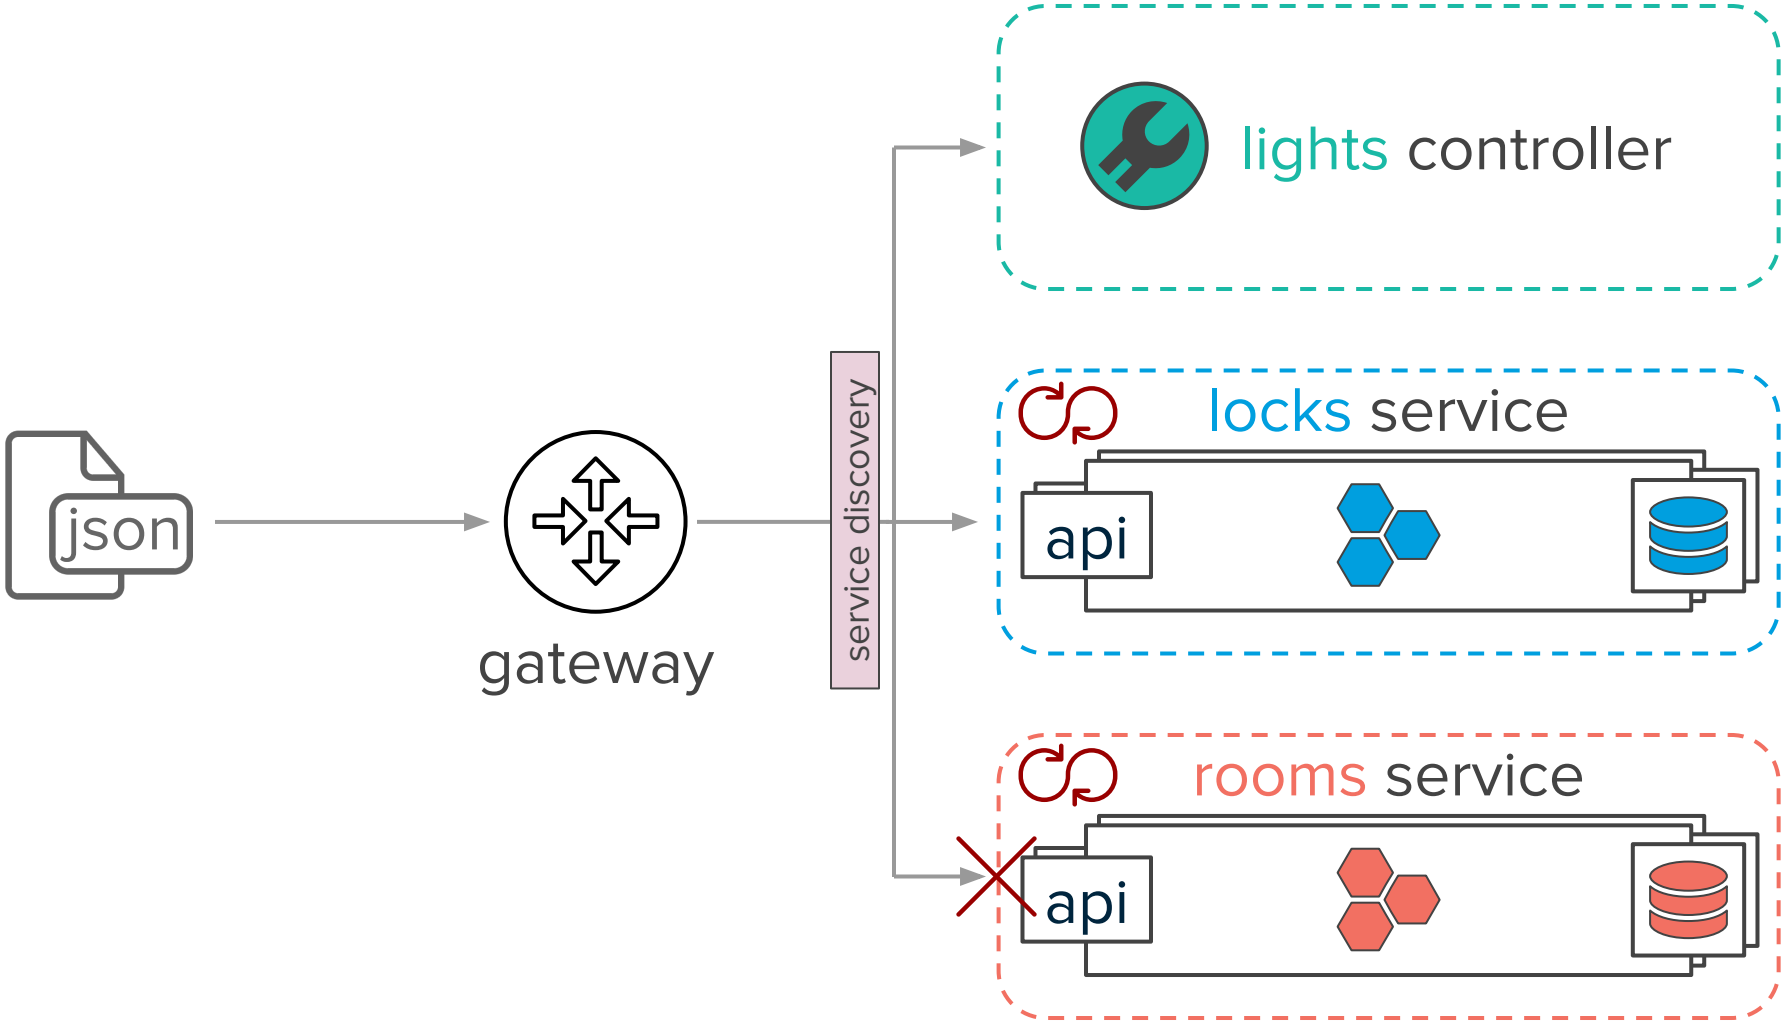
lights controller
locks service
api
rooms service
api
service discovery
json
gateway

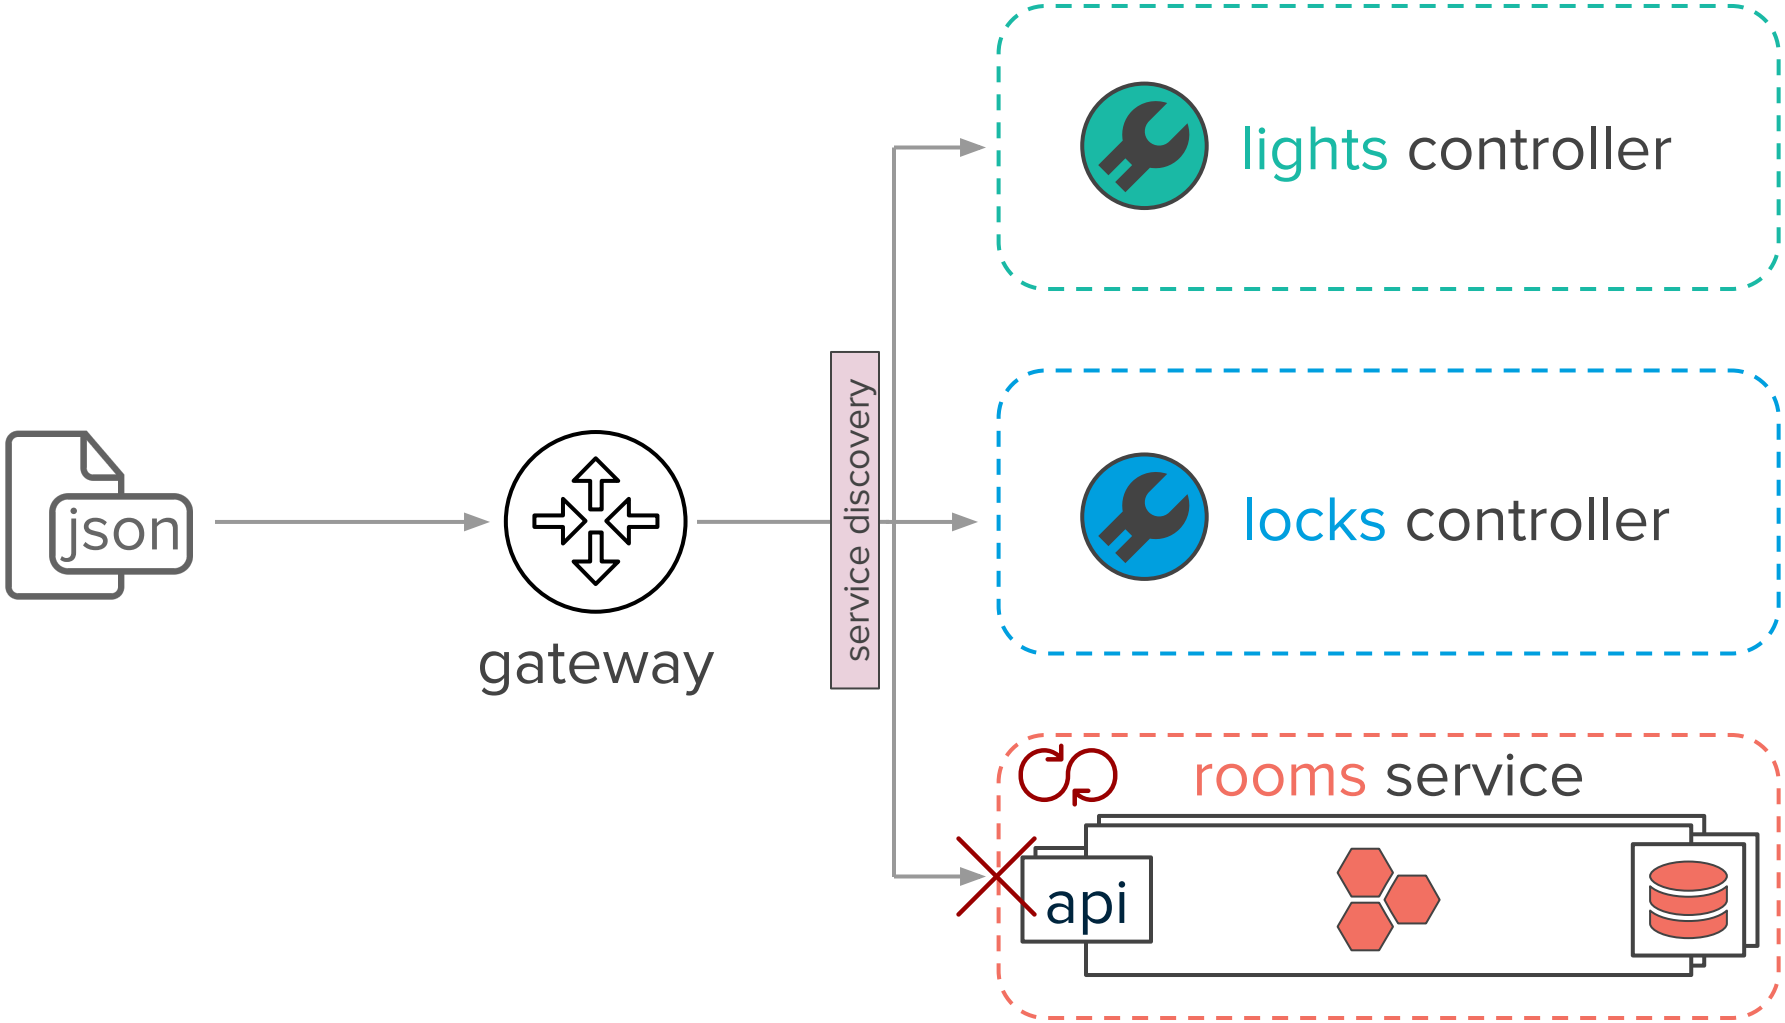
lights controller
locks controller
rooms service
api
service discovery
gateway
json

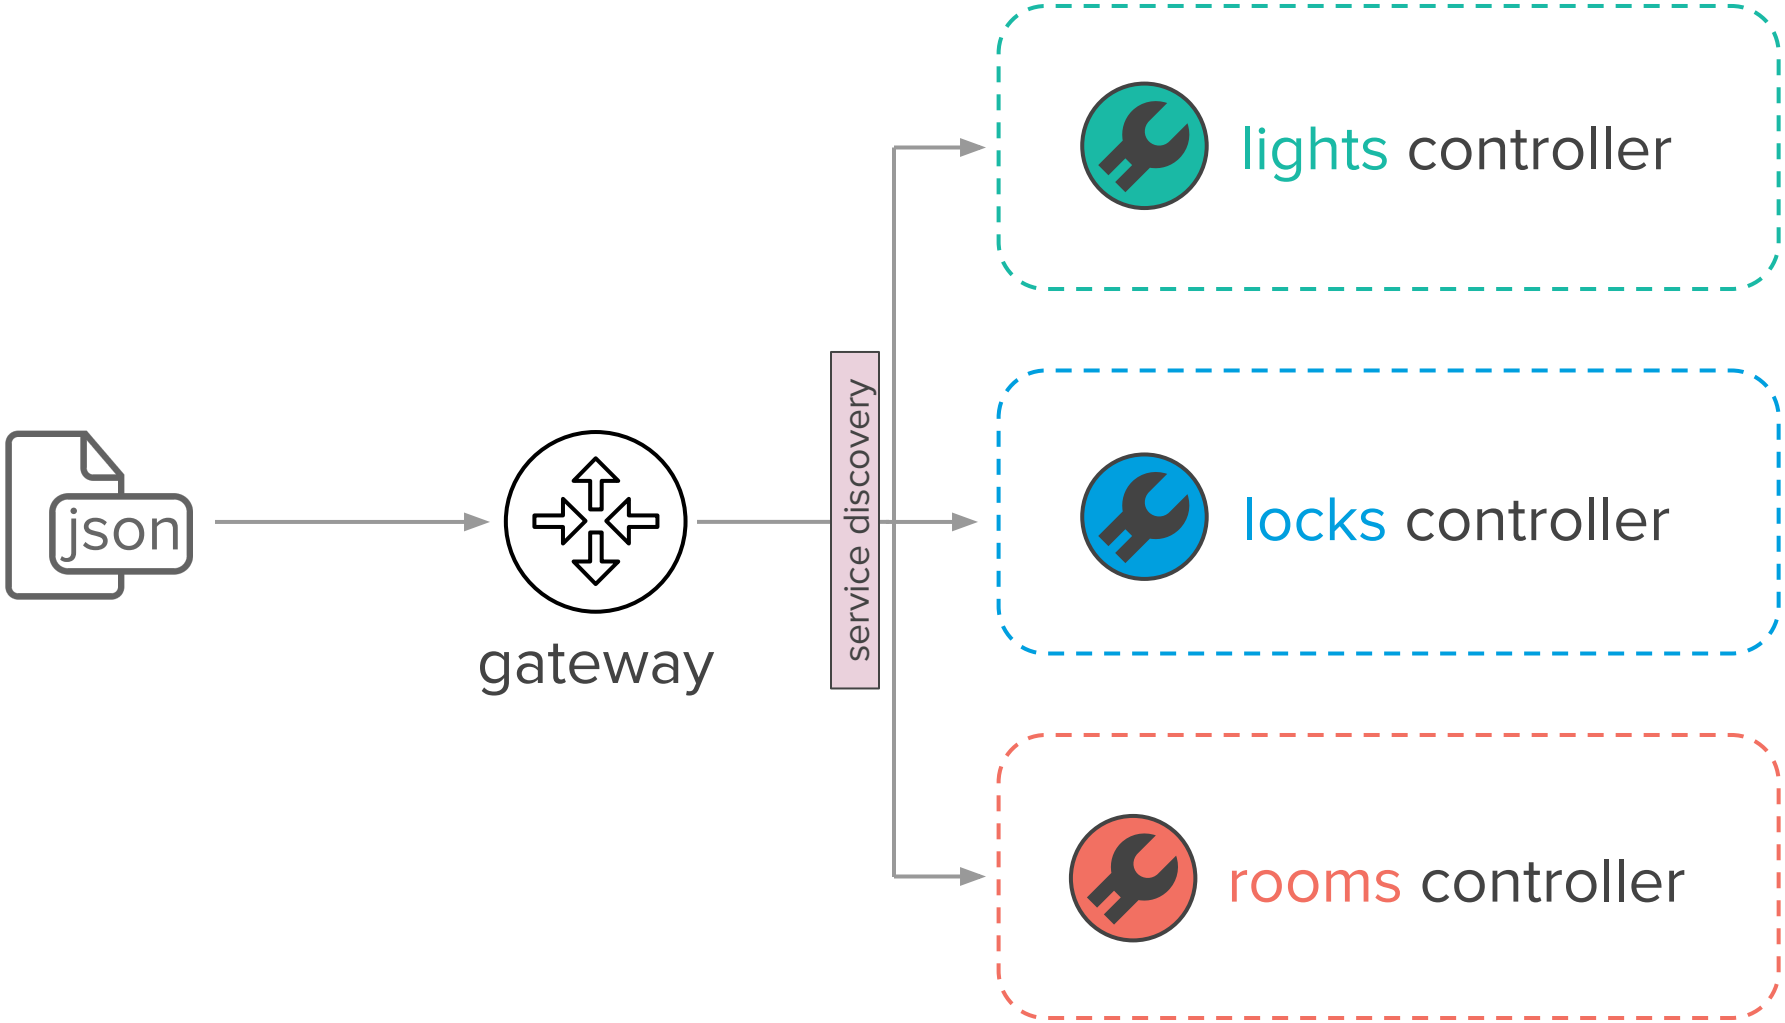
lights controller
locks controller
rooms controller
service discovery
gateway
json

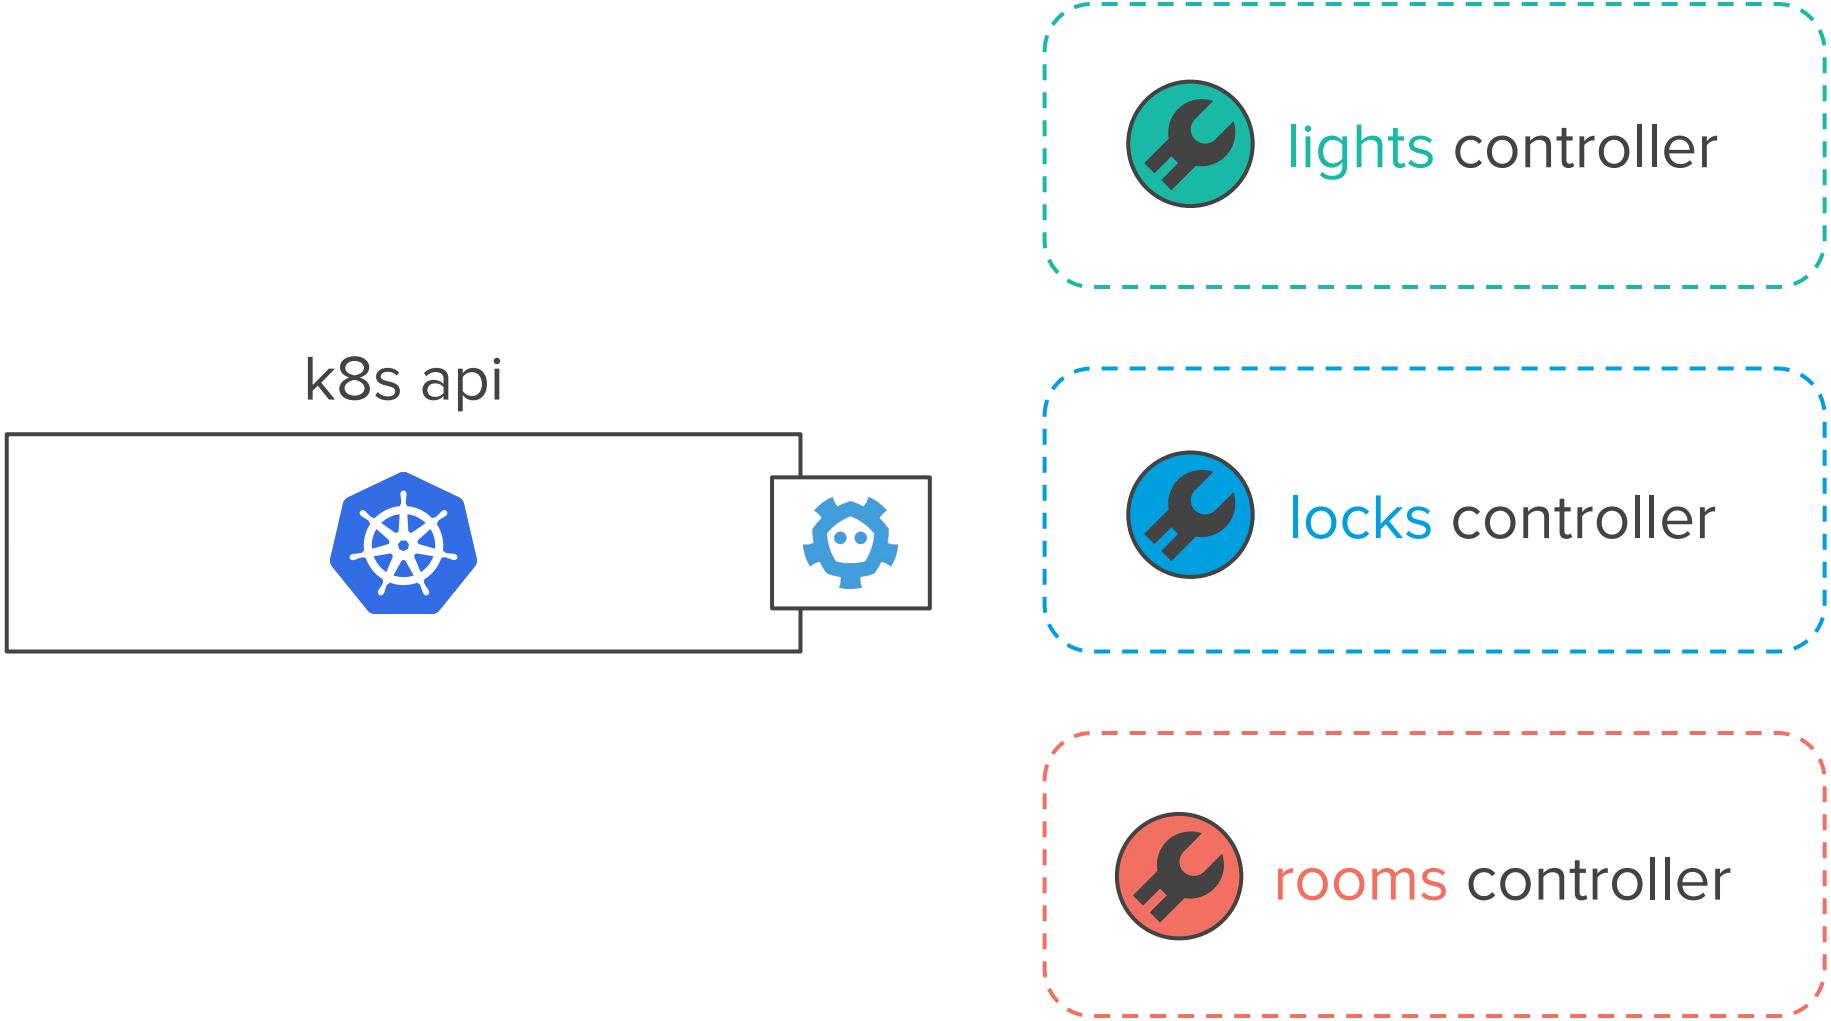
k8s api
lights controller
locks controller
rooms controller

“I want the kitchen lights to be on”

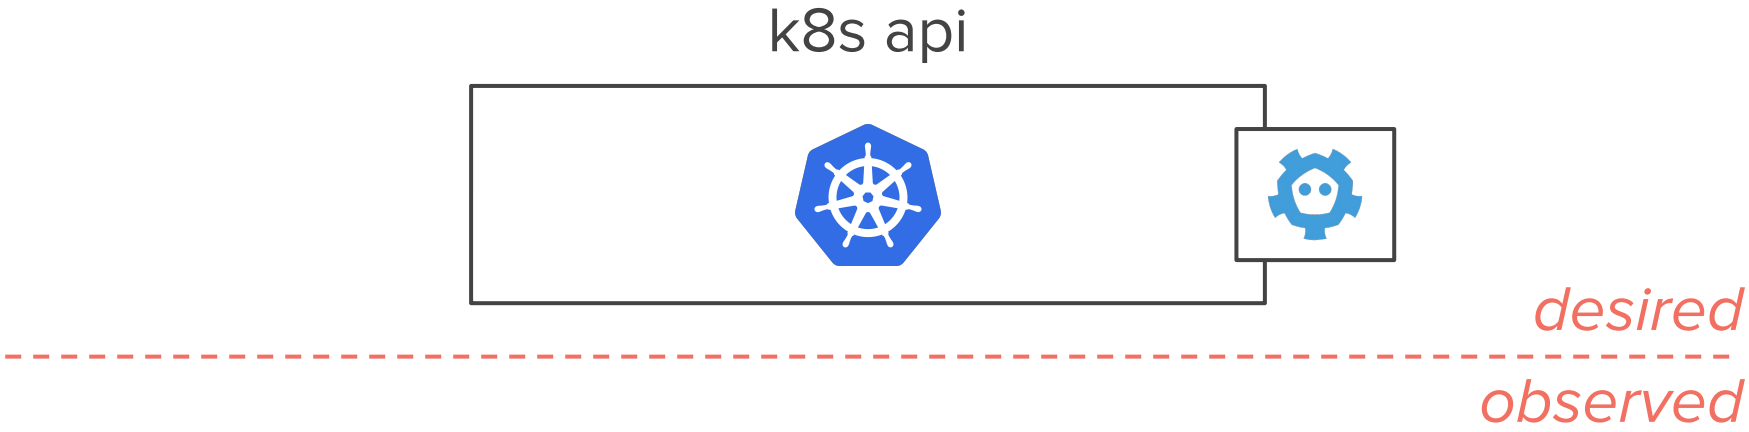

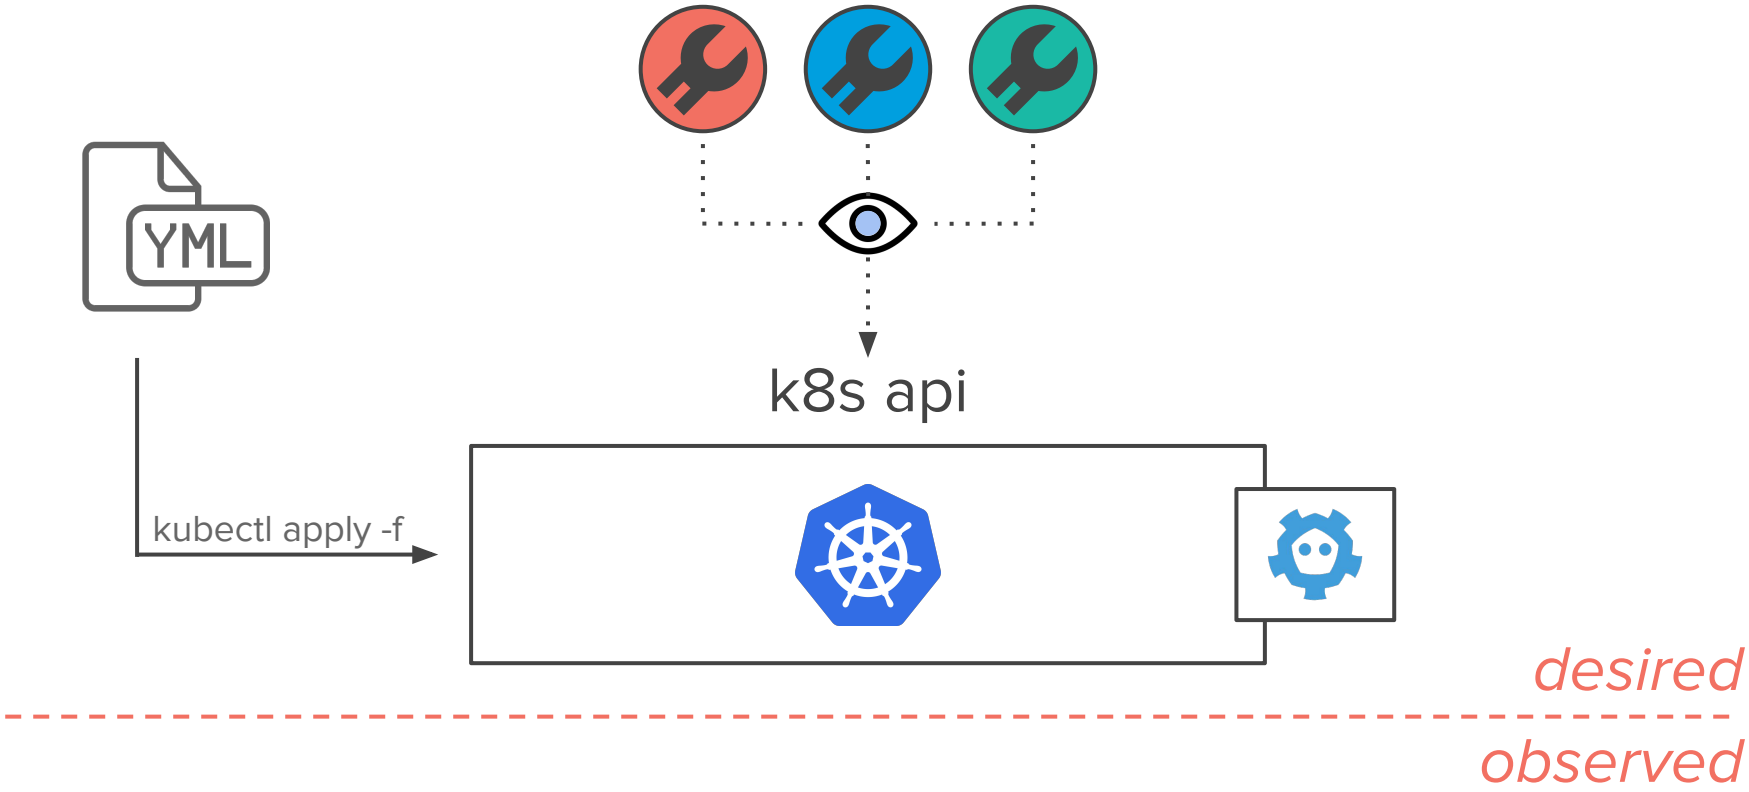
YML
kubectl apply -f
k8s api
desired
observed


```
apiVersion: v1
kind: Room
metadata:
  name: kitchen
  namespace: default
spec:
  lights:
  - name: lamp-1
    brightness: 0.5
  - name: lamp-2
    brightness: 1.0
```

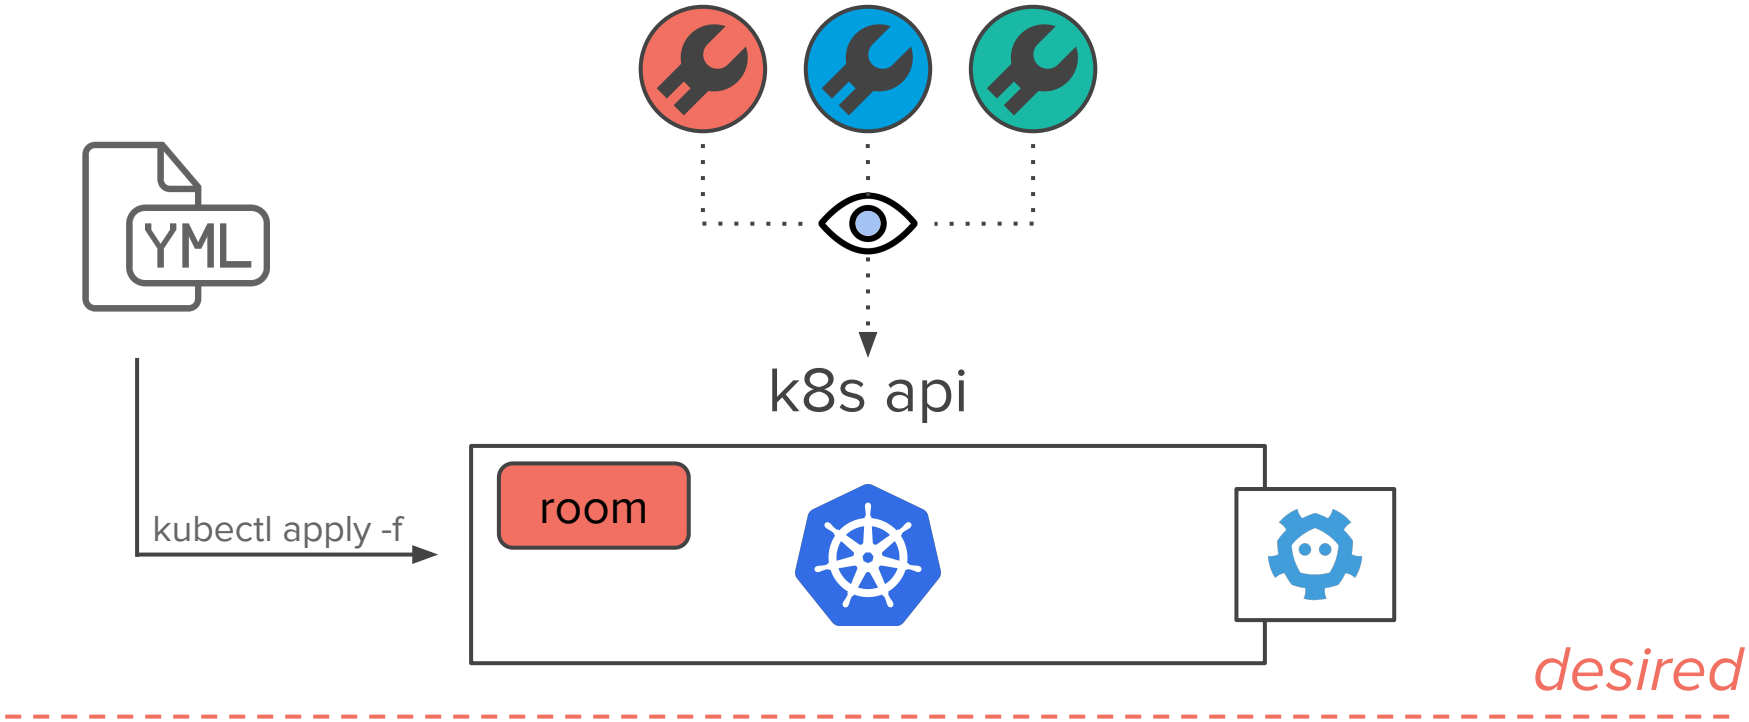
YML
kubectl apply -f
k8s api
room
desired

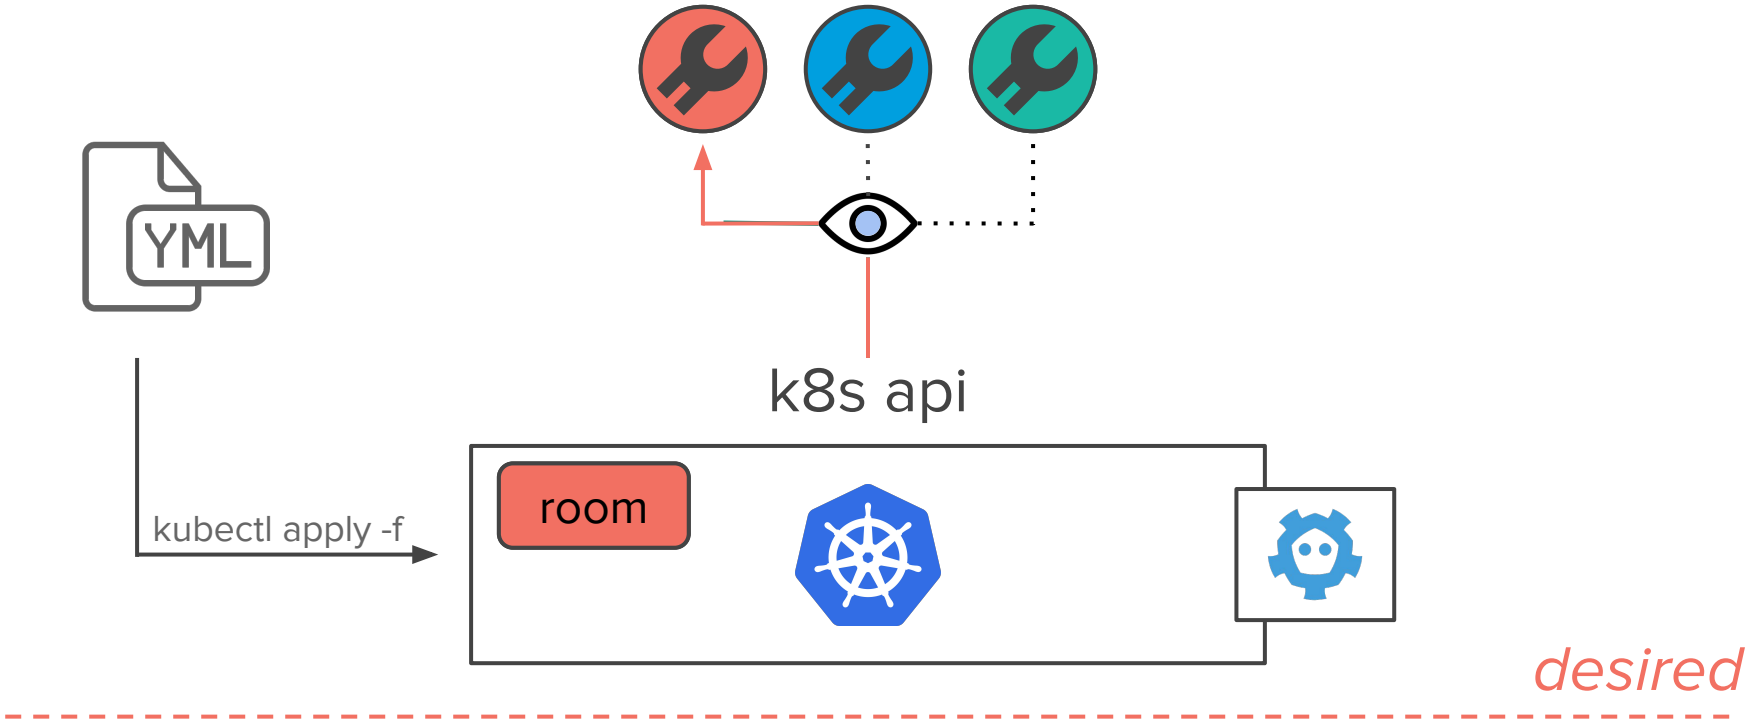
YML
kubectl apply -f
k8s api
room
desired

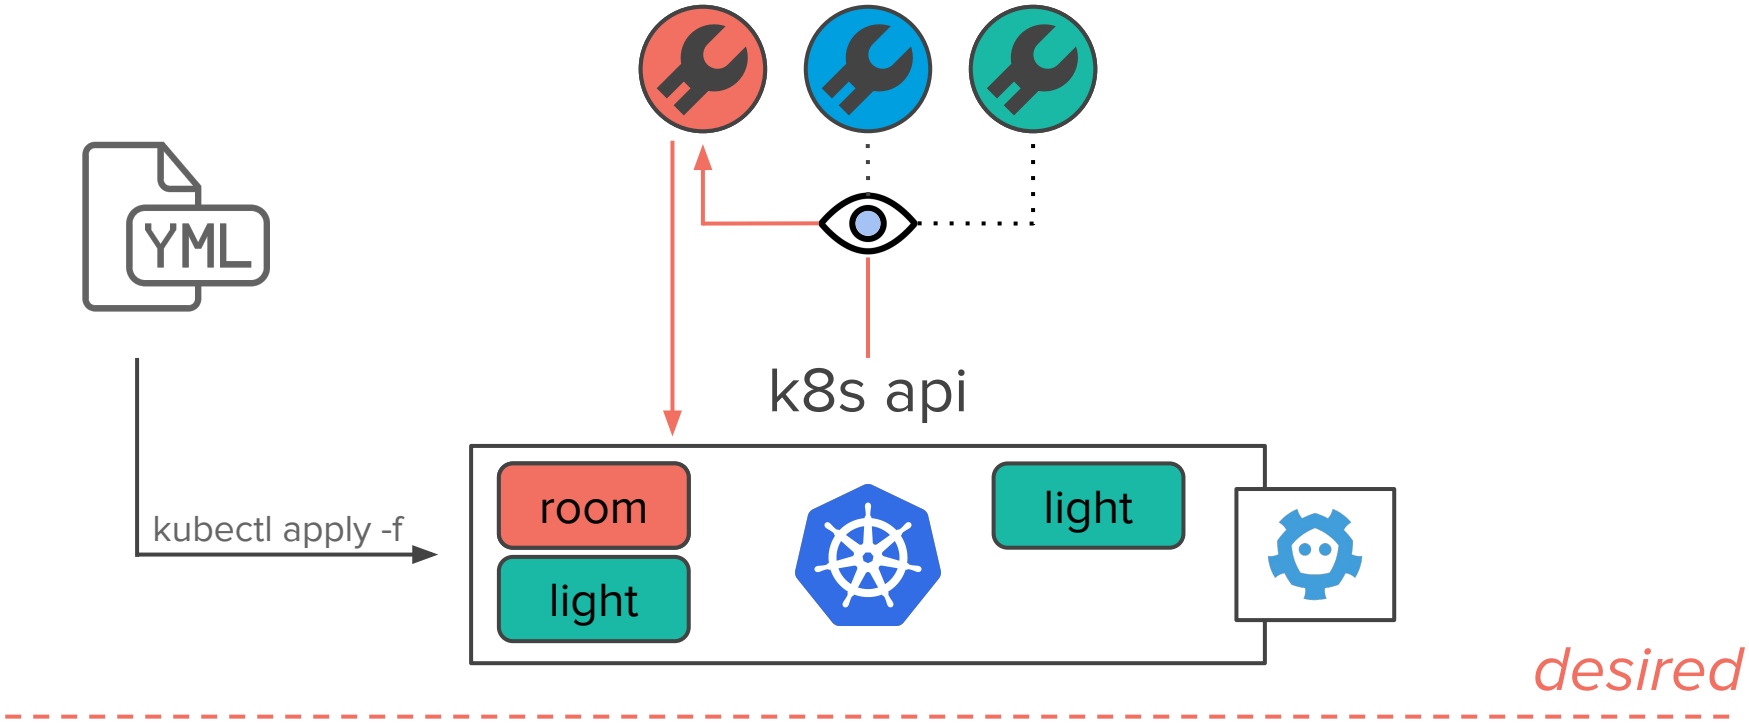
YML
kubectl apply -f
k8s api
room
light
light
desired

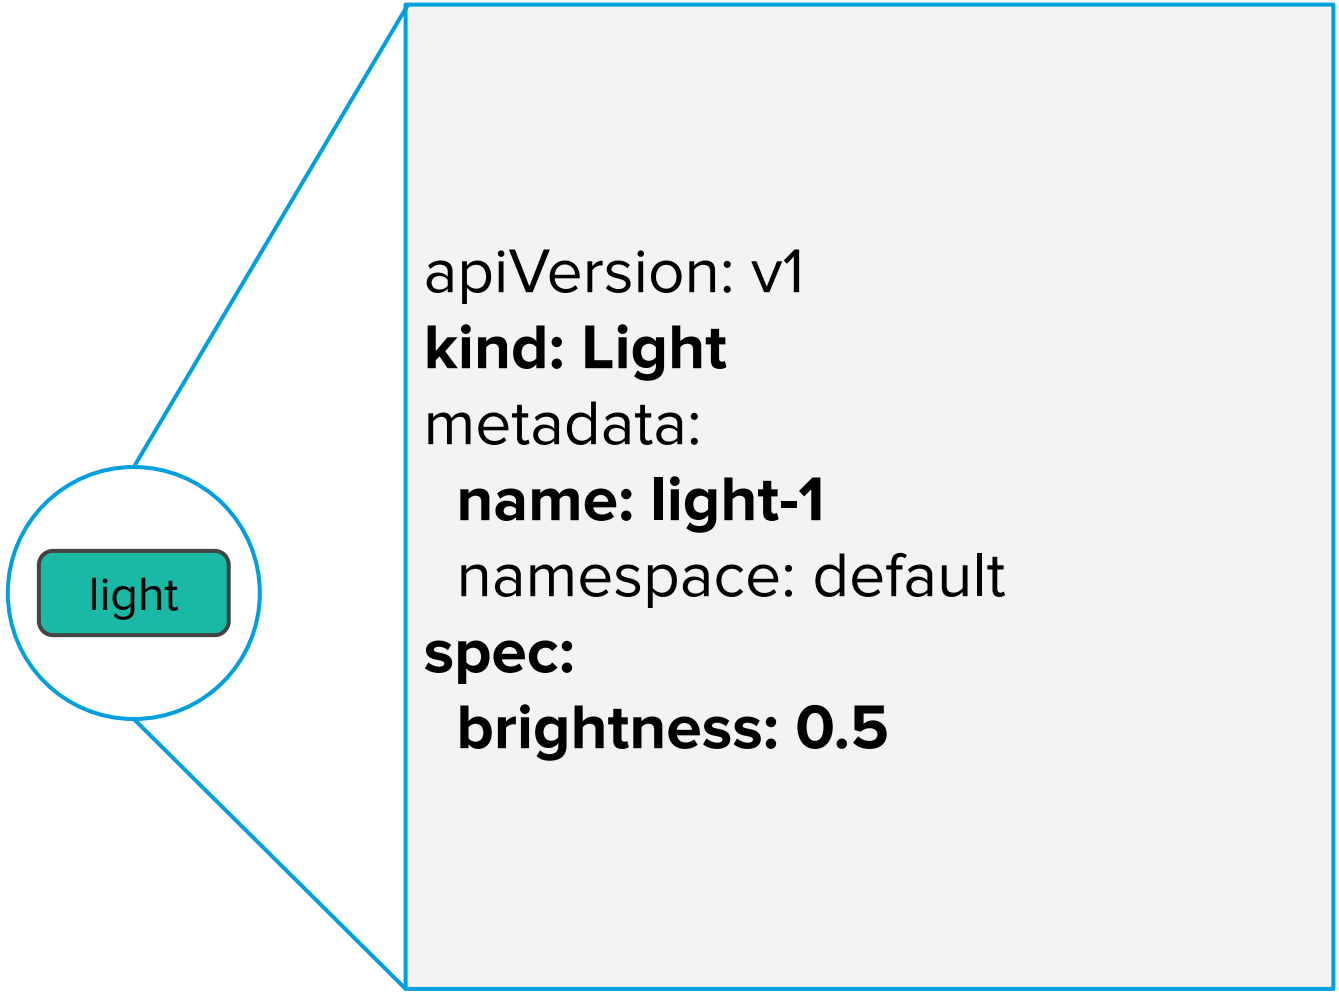
light
apiVersion: v1
kind: Light
metadata:
name: light-1
namespace: default
spec:
brightness: 0.5

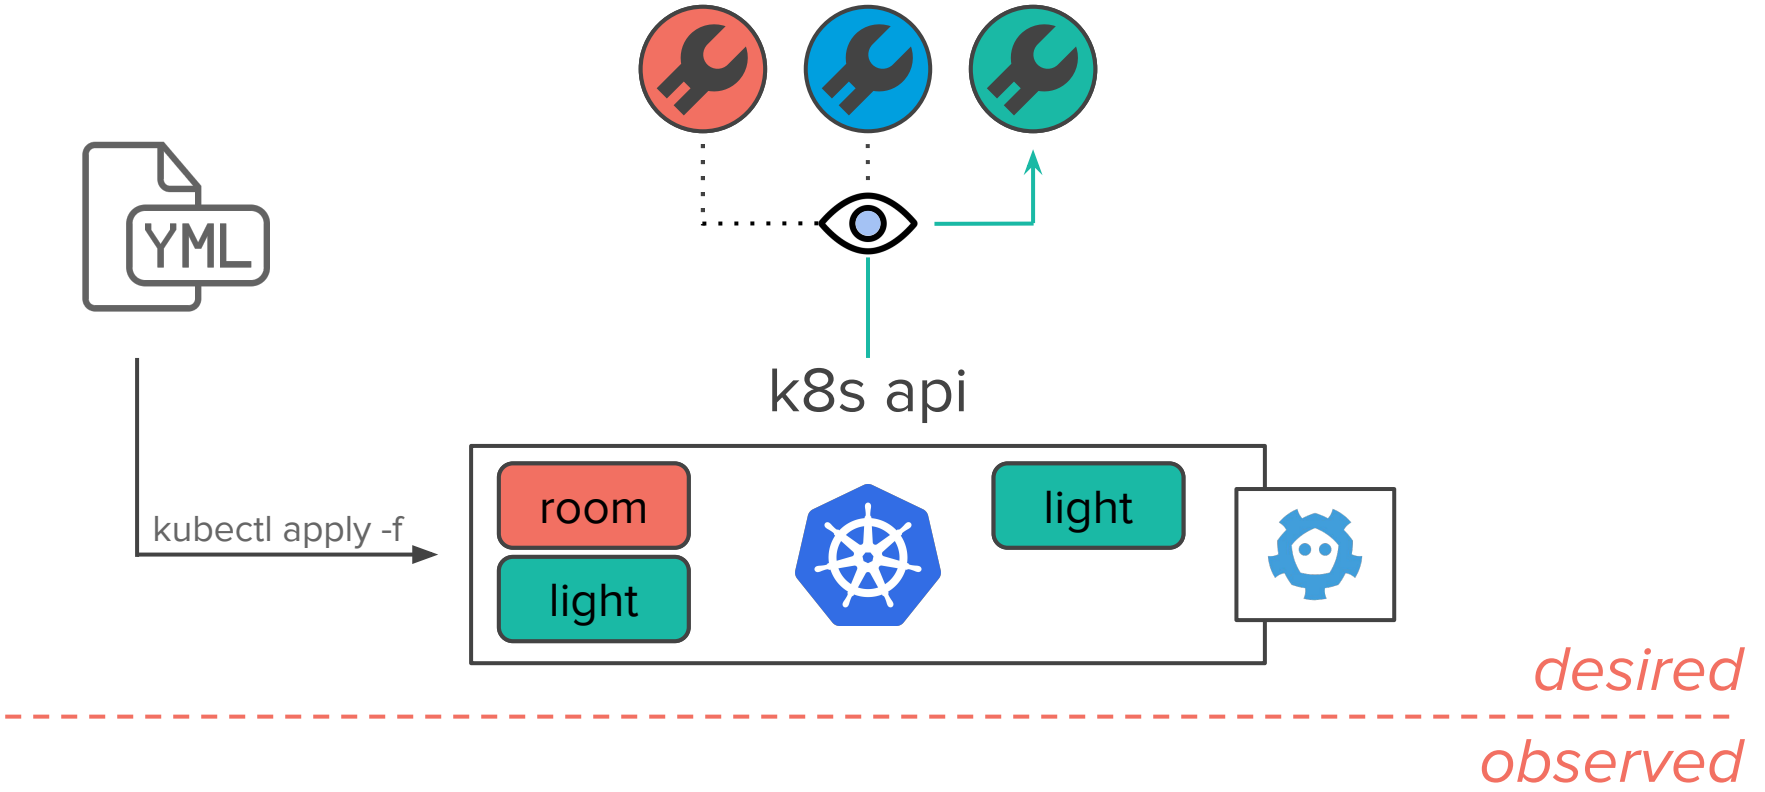
YML
kubectl apply -f
k8s api
room
light
light
desired
observed

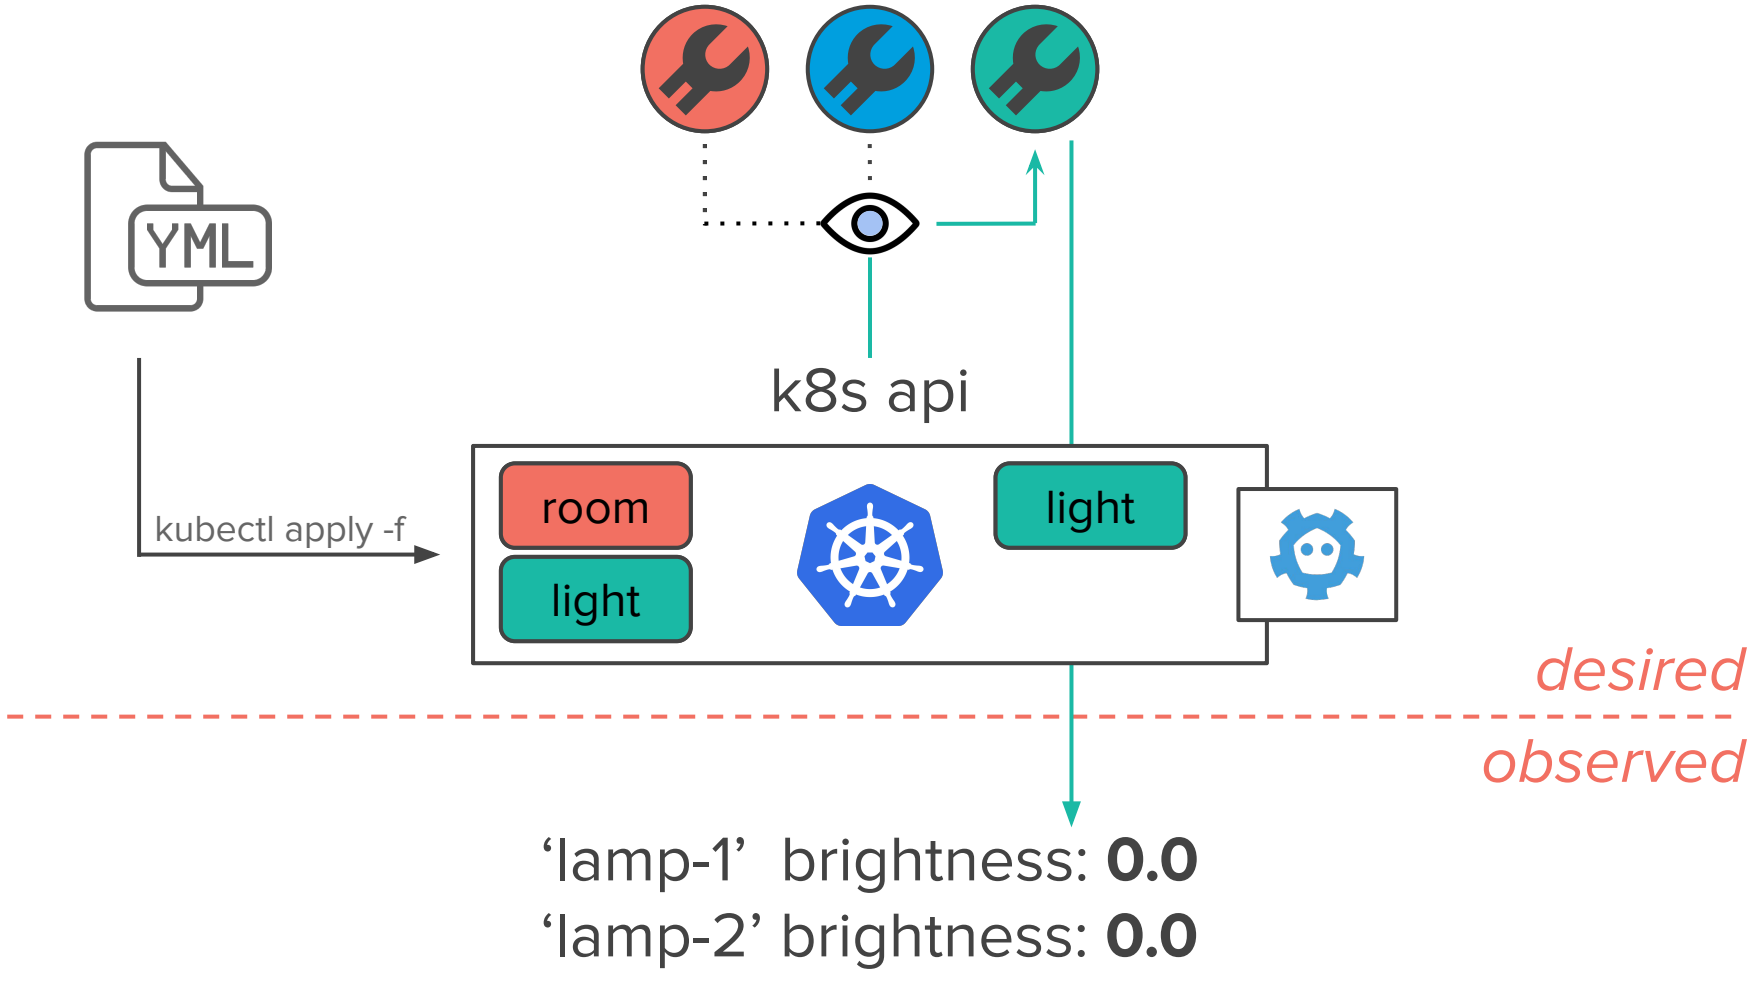
YML
kubectl apply -f
k8s api
room
light
light
desired
observed
‘lamp-1’ brightness: 0.0
‘lamp-2’ brightness: 0.0

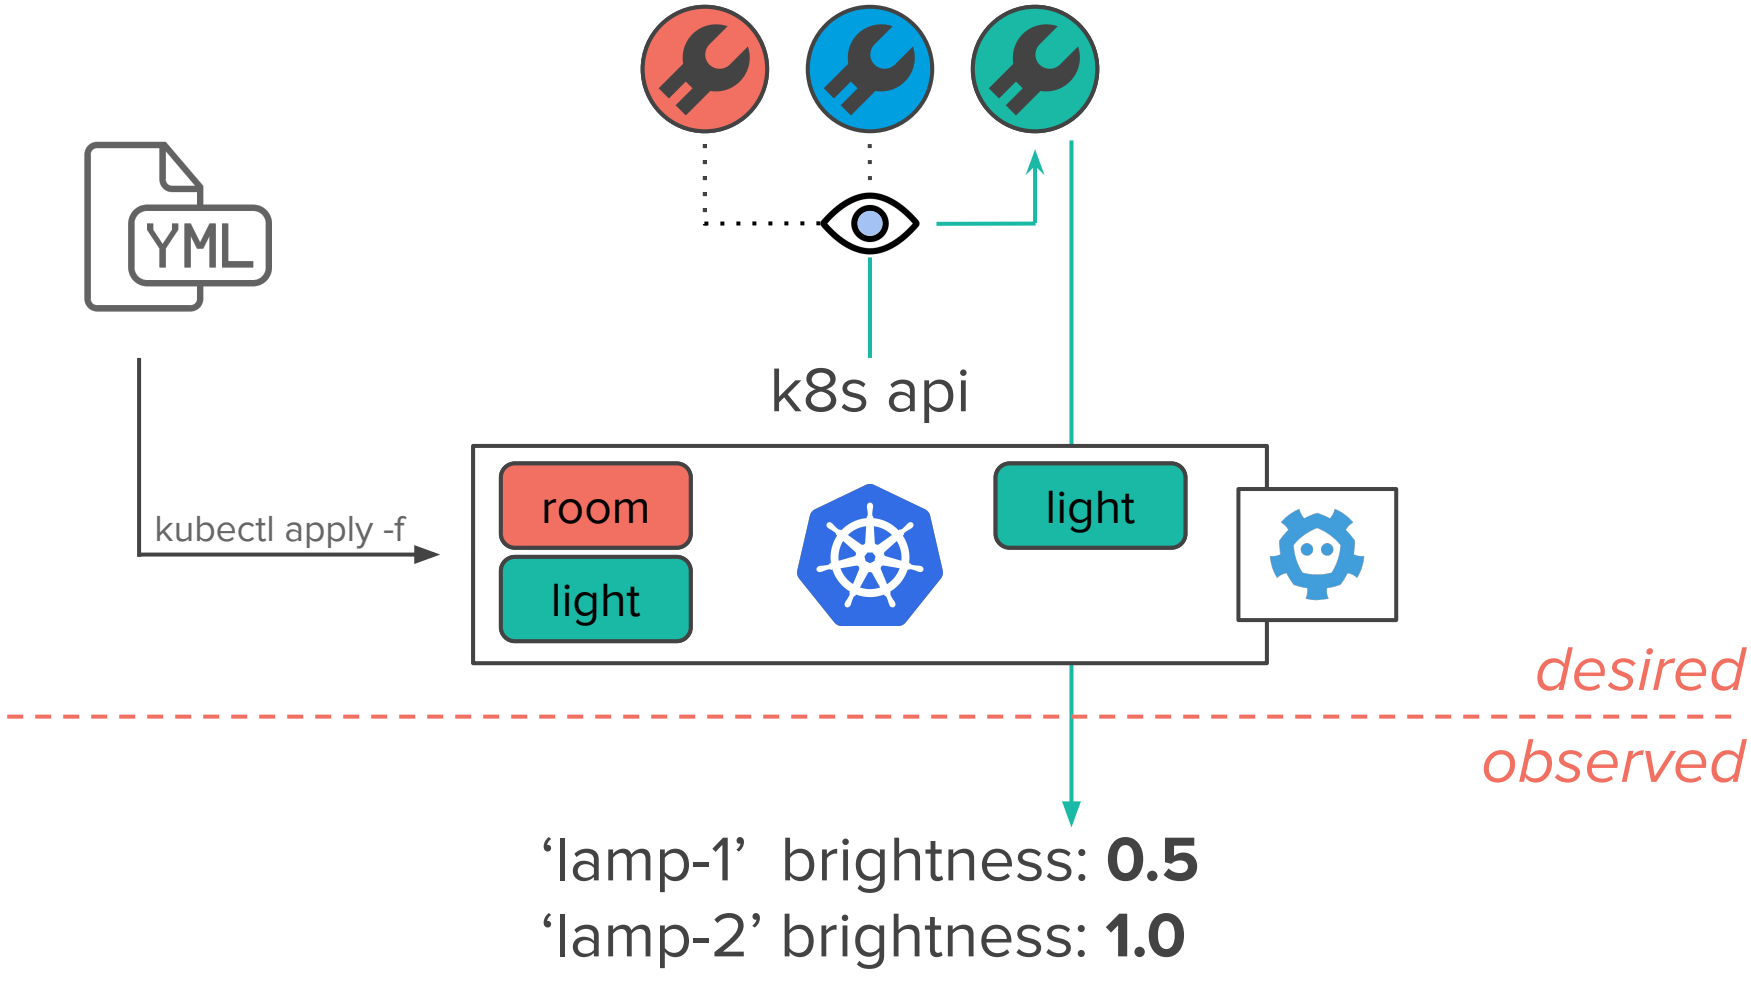
YML
kubectl apply -f
k8s api
room
light
light
desired
observed
‘lamp-1’ brightness: 0.5
‘lamp-2’ brightness: 1.0

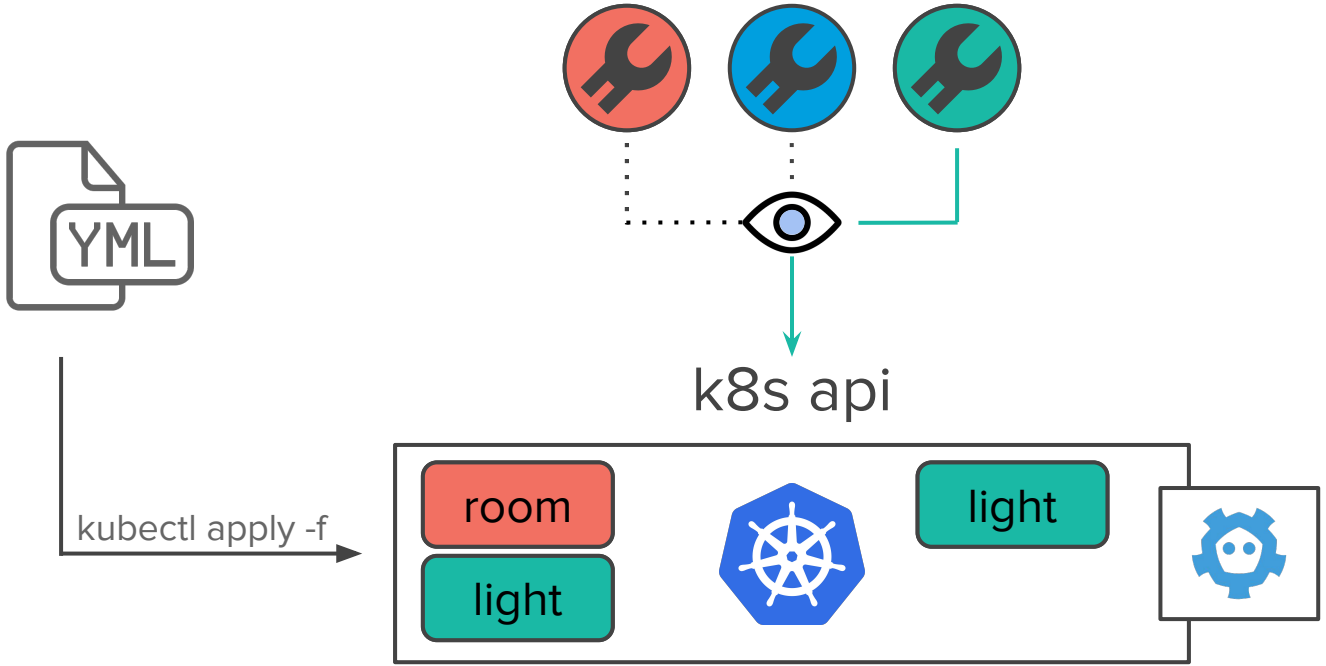

desired

observed

‘lamp-1’ brightness: **0.5**

‘lamp-2’ brightness: **1.0**

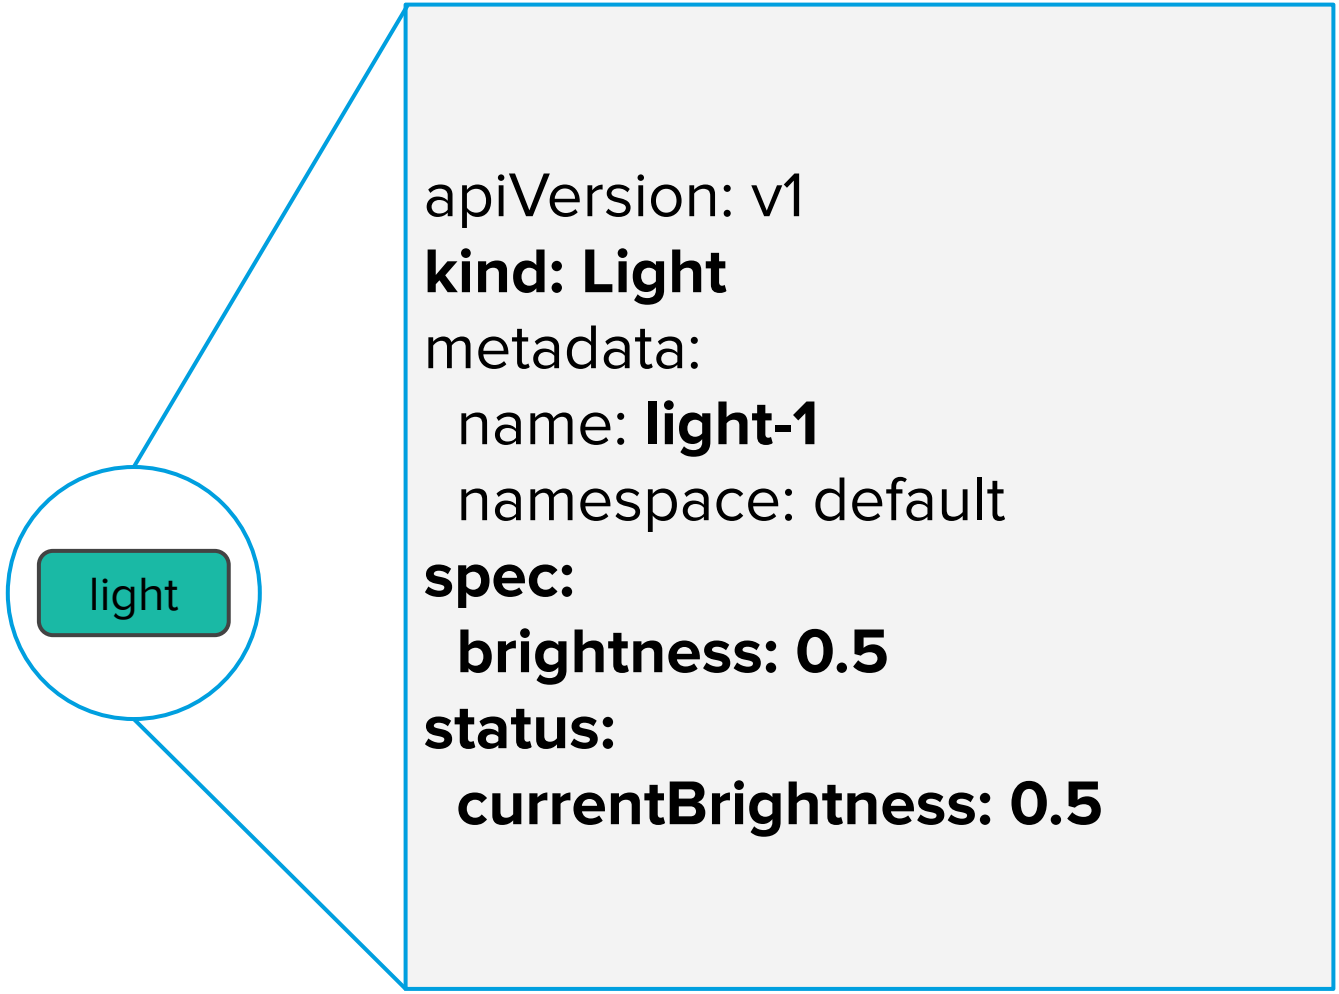
light
apiVersion: v1
kind: Light
metadata:
name: light-1
namespace: default
spec:
brightness: 0.5
status:
currentBrightness: 0.5

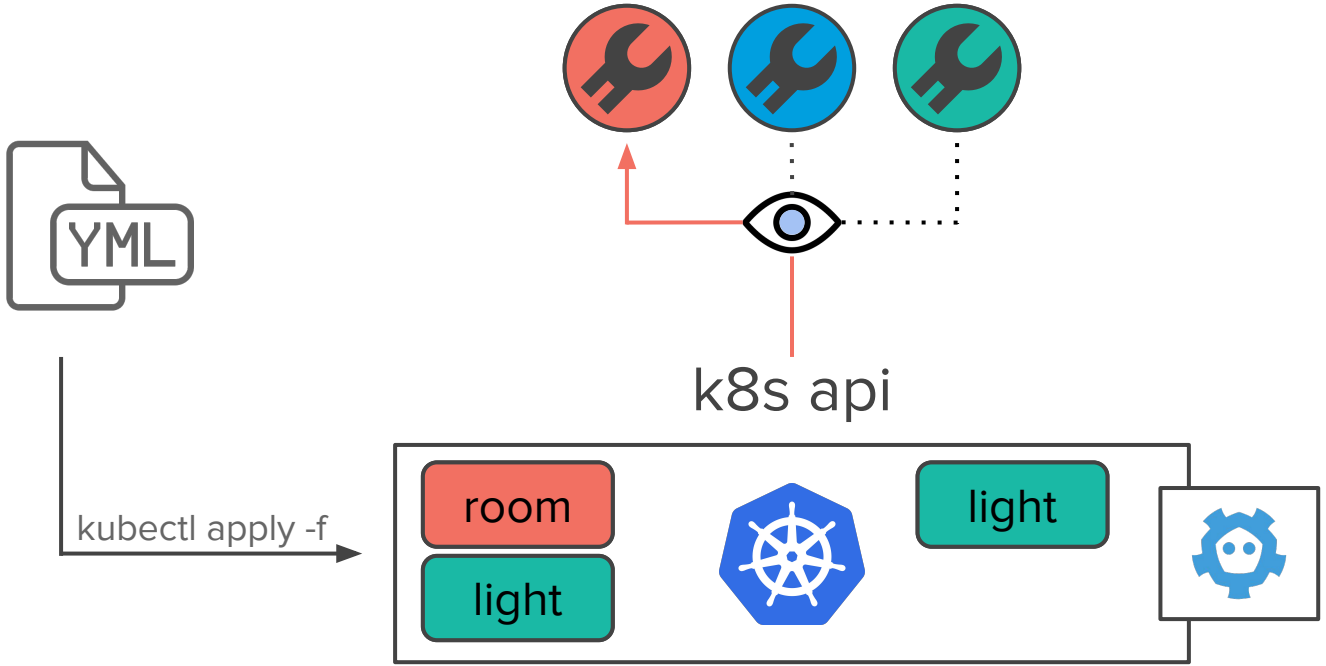

desired

observed

‘lamp-1’ brightness: **0.5**

‘lamp-2’ brightness: **1.0**

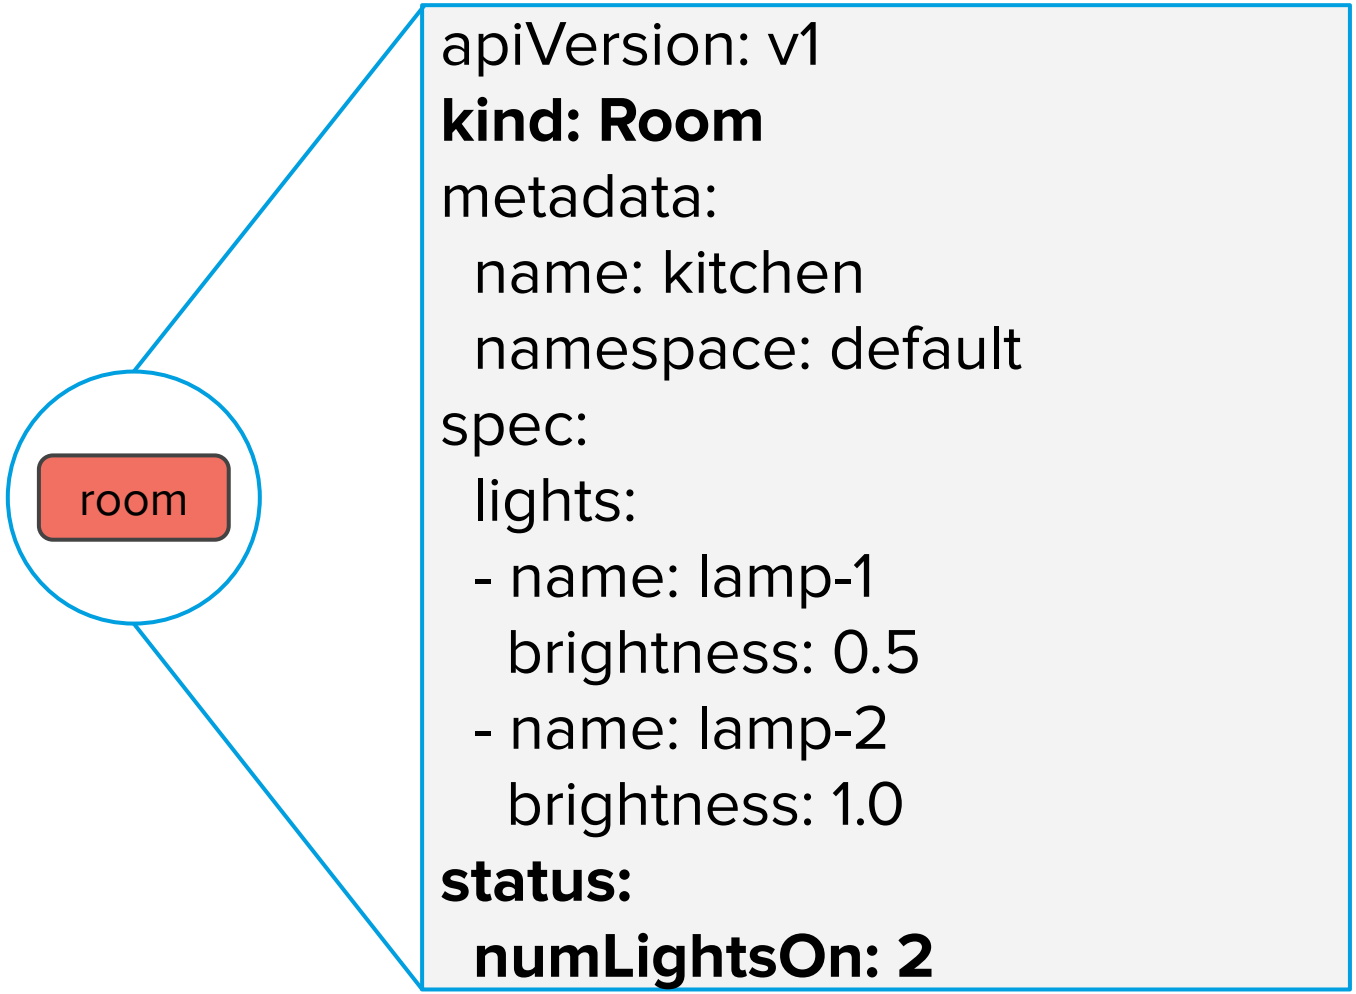
room
apiVersion: v1
kind: Room
metadata:
name: kitchen
namespace: default
spec:
lights:
- name: lamp-1
brightness: 0.5
- name: lamp-2
brightness: 1.0
status:
numLightsOn: 2

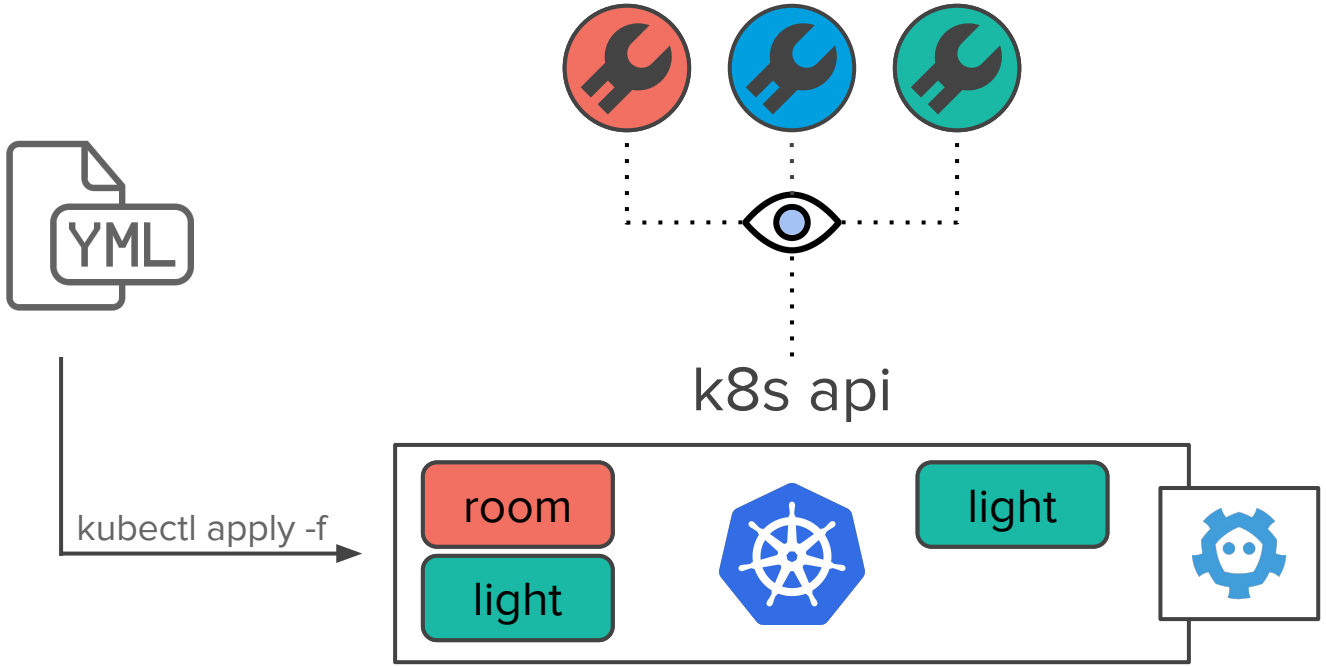

desired

observed

‘lamp-1’ brightness: **0.5**

‘lamp-2’ brightness: **1.0**

Pros, Cons & Considerations

Storage

Access to an etcd datastore for “free”!

- Developers do not need to worry about operational overheads
- … But etcd is not a relational database

High Availability

API is highly available for “free”!

- When deployed in a multi-master configuration, the k8s API is Highly Available
- Developers do not need to worry about operational overheads!

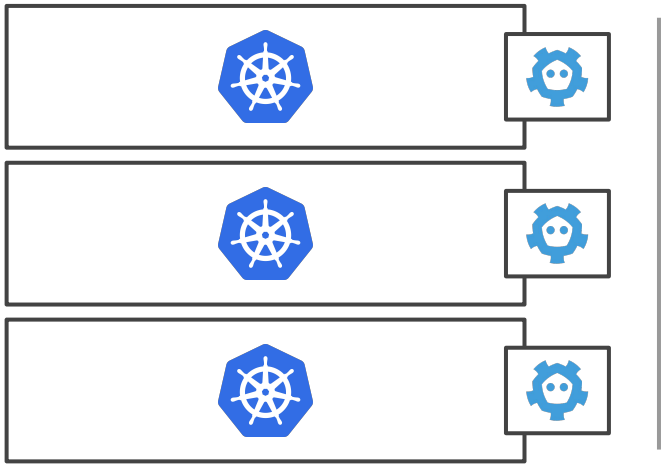

multi-master

Performance

Performance is largely dependent on etcd

- “Noisy” services could have a negative impact on performance
- 1 huge k8s cluster vs lots of little ones
  - Perhaps not so much of an issue in smaller clusters
- API machinery are thinking about scaling targets
  - Fill in the survey!

Programming Model

Declarative vs Imperative

- The k8s API is entirely declarative and eventually consistent
- Great for stability and reliability of the system!
- Writing reconciliation
- Not everything fits!

Team Collaboration

CRDs as the standard interface between teams

- Teams ship controllers and CRDs
- Team A's controller could watch for changes to Team B's Custom Resources
- Part of what it means to be "kubernetes-native"

Other API Features

What other API features should you consider?

- AuthN/Z
- Pagination
- Querying
- Binary Data
- Resource relationship
- Versioning
- Quotas
- Tooling (kubectl/UIs etc)

To CRD, or Not to CRD?

Kelsey Hightower
@kelseyhightower

To the entire Kubernetes community:

> 
> **Abby Fuller** @abbyfuller · Jun 18
> I don't know who needs to hear this but not everything needs to be a Kubernetes CRD

3:12 AM · Jun 18, 2019 · Twitter for iPhone

86 Retweets **398** Likes

Thank You!